

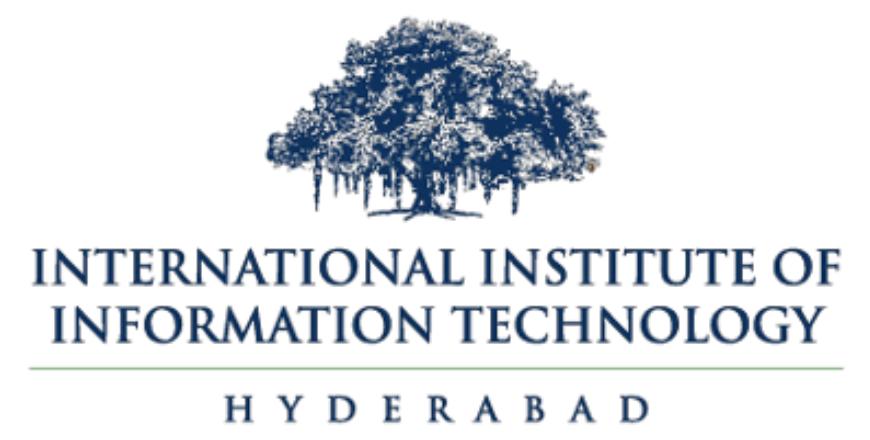
CS3.301 Operating Systems and Networks

OS and Networks: Run Down and Concluding Thoughts!

Karthik Vaidhyanathan

<https://karthikvaidhyanathan.com>

1



Acknowledgement

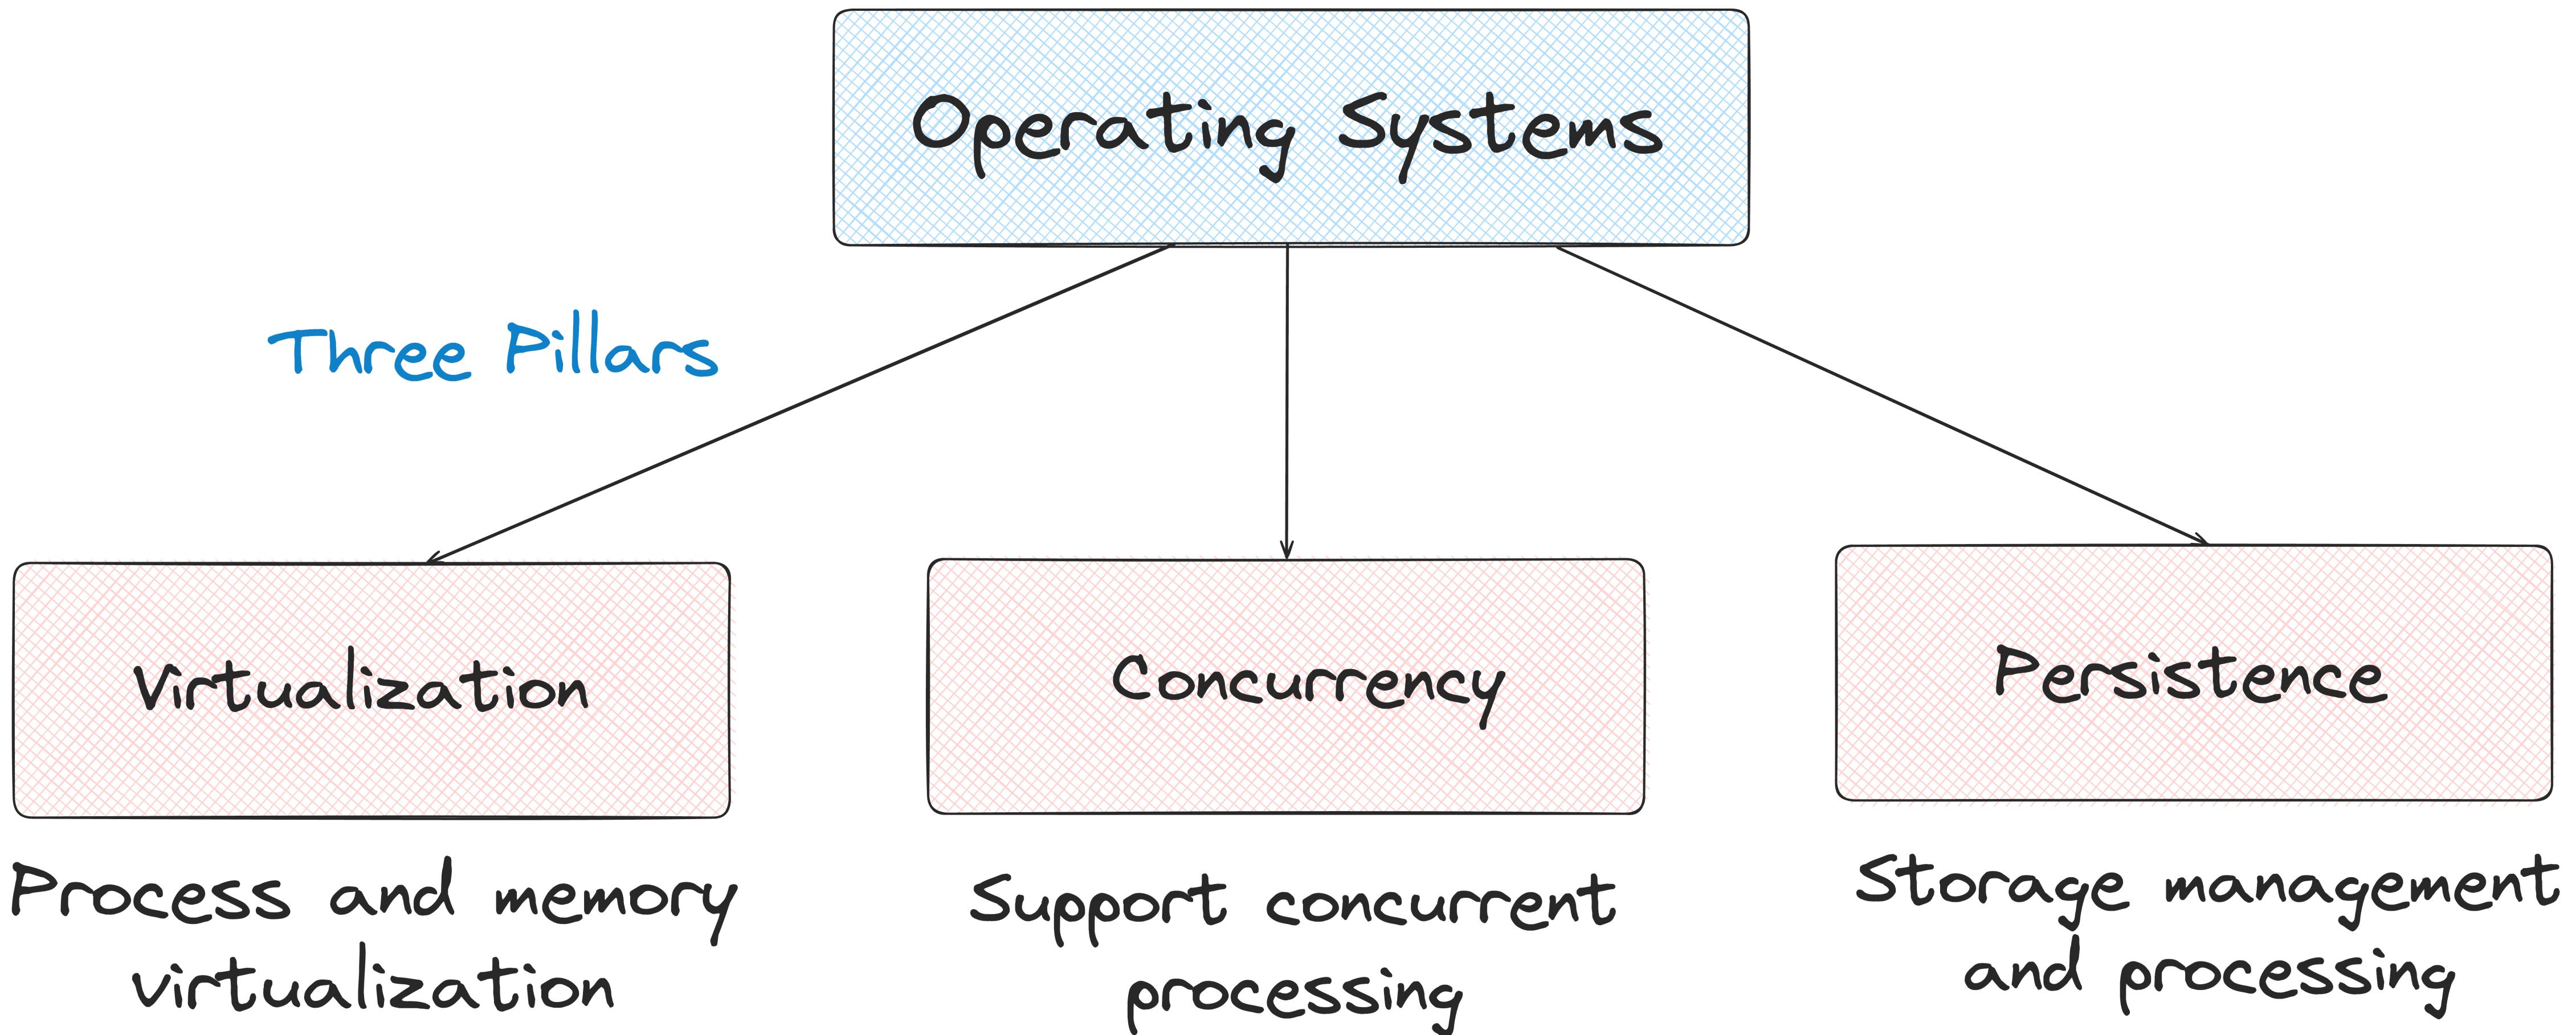
The materials used in this presentation have been gathered/adapted/generate from various sources as well as based on my own experiences and knowledge -- Karthik Vaidhyanathan

Sources:

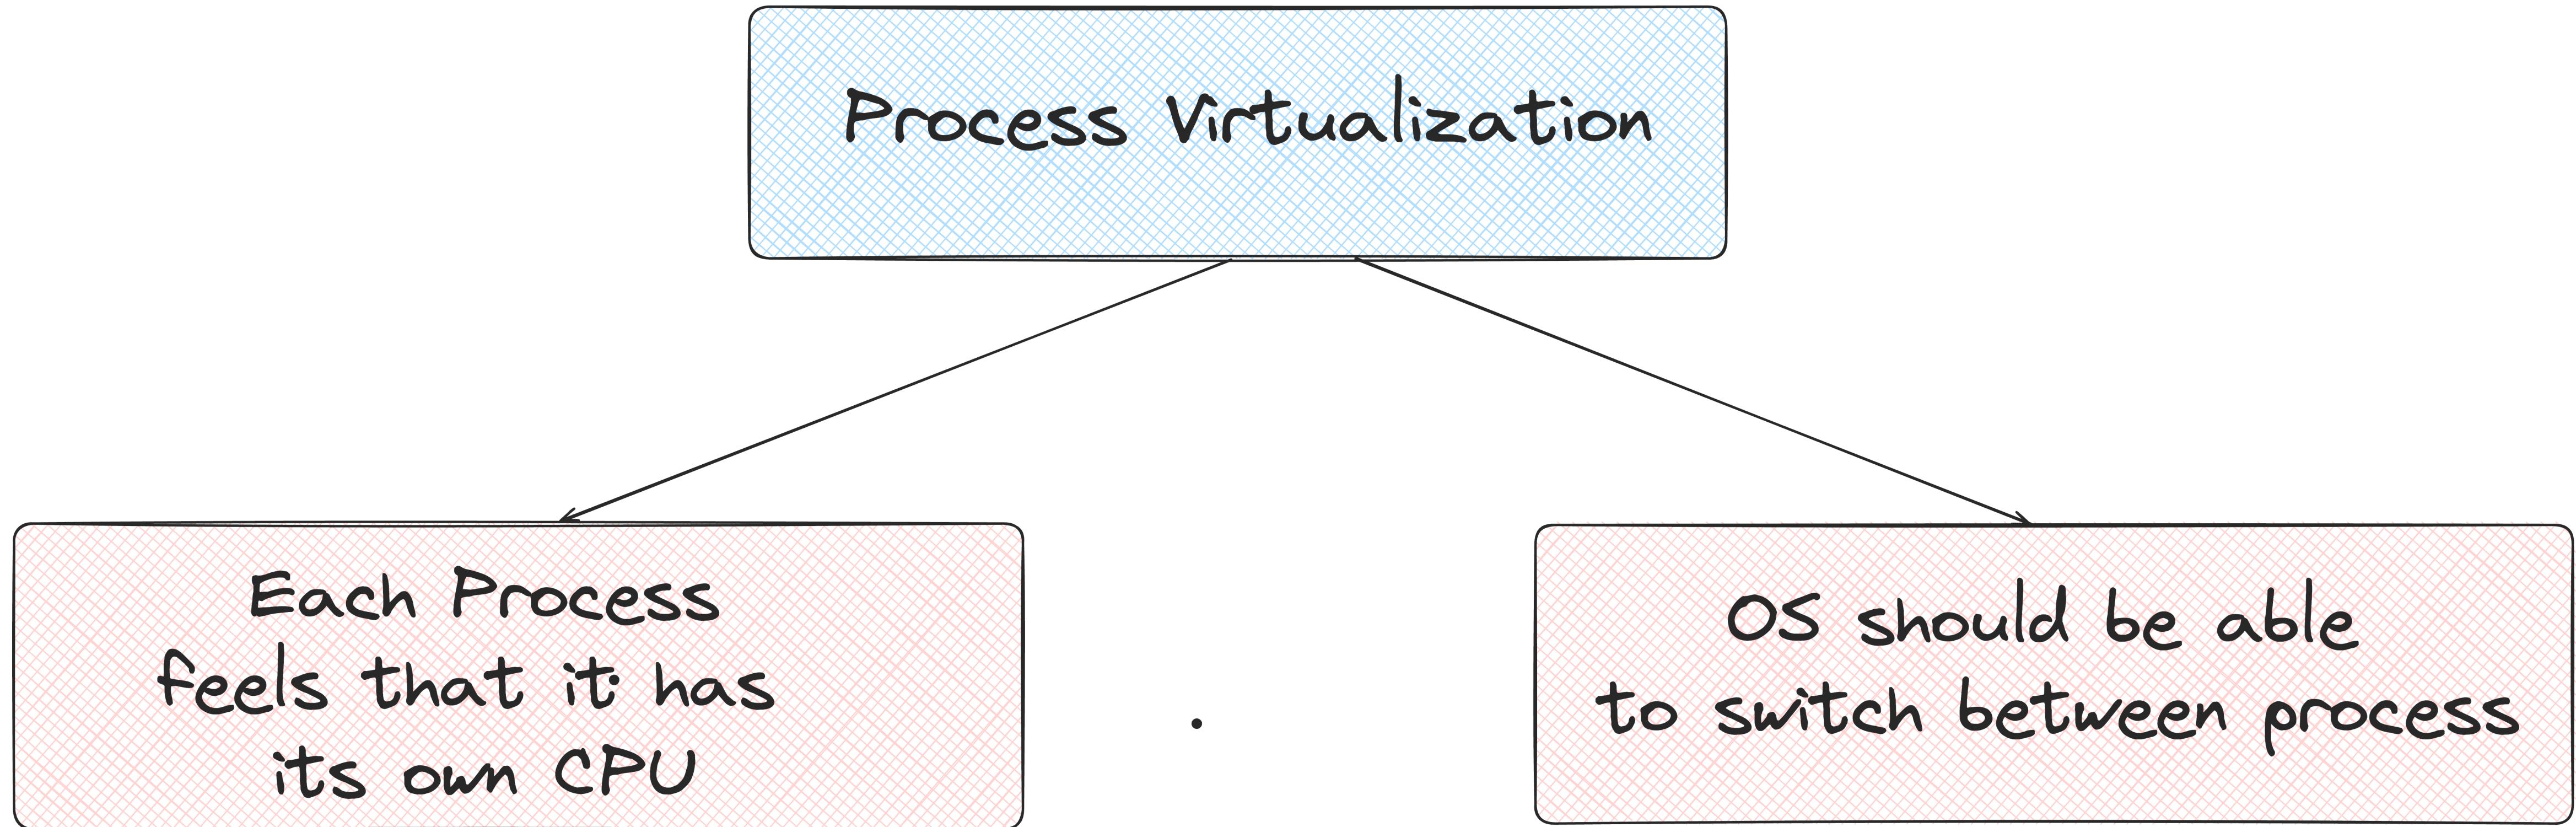
- Operating Systems in Three Easy Pieces by Remzi et al.
- Computer Networks, 6e by Tanbaum, Teamster and Wetherall
- Computer Networks: A Top Down Approach by Kurose and Ross
- Computer Networking essentials, Youtube Channel
- Other online sources which are duly cited
- Different materials used throughout the course!



OS: An Overview

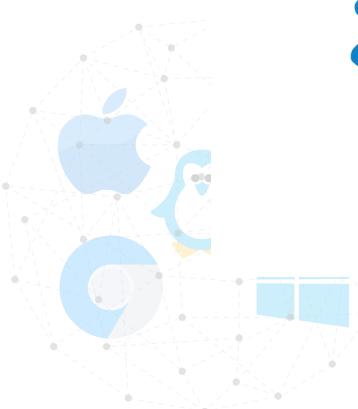


Process Virtualization



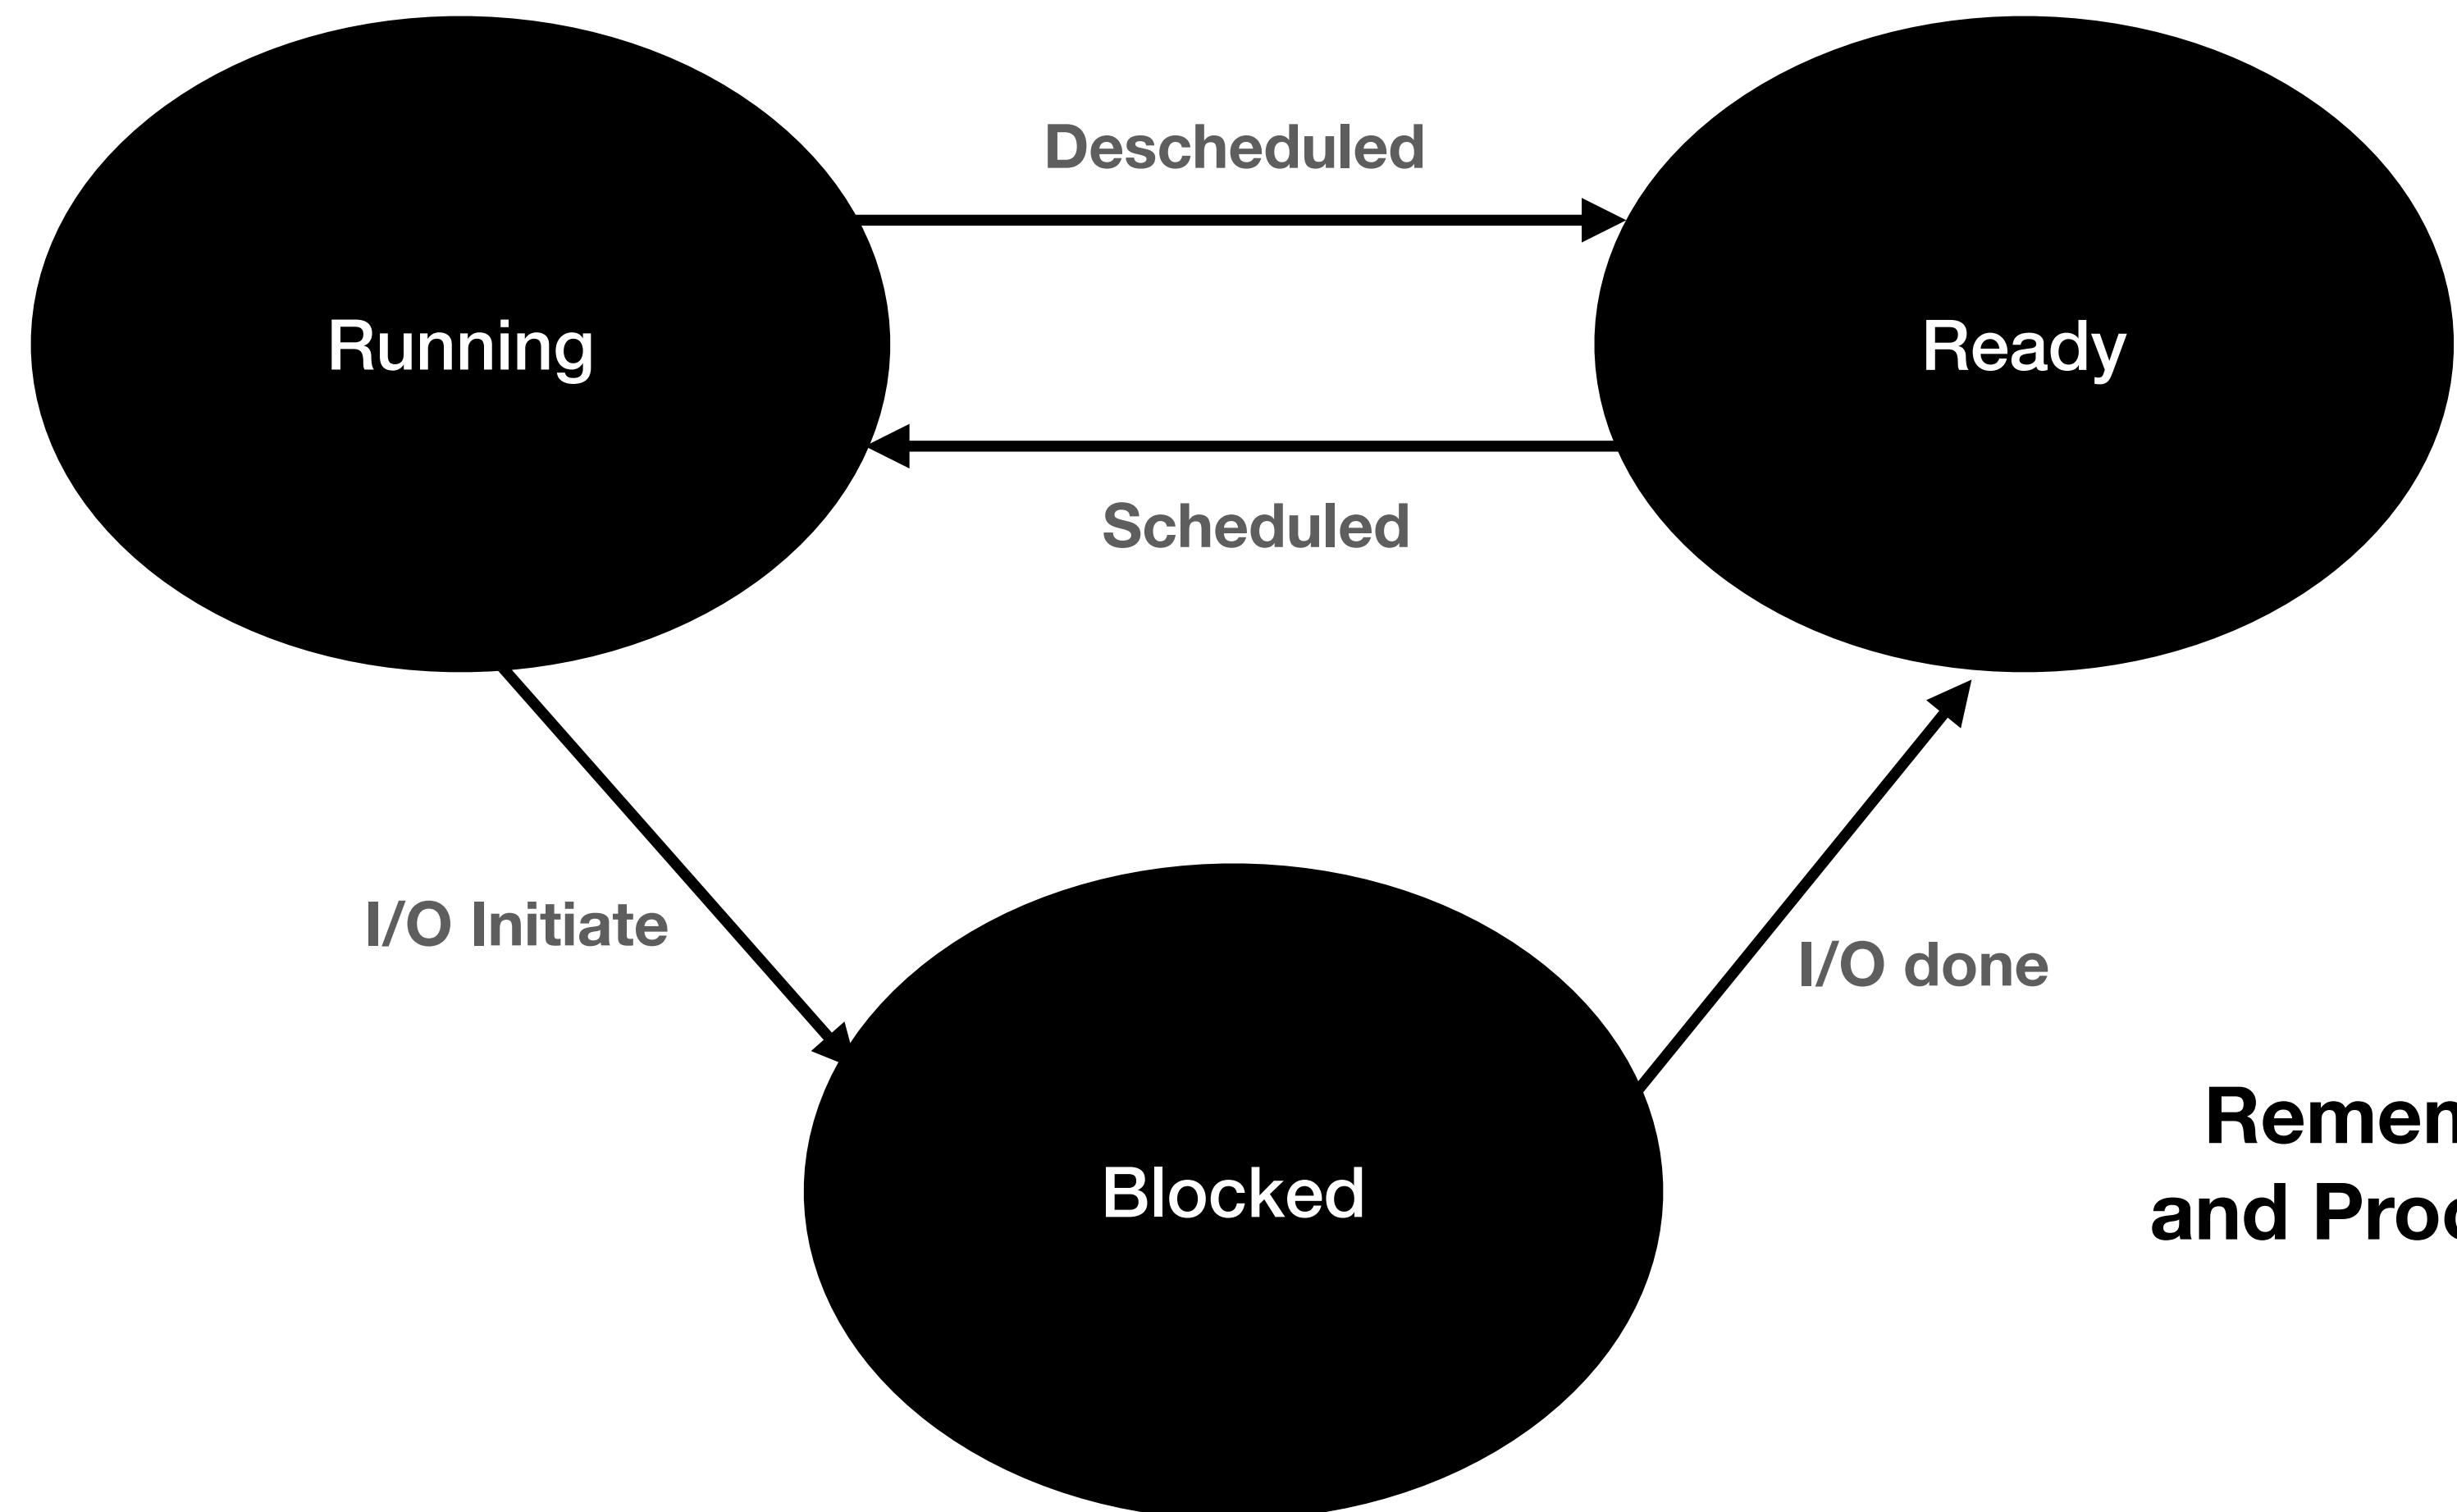
- 1. Process has different states
- 2. Process management API

- 1. Support for context switch
- 2. Process Scheduling



States of the Process

Process State Transitions

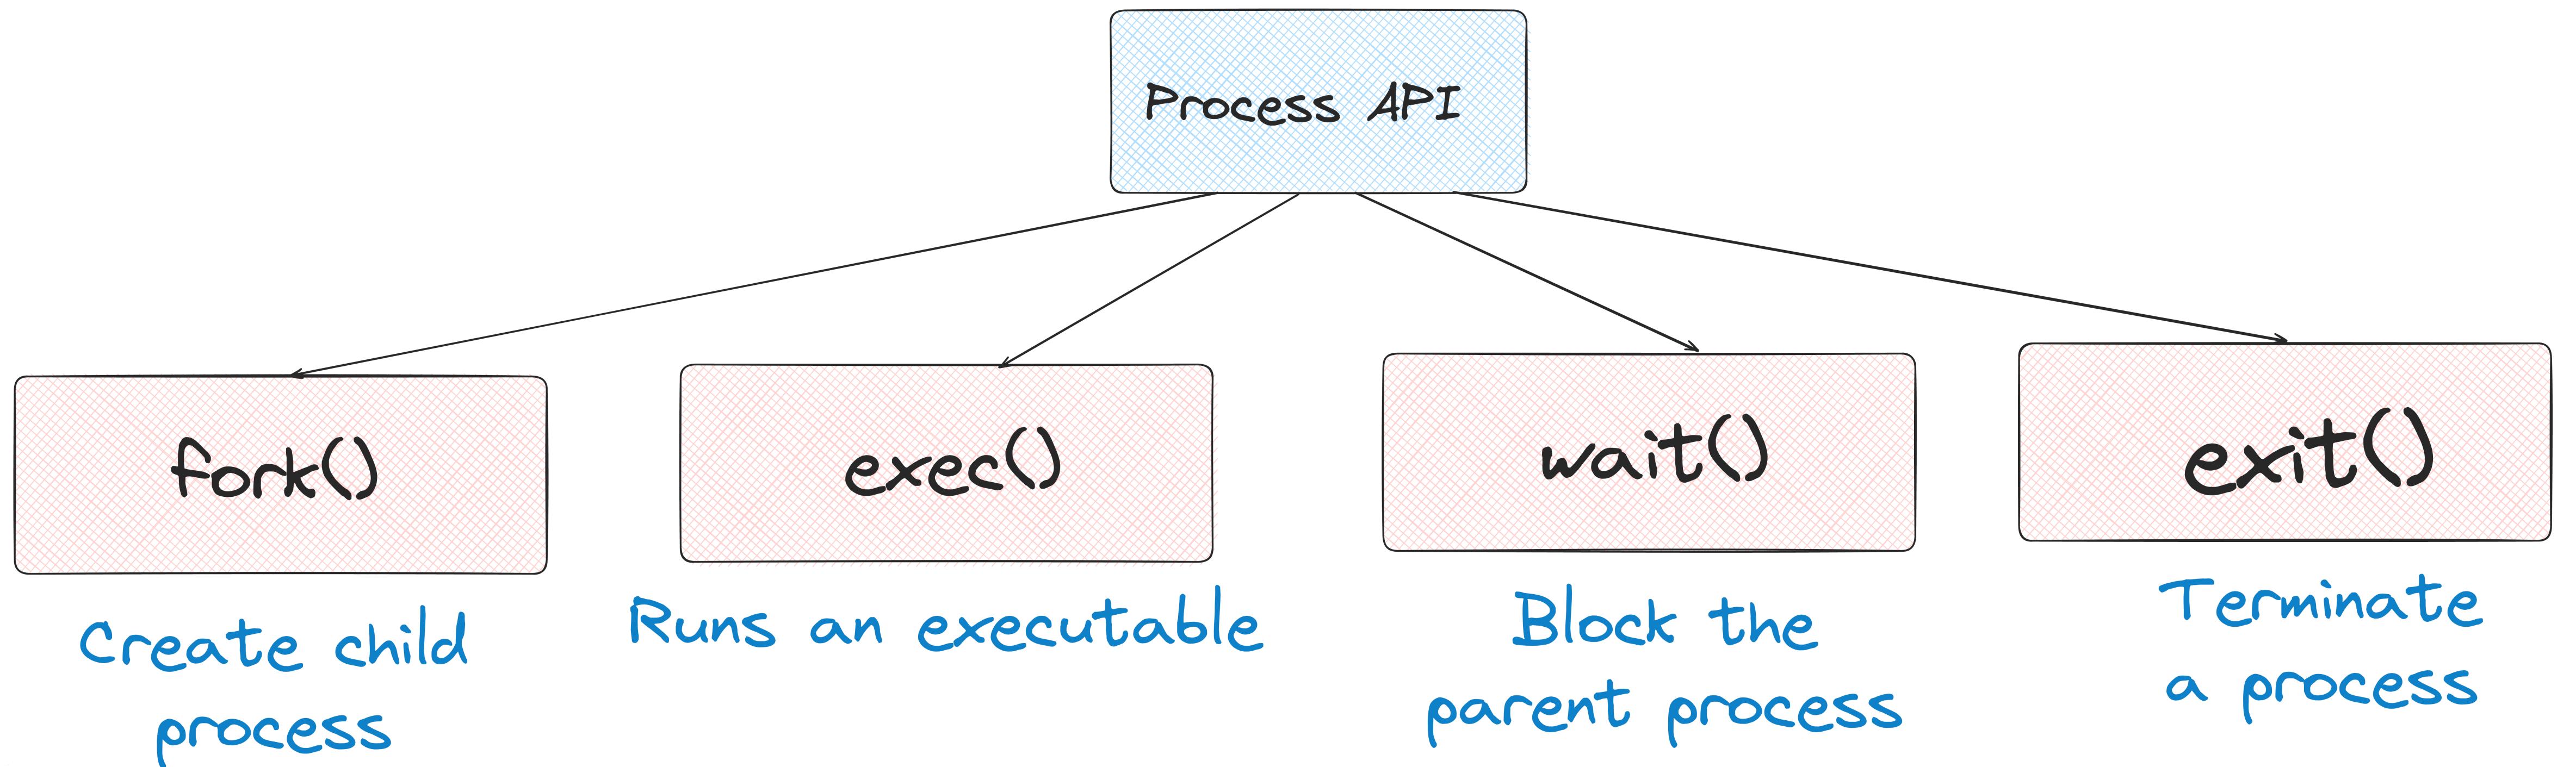


**Remember: Process lists
and Process Control Blocks**

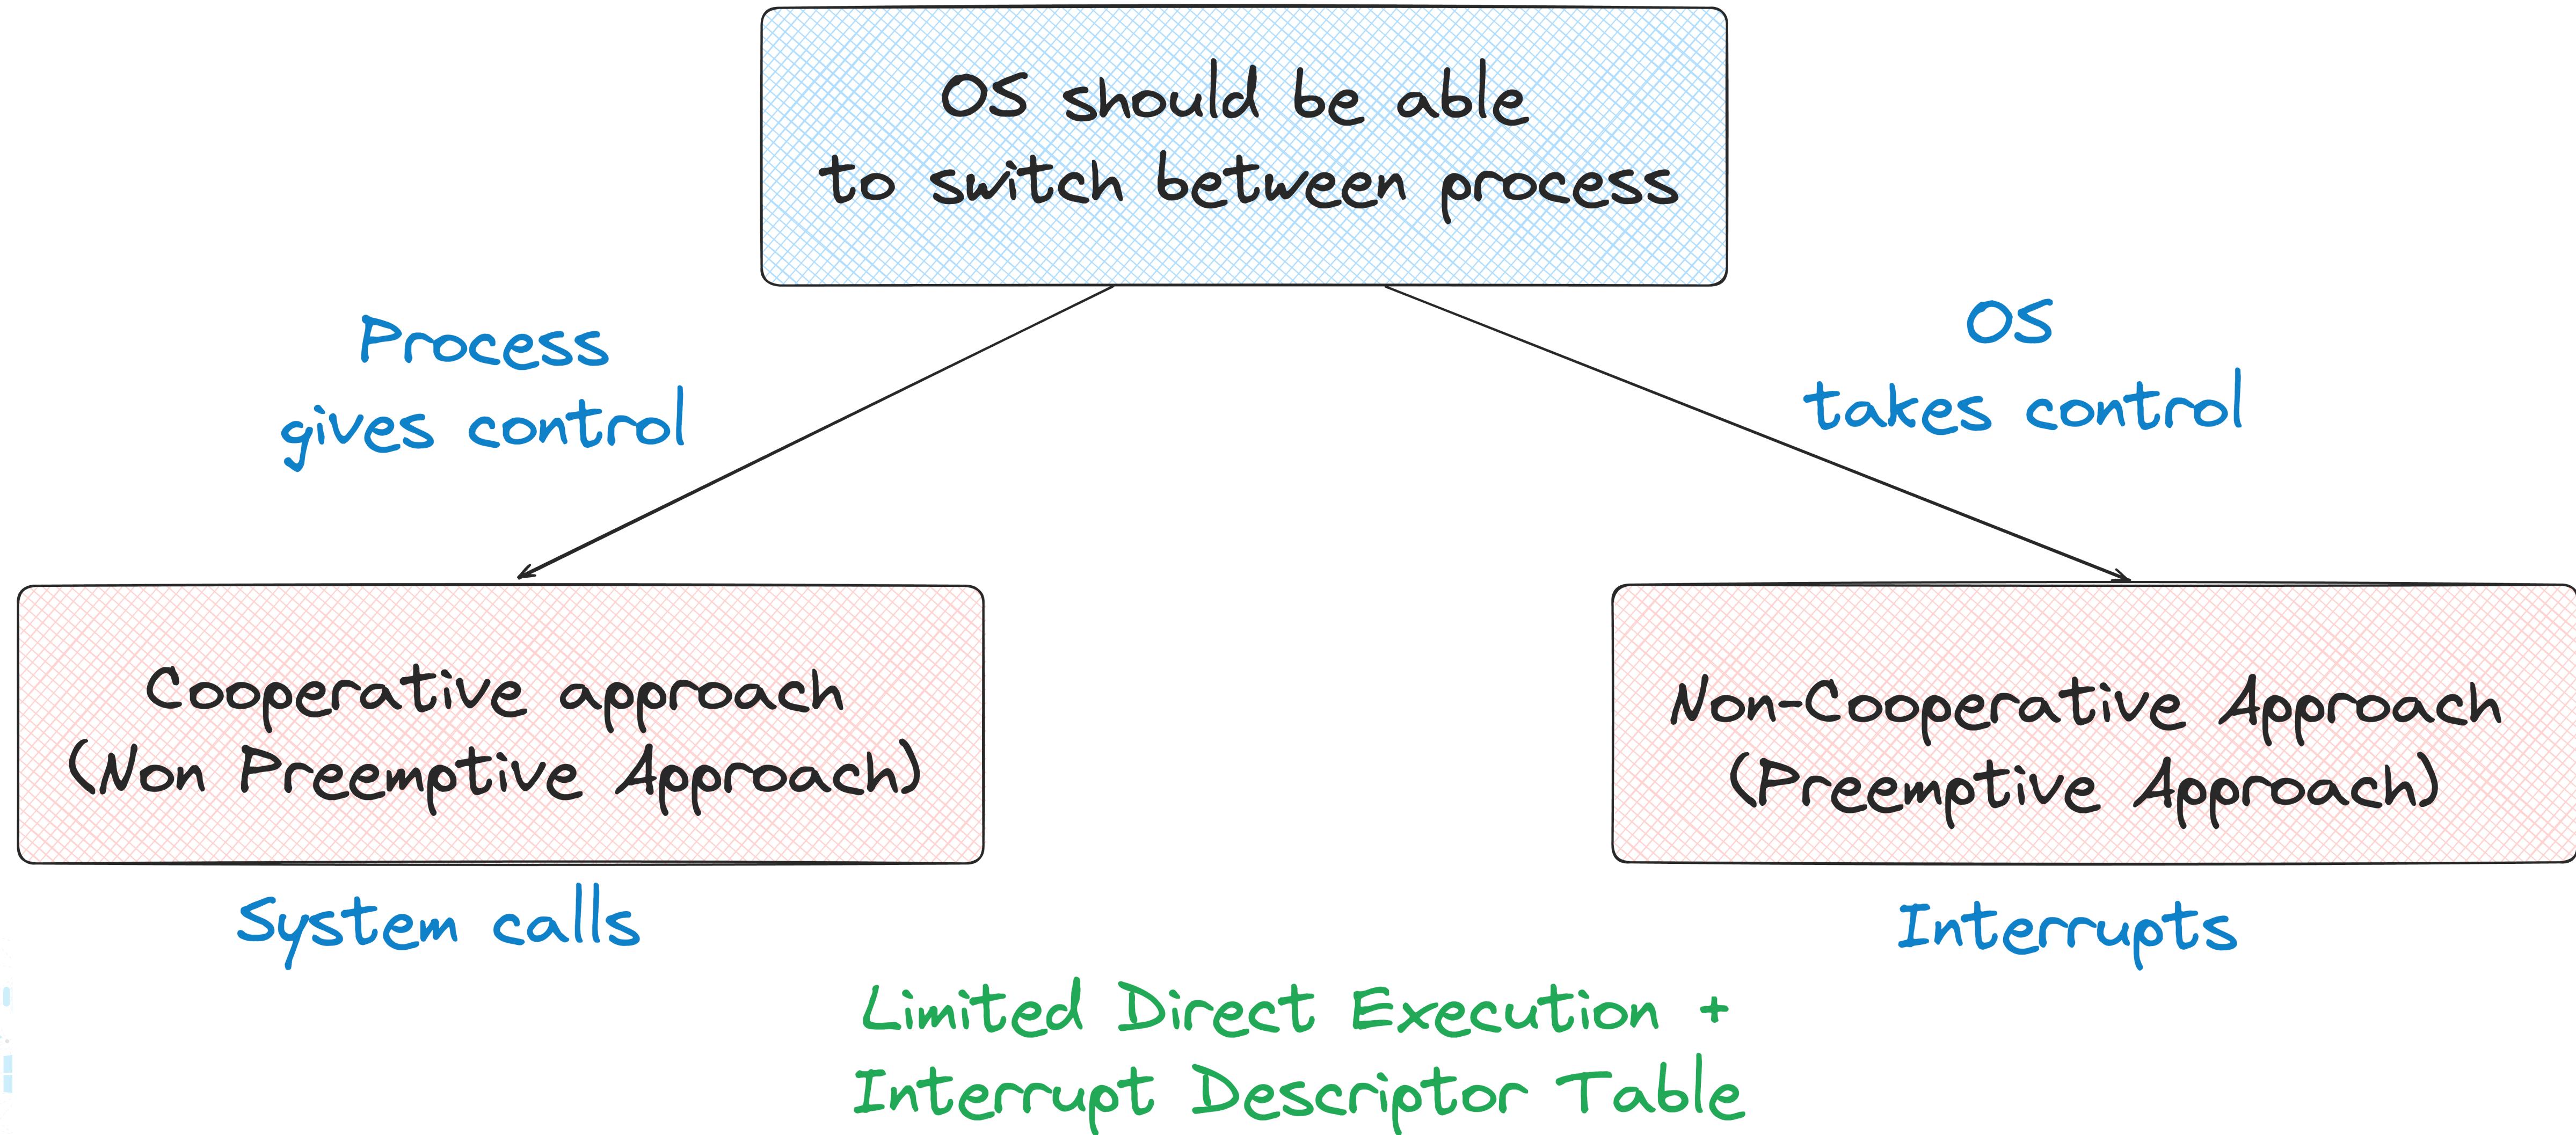


Process Management API

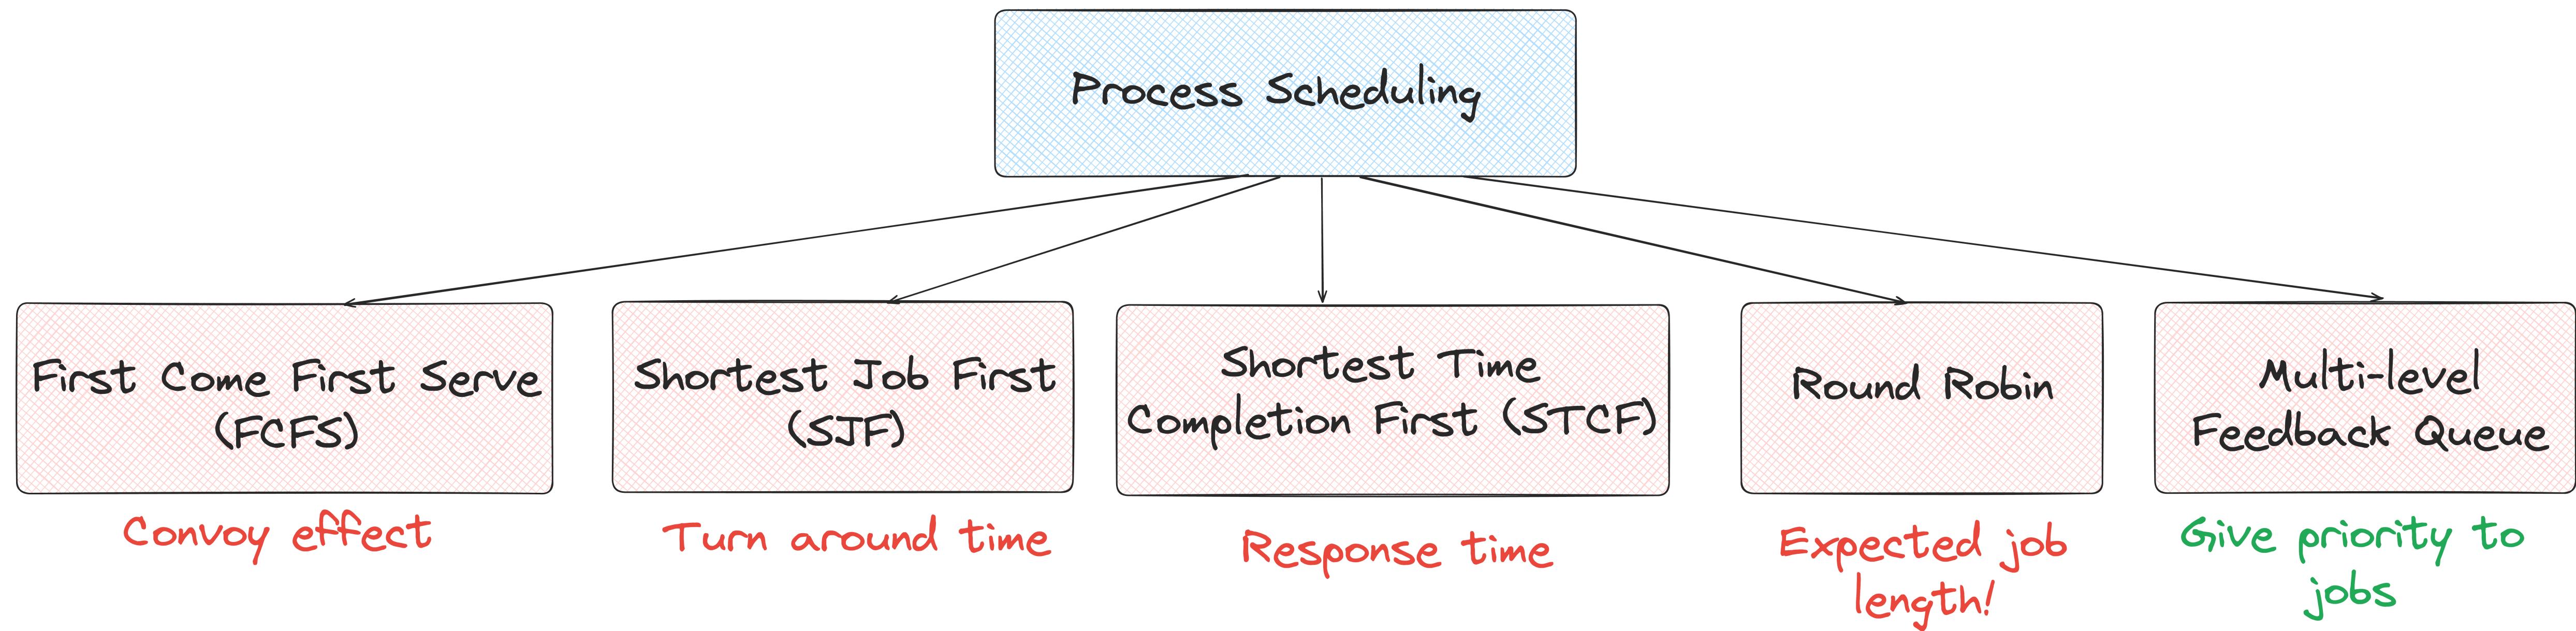
Memory image of a process - **Code, data, stack, and heap**



Switching Between Process



Process Scheduling

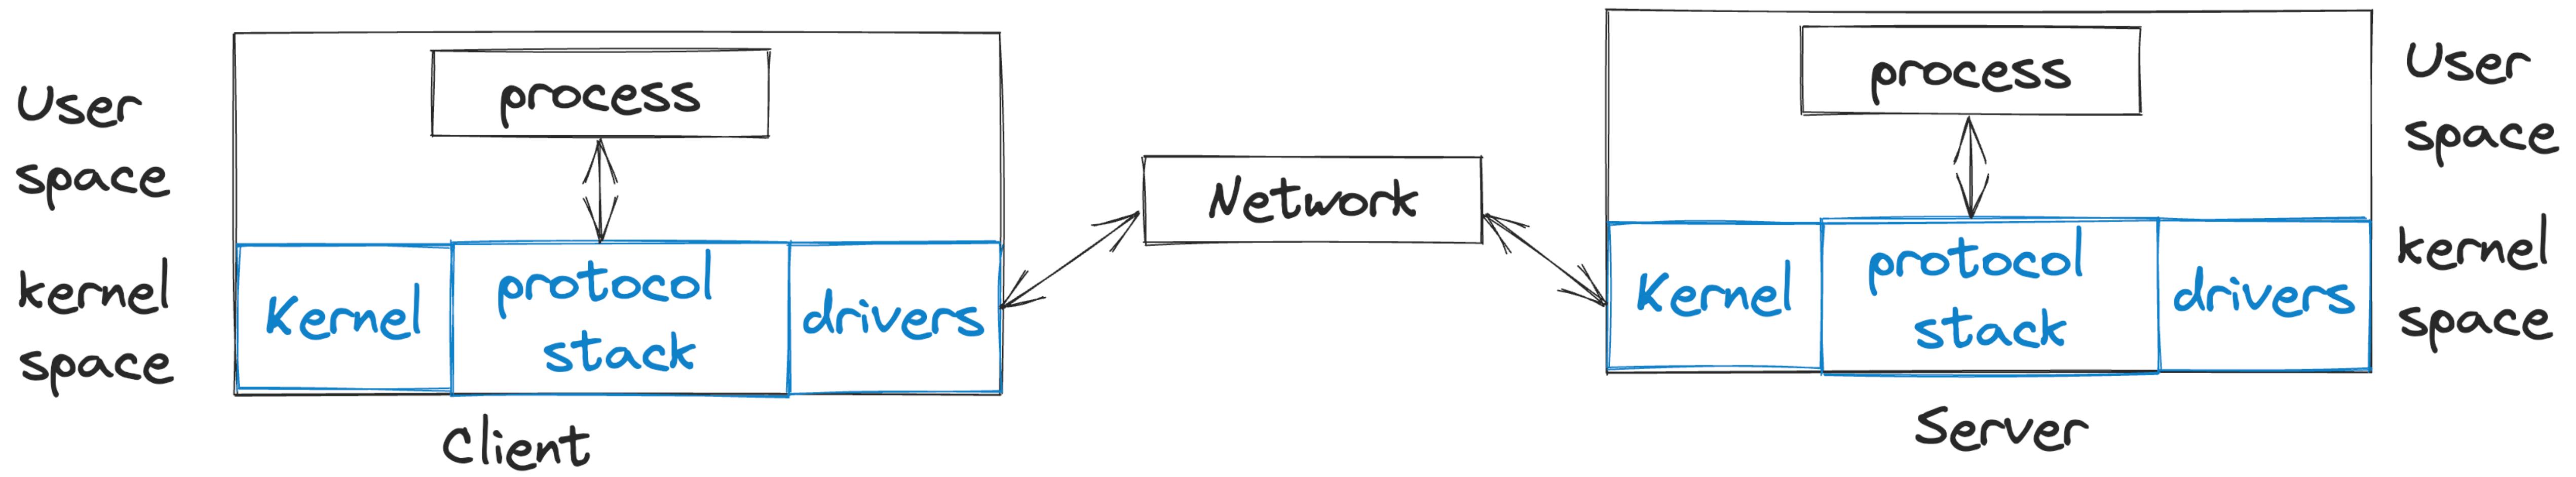


What if Process wants to communicate?

- Process A (eg: Whatsapp) is executing in Host 1
 - Process B (Whatsapp) is executing in Host 2
- Host 1 will have an address, same is the case with host 2
- How to ensure the data reaches from Host 1 to Host 2?
 - What all needs to be considered?
 - Remember: There will be multiple processes that are executing in a host



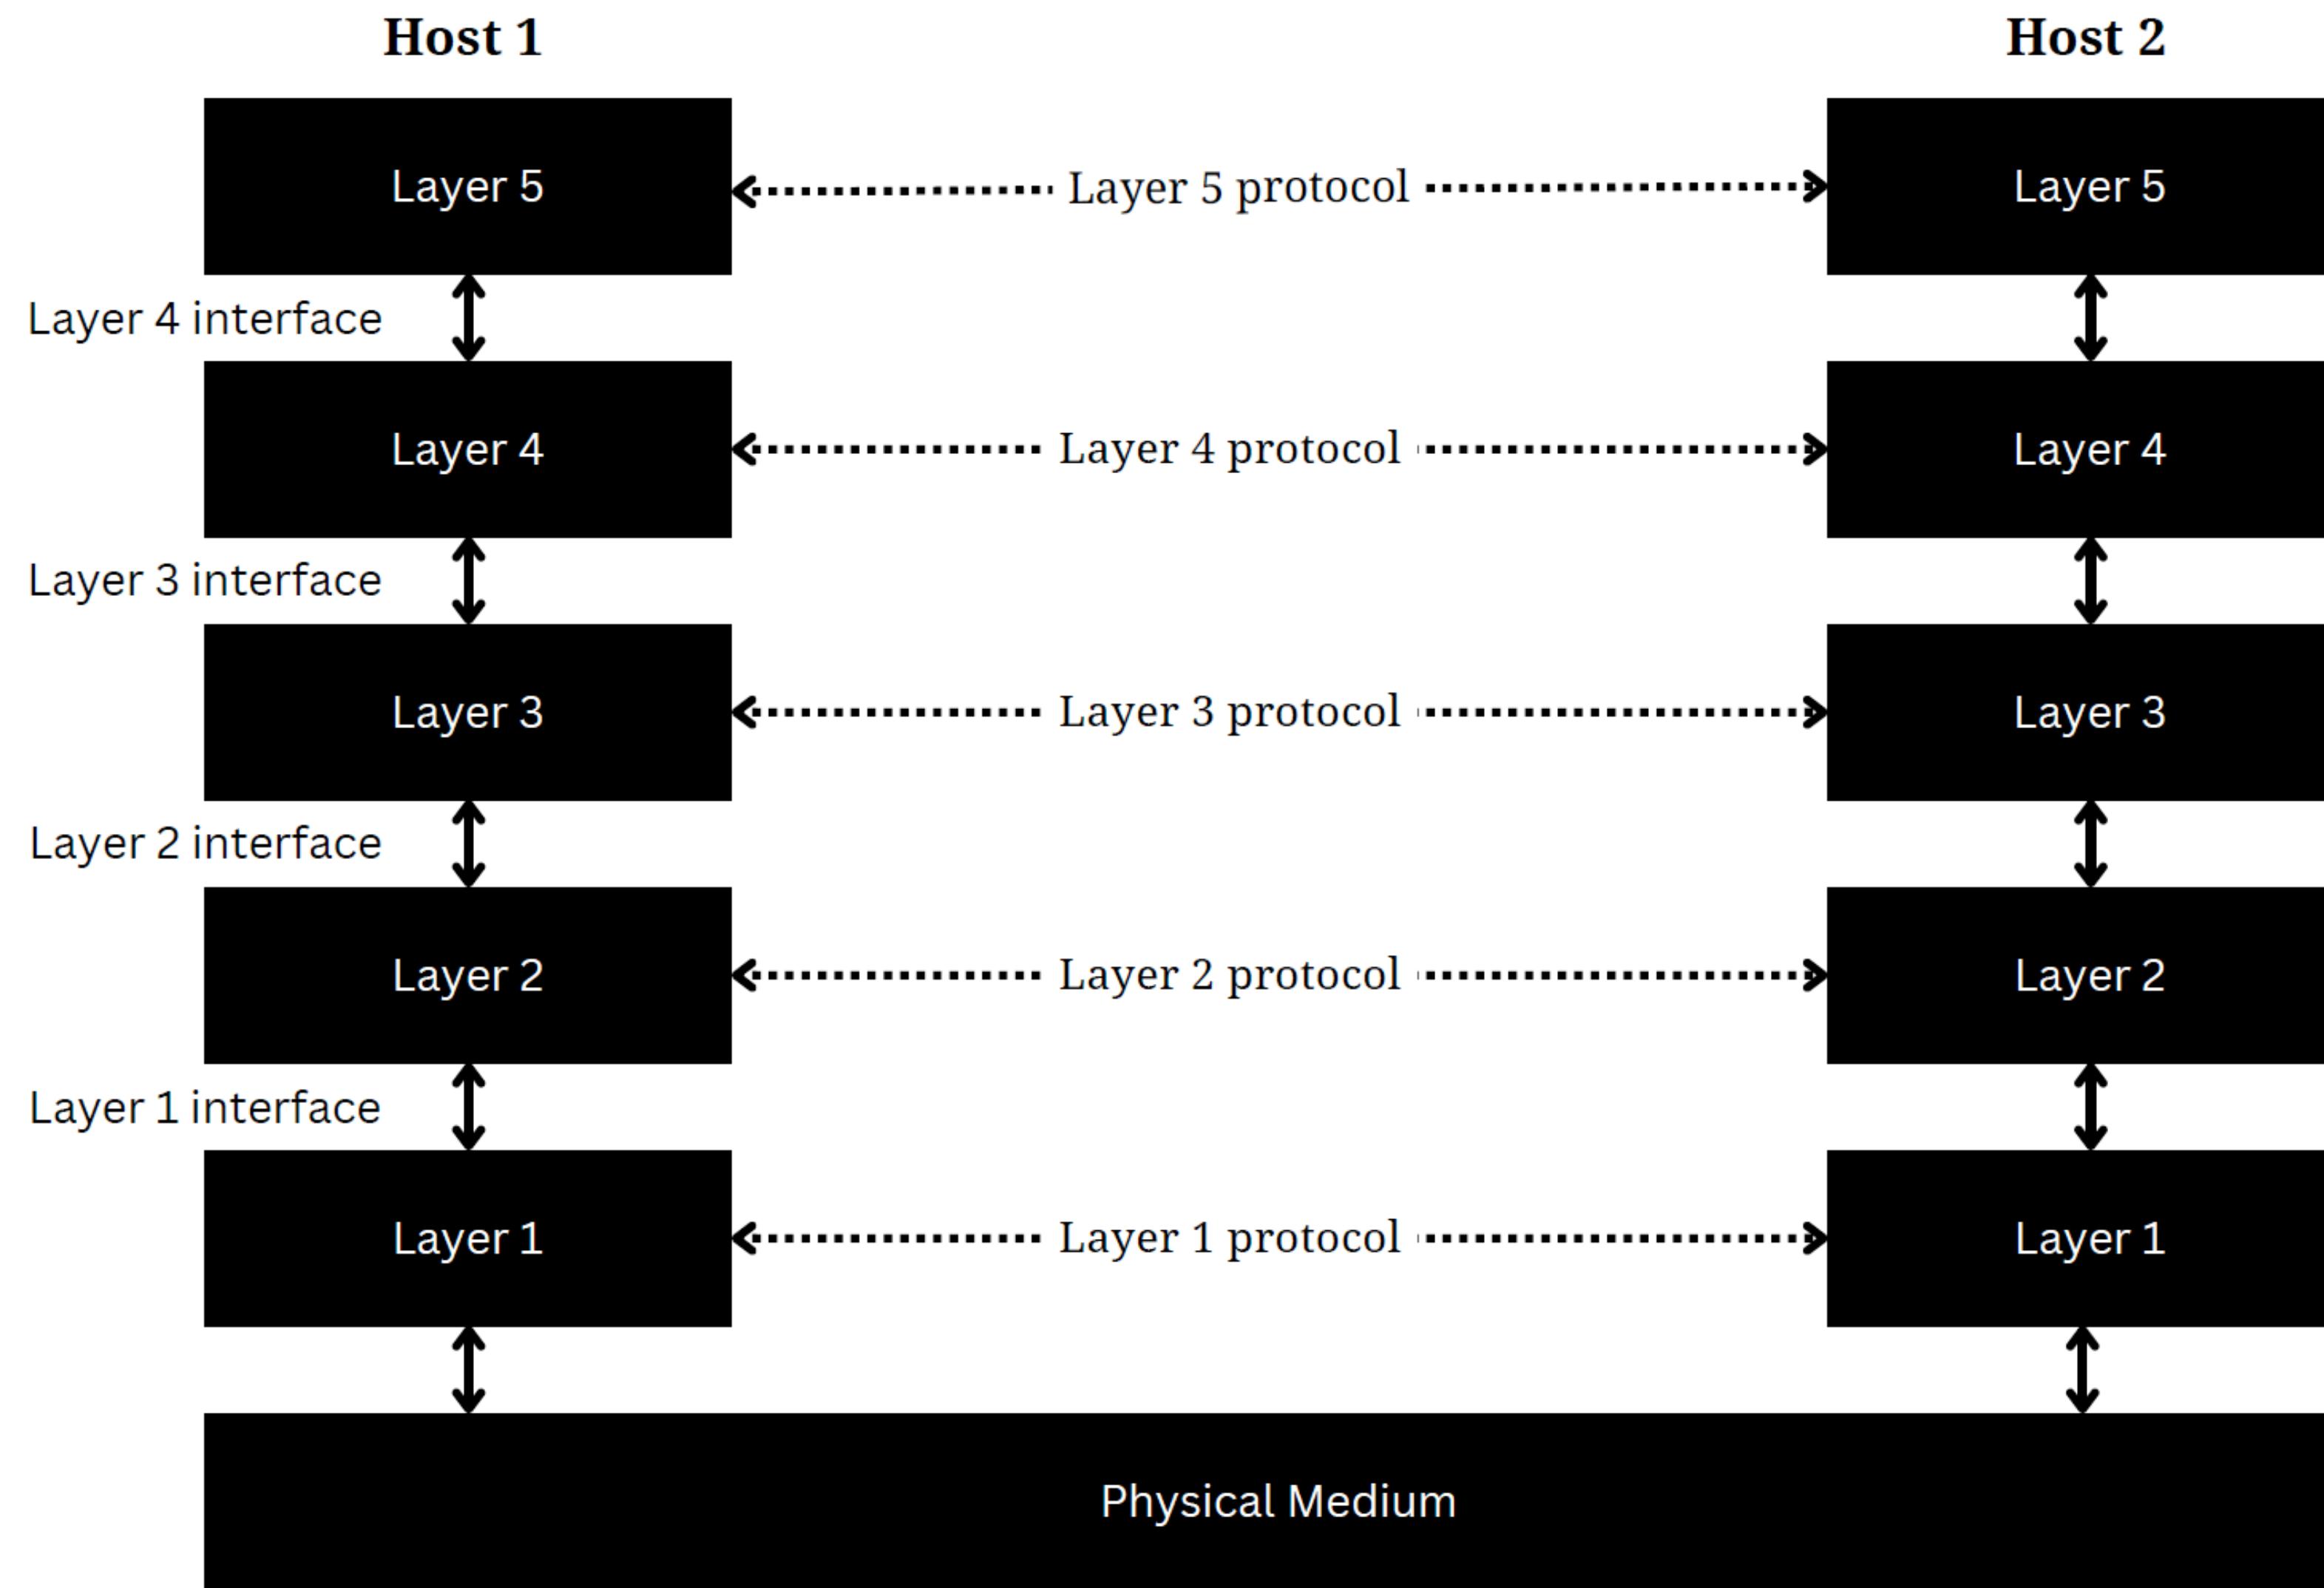
The role of Operating System



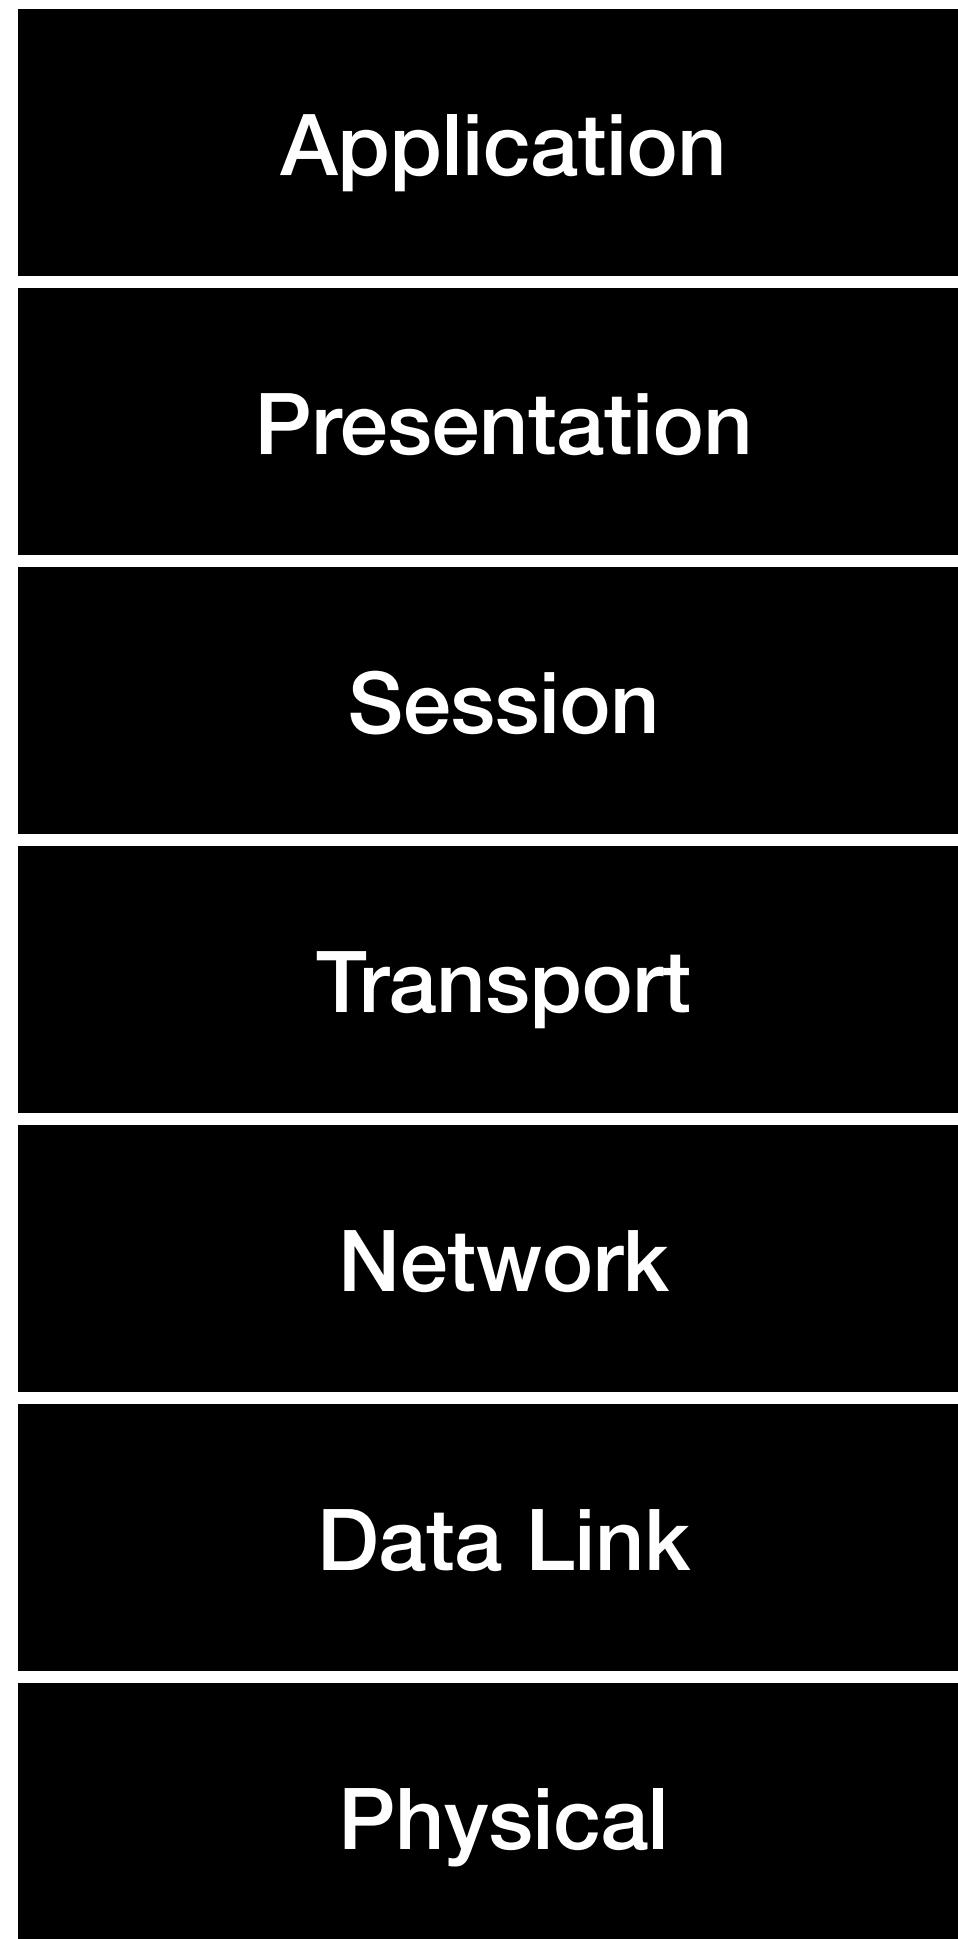
- Software component in the OS that supports network calls - **Protocol stack**
- Provides Service primitives which are nothing but system calls - **Some API?**



Networking Layers



The OSI Model



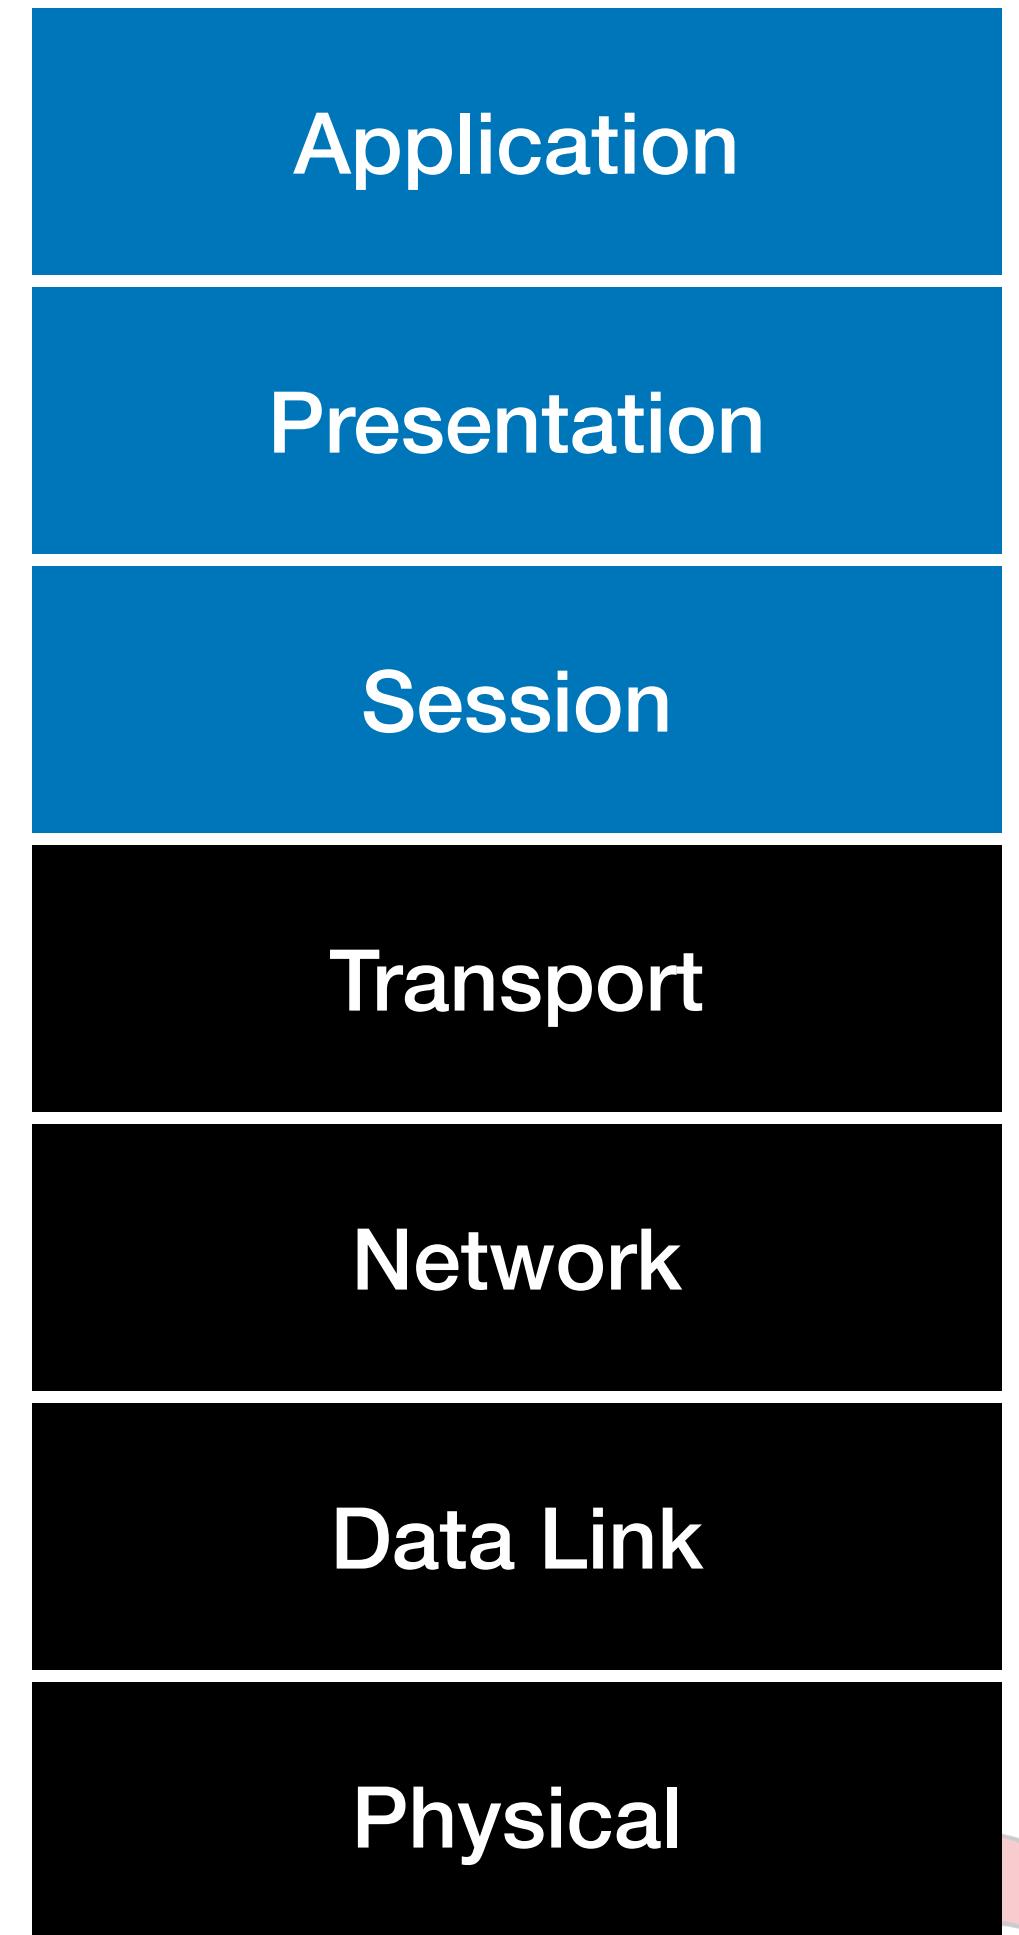
- Open System Interconnection (OSI)
- A Conceptual framework used to understand how communication works through different layers
- Divides the network communication process into seven layers
- Developed to facilitate interoperability between different technologies
- Each layer has a specific function. If they all do what they are supposed to do => sharing of data



Session, Presentation and Application

Application to Application

- **Session Layer (L5)**
 - Manages connection between different devices
 - Establishing, maintaining and terminating connections
- **Presentation Layer (L6)**
 - Ensures that data is in format that sender and receiver can understand
 - Manages data encryption, compression
- **Application Layer (L7)**
 - Provides network services to the application processes
 - Eg: web browser, email clients, other softwares/apps



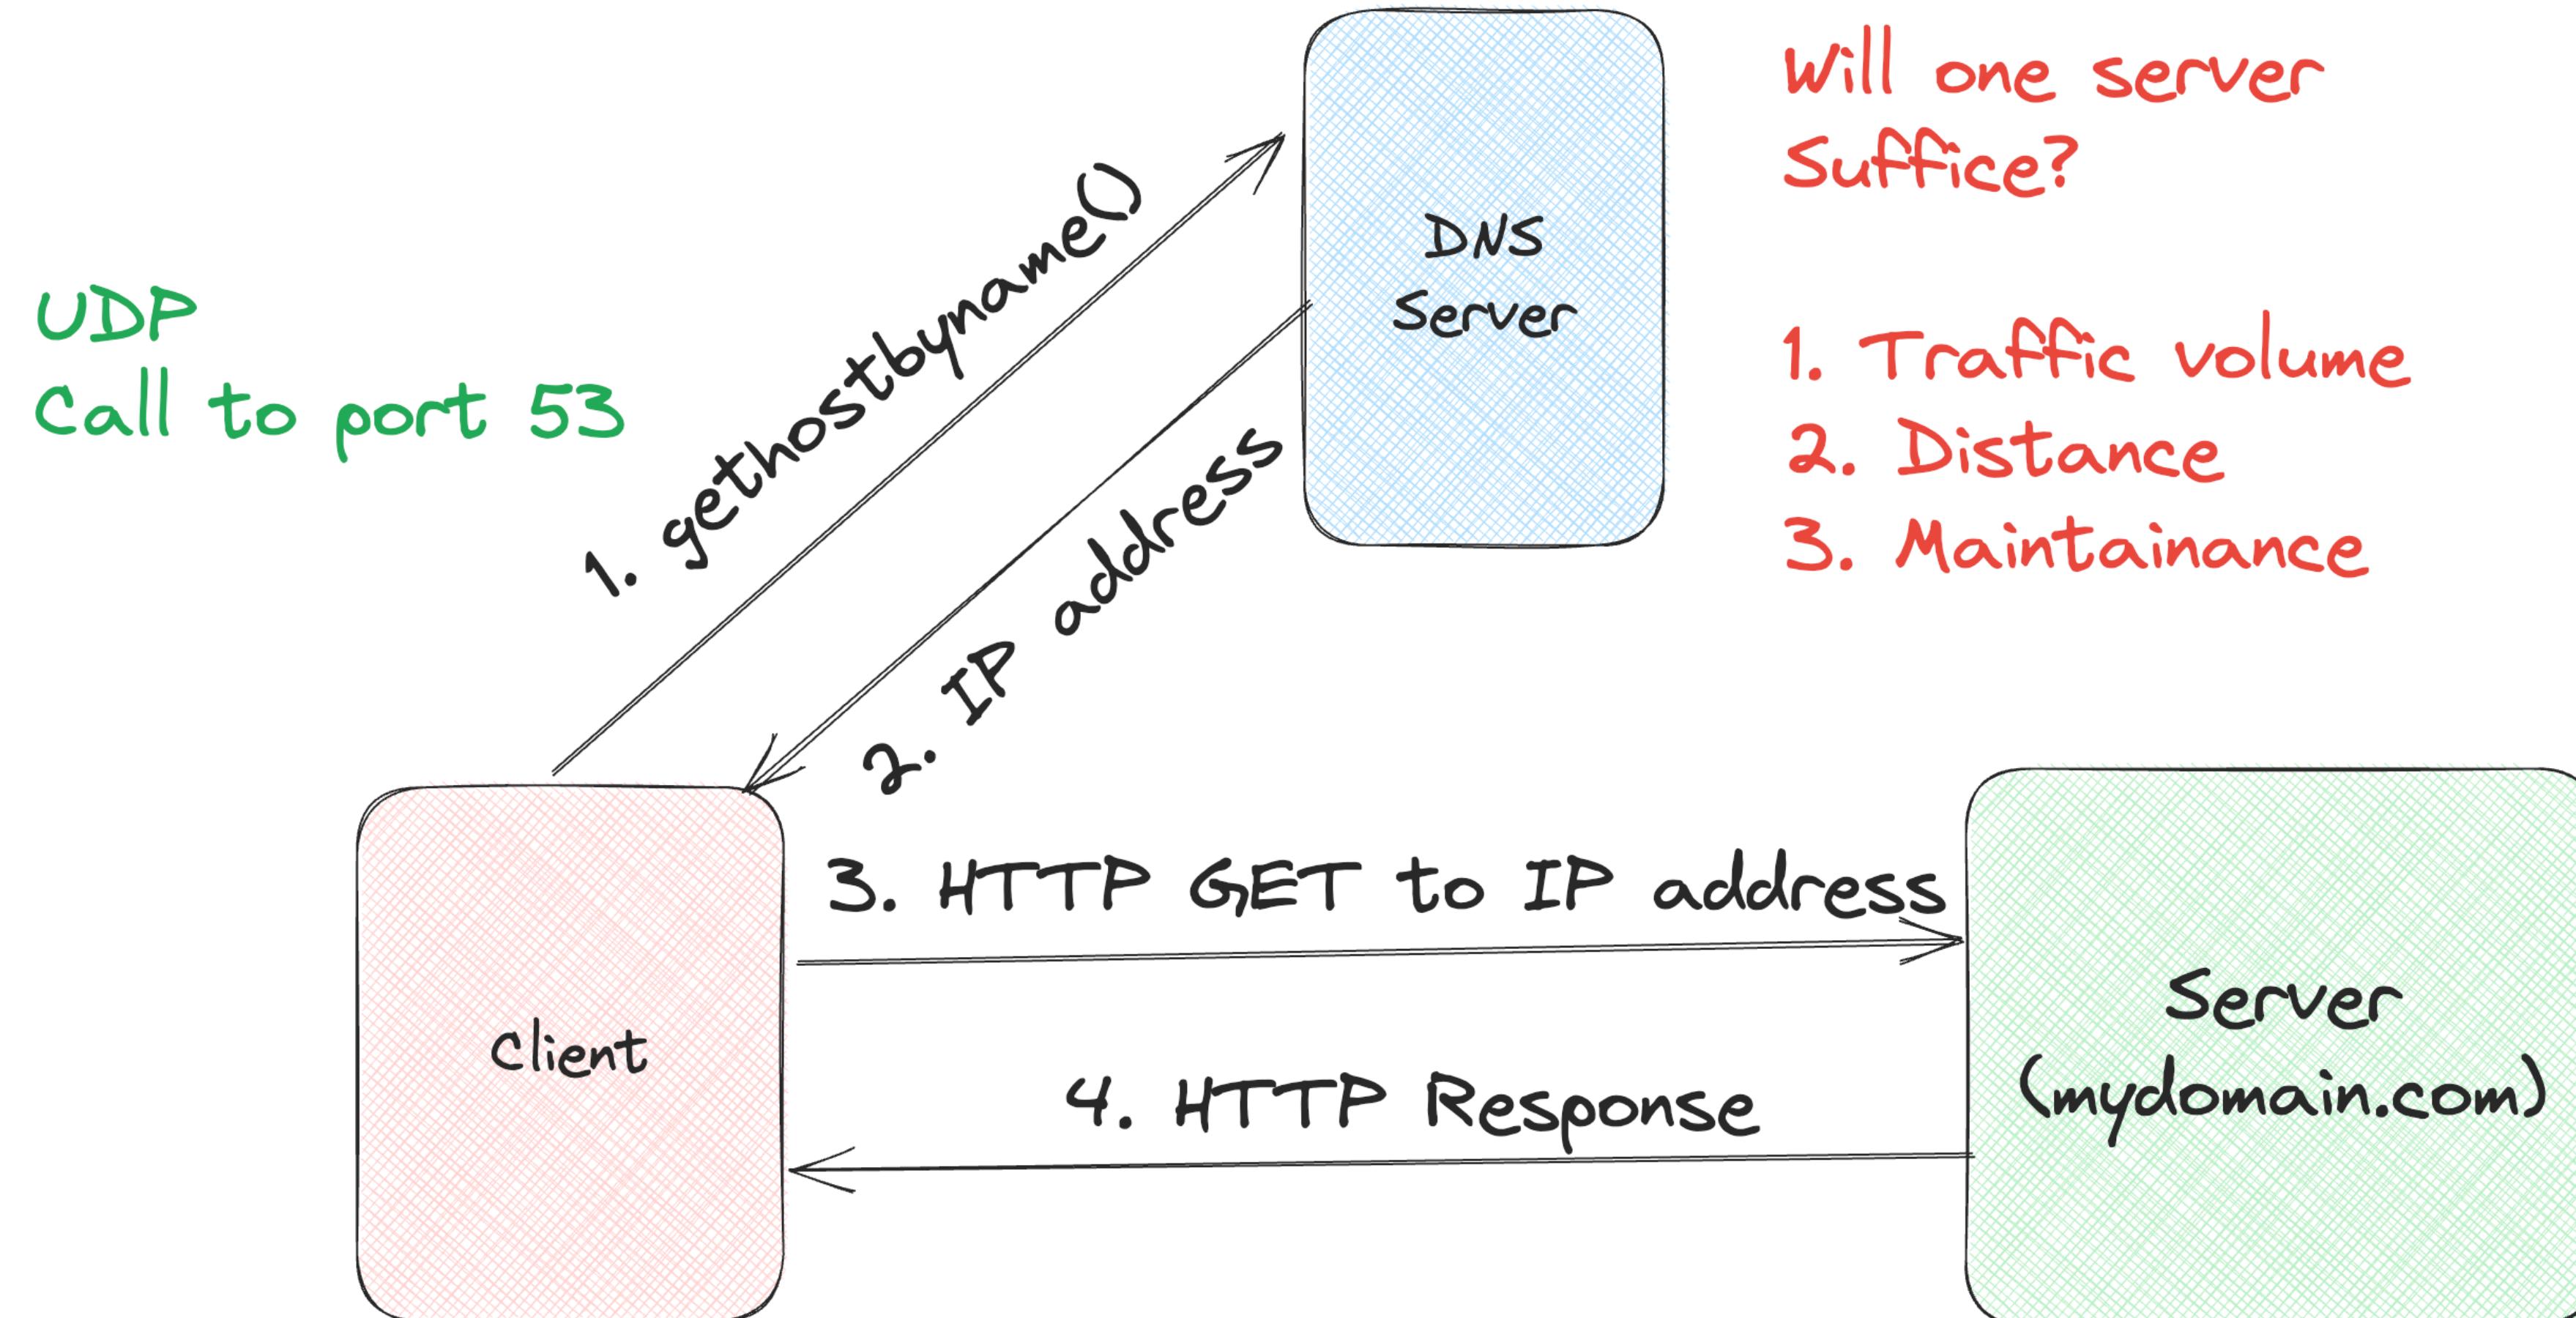
Application Layer Protocol

What does it mean?

- Application layer protocol defines the following:
 - Types of message exchanges (request/response)
 - Syntax of various message types
 - Semantics of the fields
 - When and how the process sends and responds to messages
- Some protocols: HTTP, SMTP, DNS, etc.



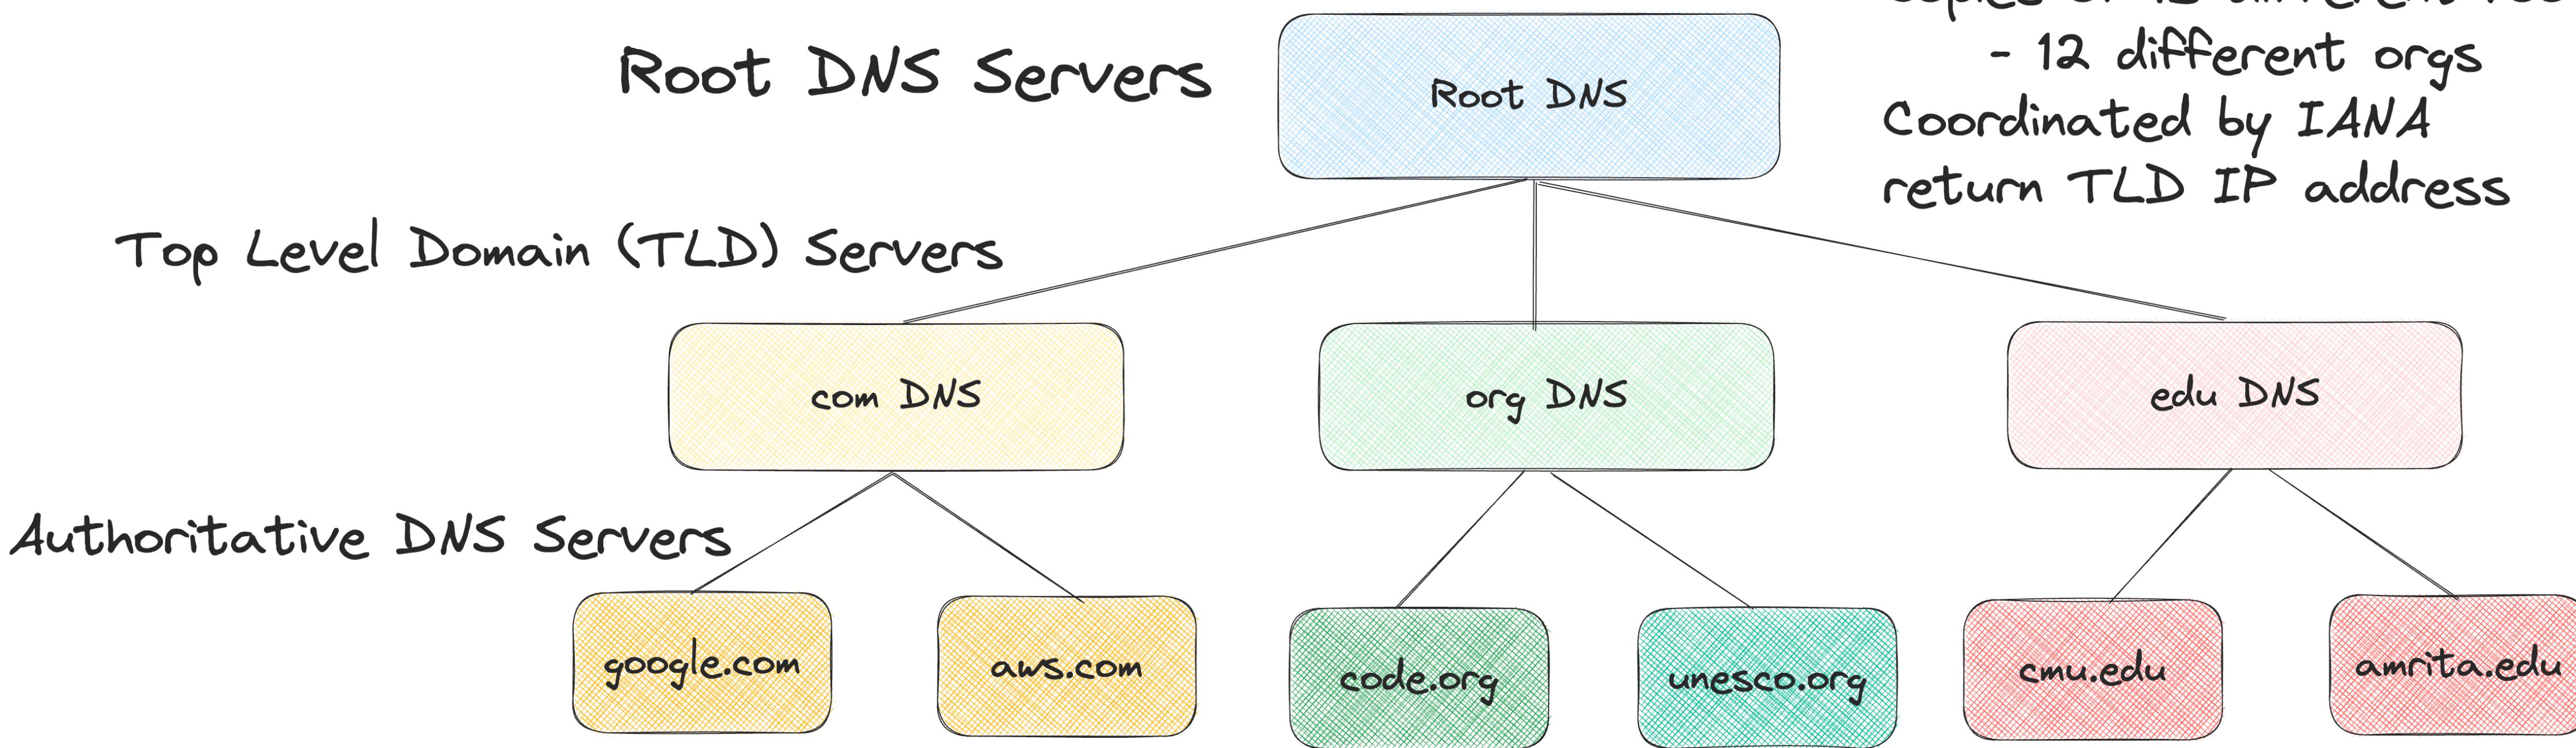
Mapping domain name to IP: DNS



wants to
send HTTP request
to mydomain.com



DNS: Distributed Hierarchical Database



More than 1000 root servers
Copies of 13 different root servers

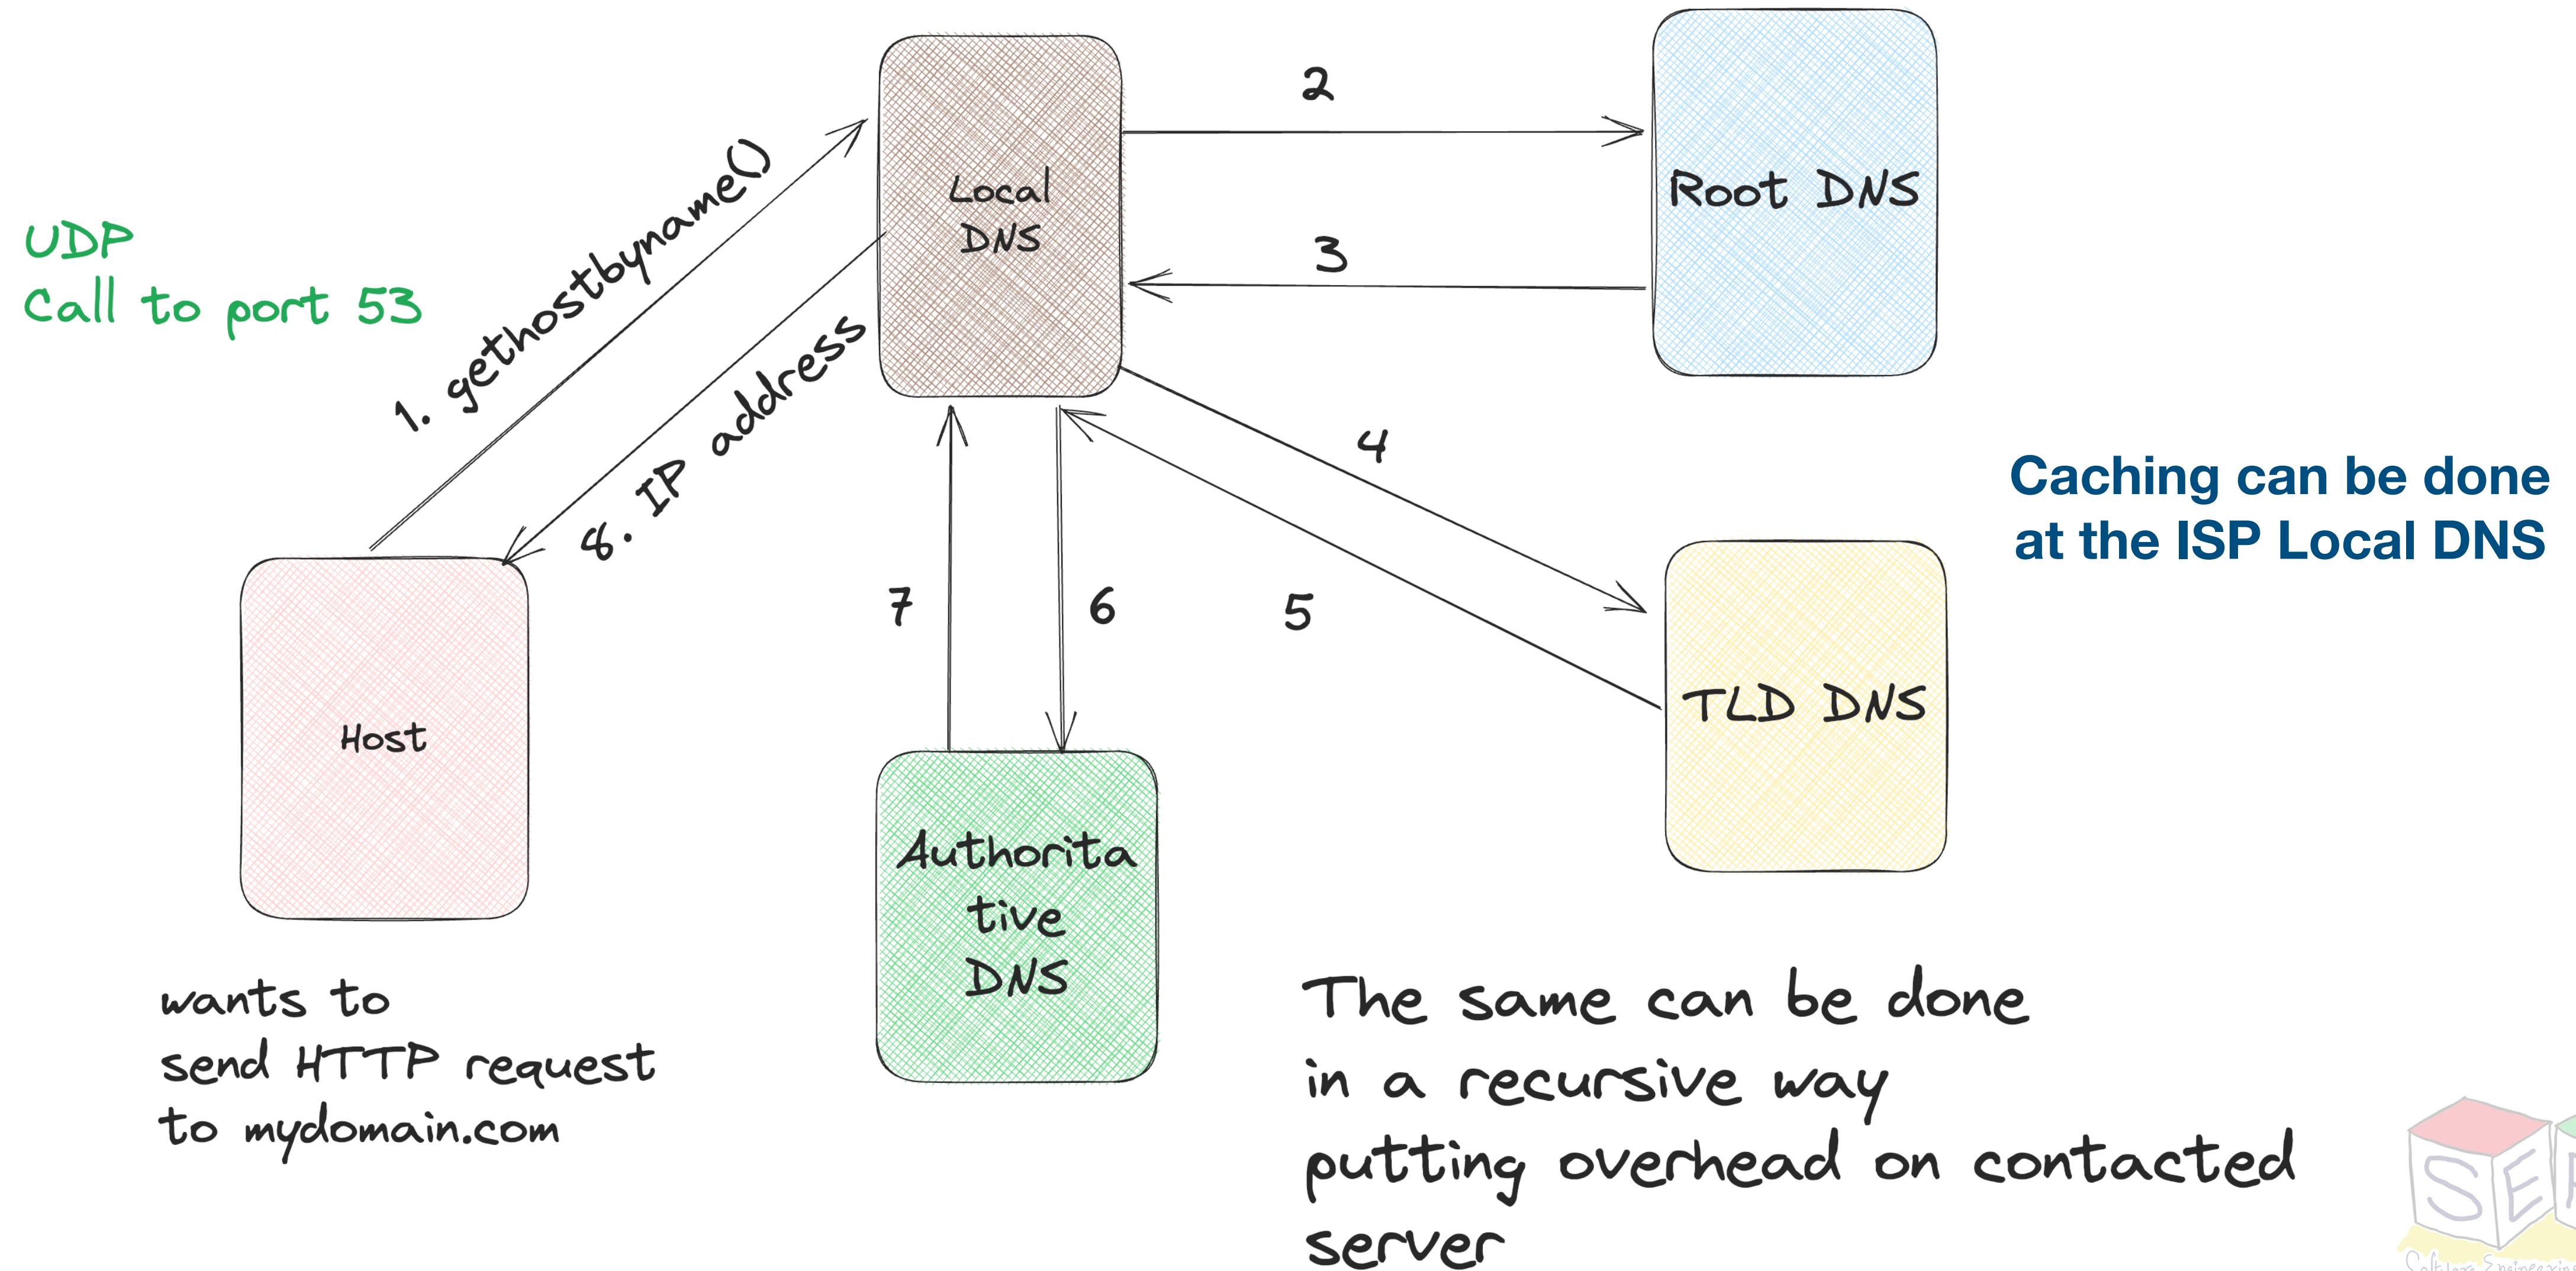
- 12 different orgs

Coordinated by IANA
return TLD IP address

TLD Servers - can be maintained by orgs, provide IP of authoritative DNS Servers
Authoritative DNS servers - Orgs can choose to implement their own or go for third party
All DNS records have to be made public - that maps hosts to IP address

Local DNS

Each ISP can have DNS and clients can connect to that



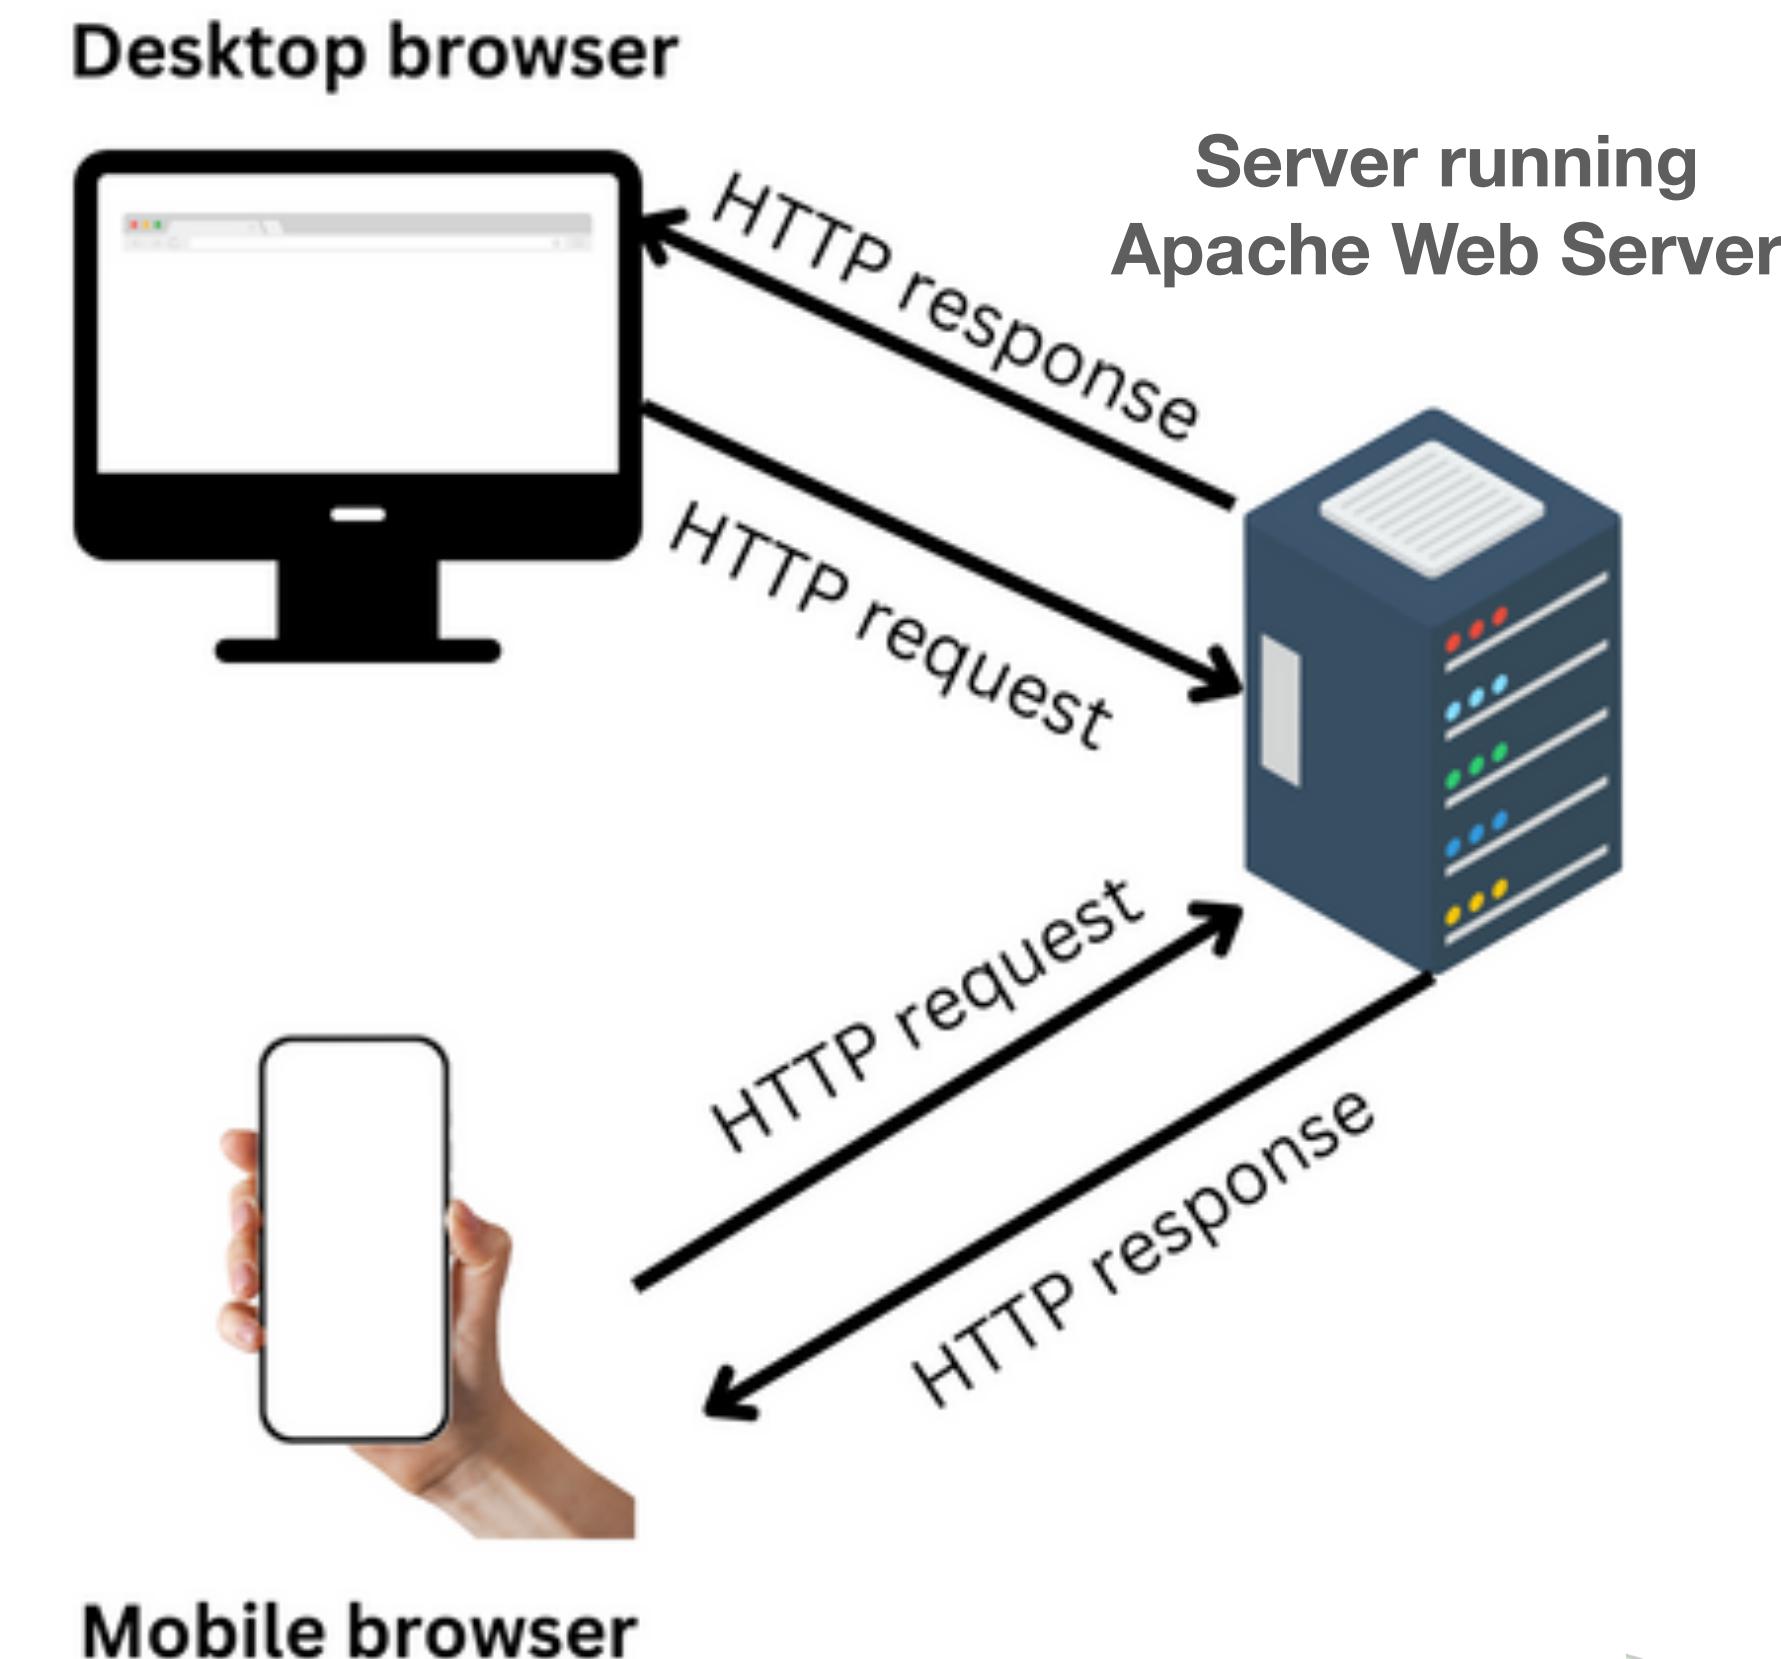
HTTP: Hyper Text Transfer Protocol

- Application layer protocol of the web
- Implemented in two programs: Client and Server
- HTTP protocol defines structure of messages
- **Client:** browser that sends requests, receives and displays web objects (using HTTP protocol)
- **Server:** Web server that sends objects in response to requests (using HTTP protocol)
- Uses **TCP** and it is stateless

<https://iiit.ac.in/samplePage.html>

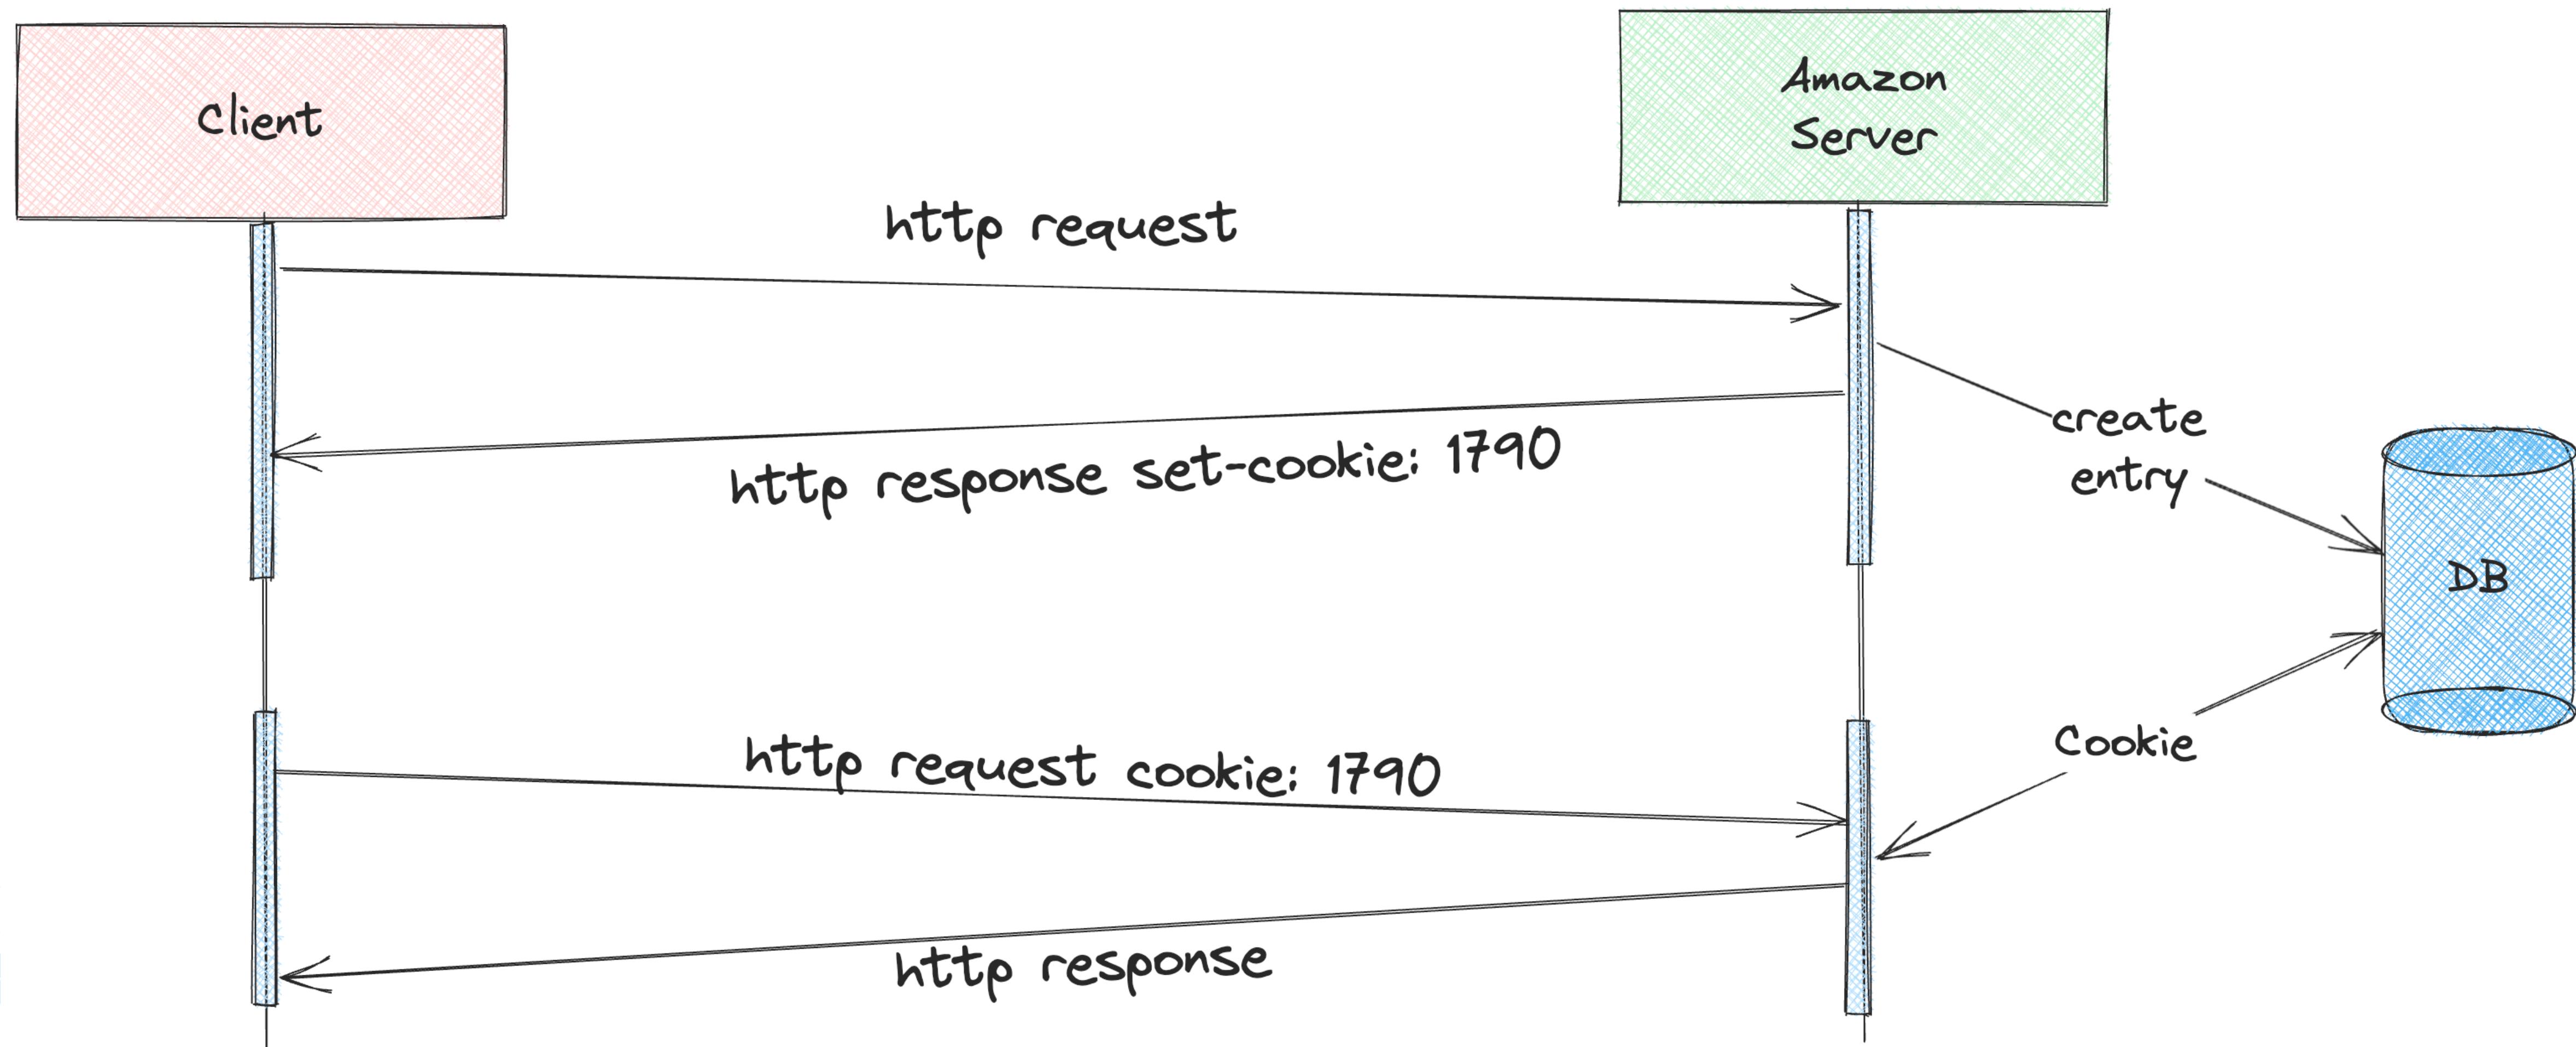
Host name: iiii.ac.in

Object: samplePage.html



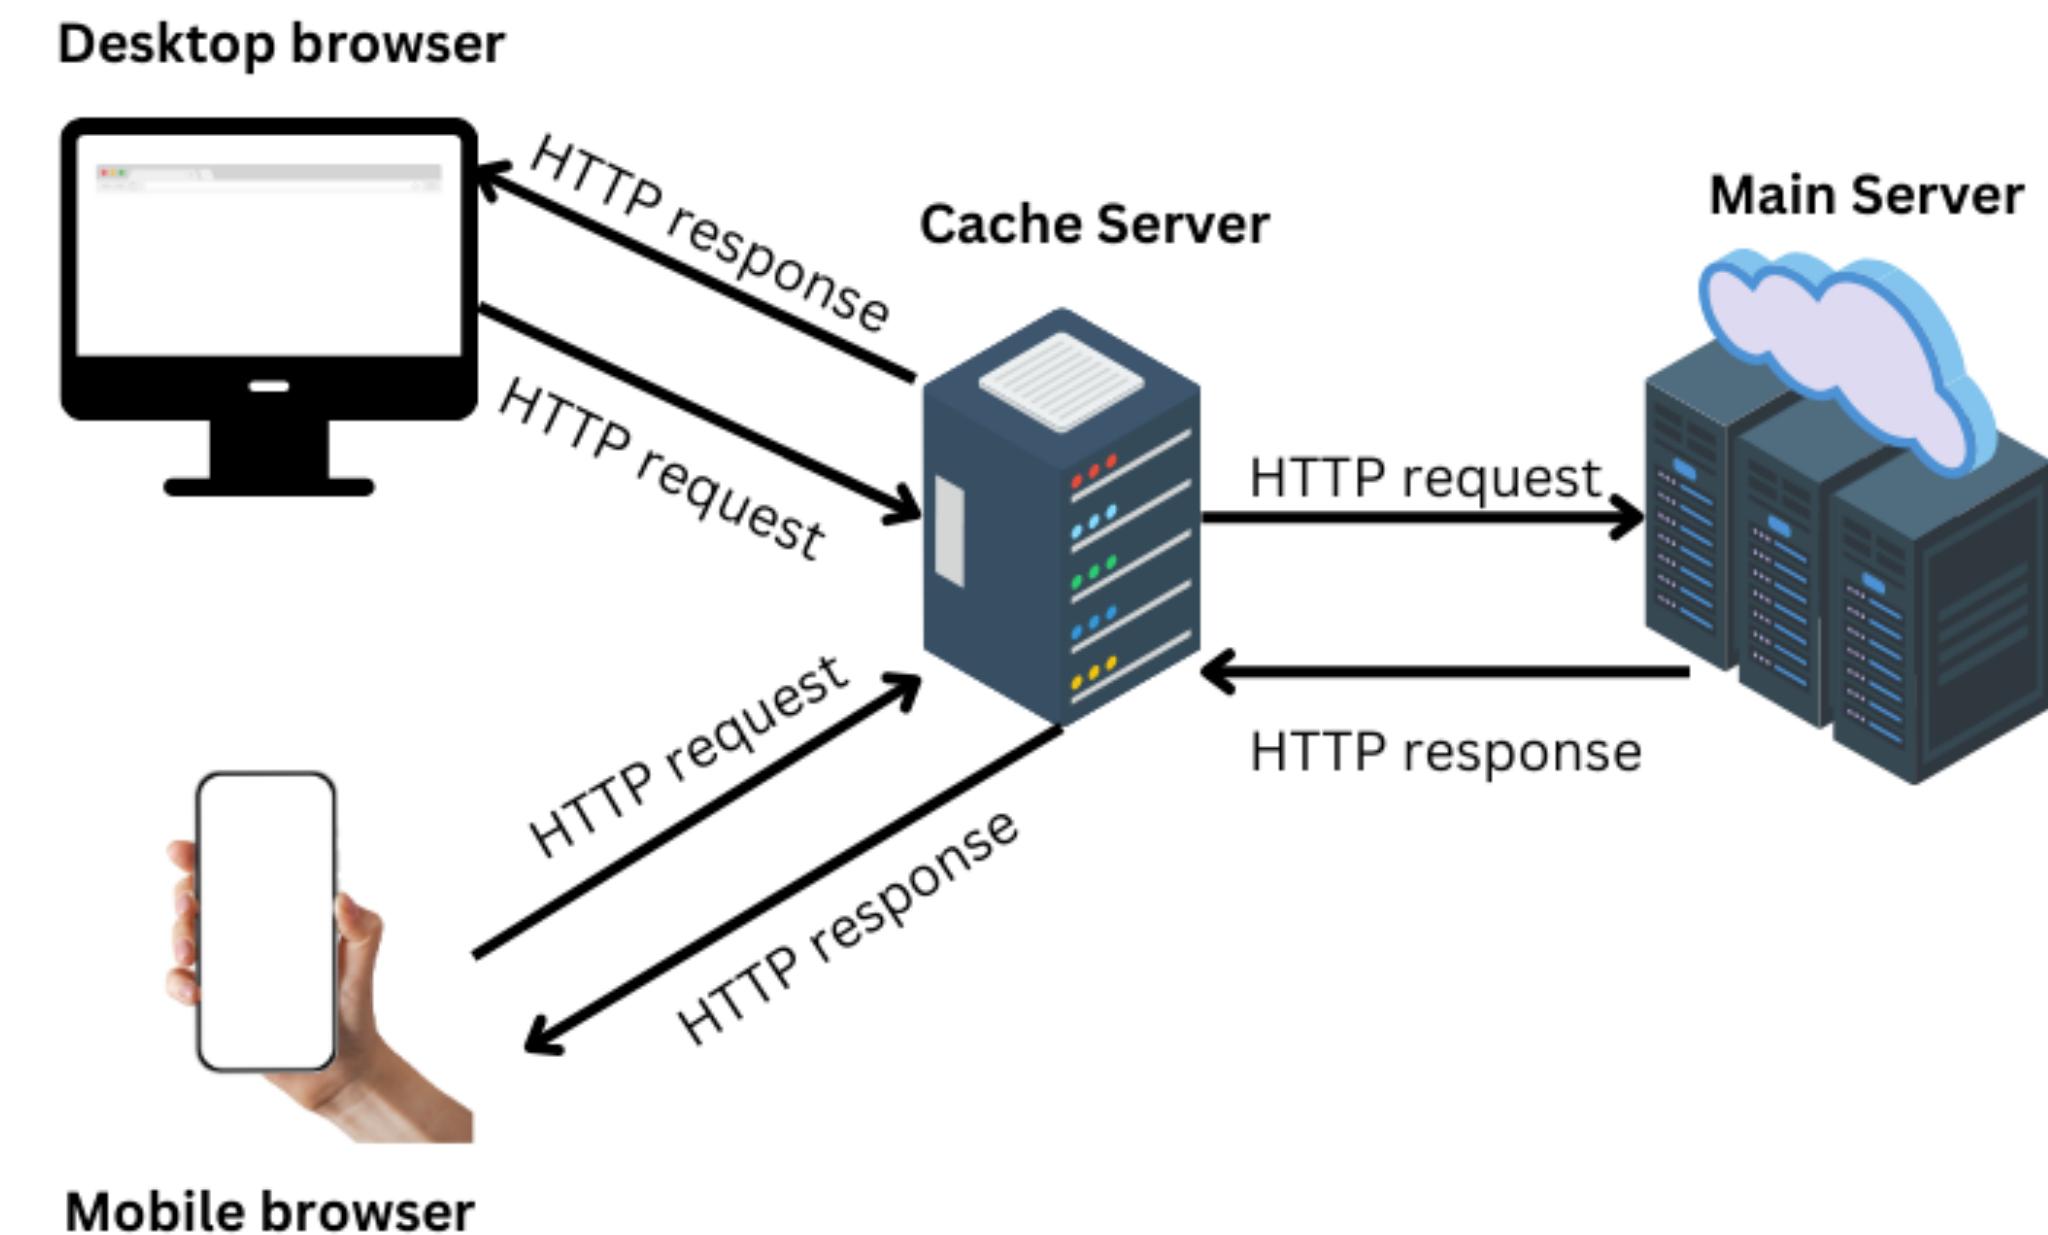
How to store information in stateless Protocol?

Using Cookies



Web Caches

- Not every time we need to access the main (original) web server
- We can have proxy server that satisfies request on behalf of main server
- Browser can be controlled to point towards a cache (mentioned in response header)
 - If cache hit: return object from cache
 - Else cache request object from main server and returns it
- Conditional GET is used to update Cache (“if-modified-since”)

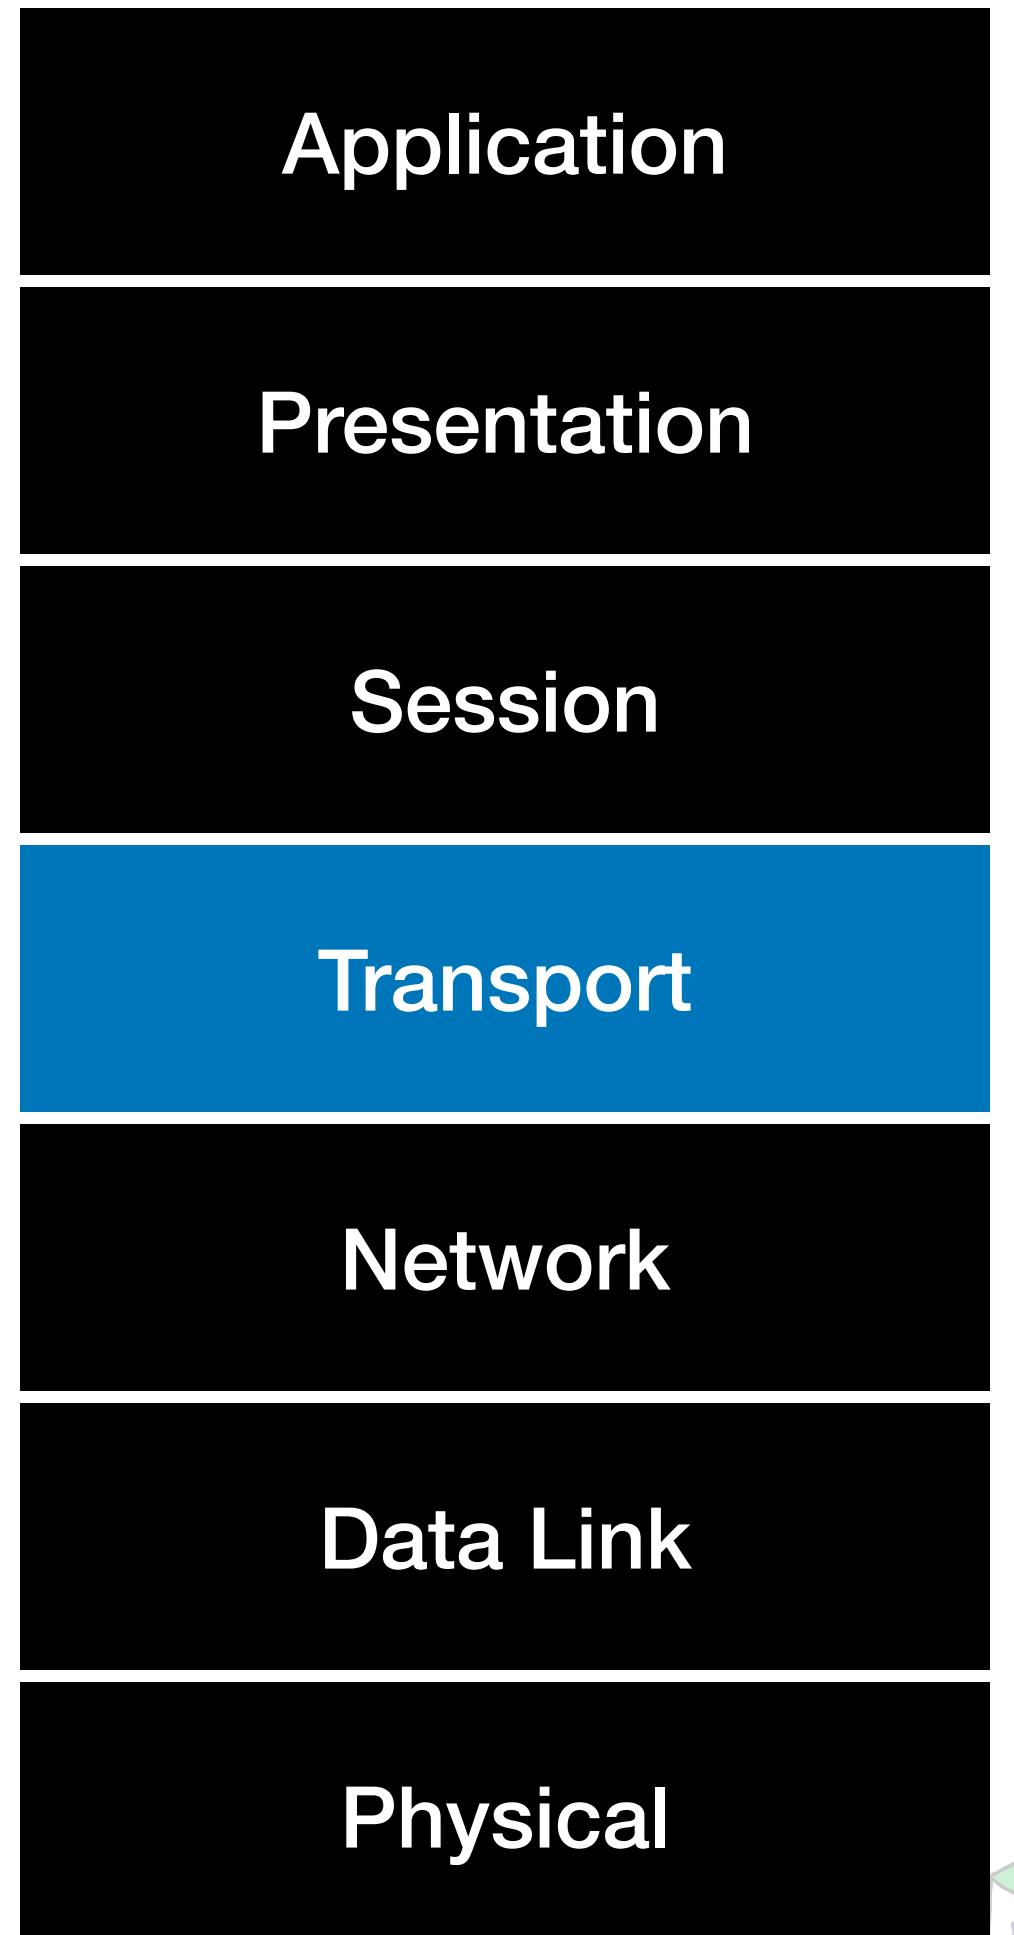


Cache-Control: max-age=<seconds>

Cache-Control: no-cache

On to Transport Layer (L4)

- Process gets the data delivered through support of transport layer
- Addressing scheme: Ports
- Layer 4 has an addressing scheme to guarantee message delivery
 - Ports! (0 - 65535), Privileged: 0-1023, Registered: 1024 - 49151
- Two strategies/protocols that allows this
 - **Transmission Control Protocol (TCP)** - favours reliability
 - **User Datagram Protocol (UDP)** - favours efficiency



TCP vs UDP

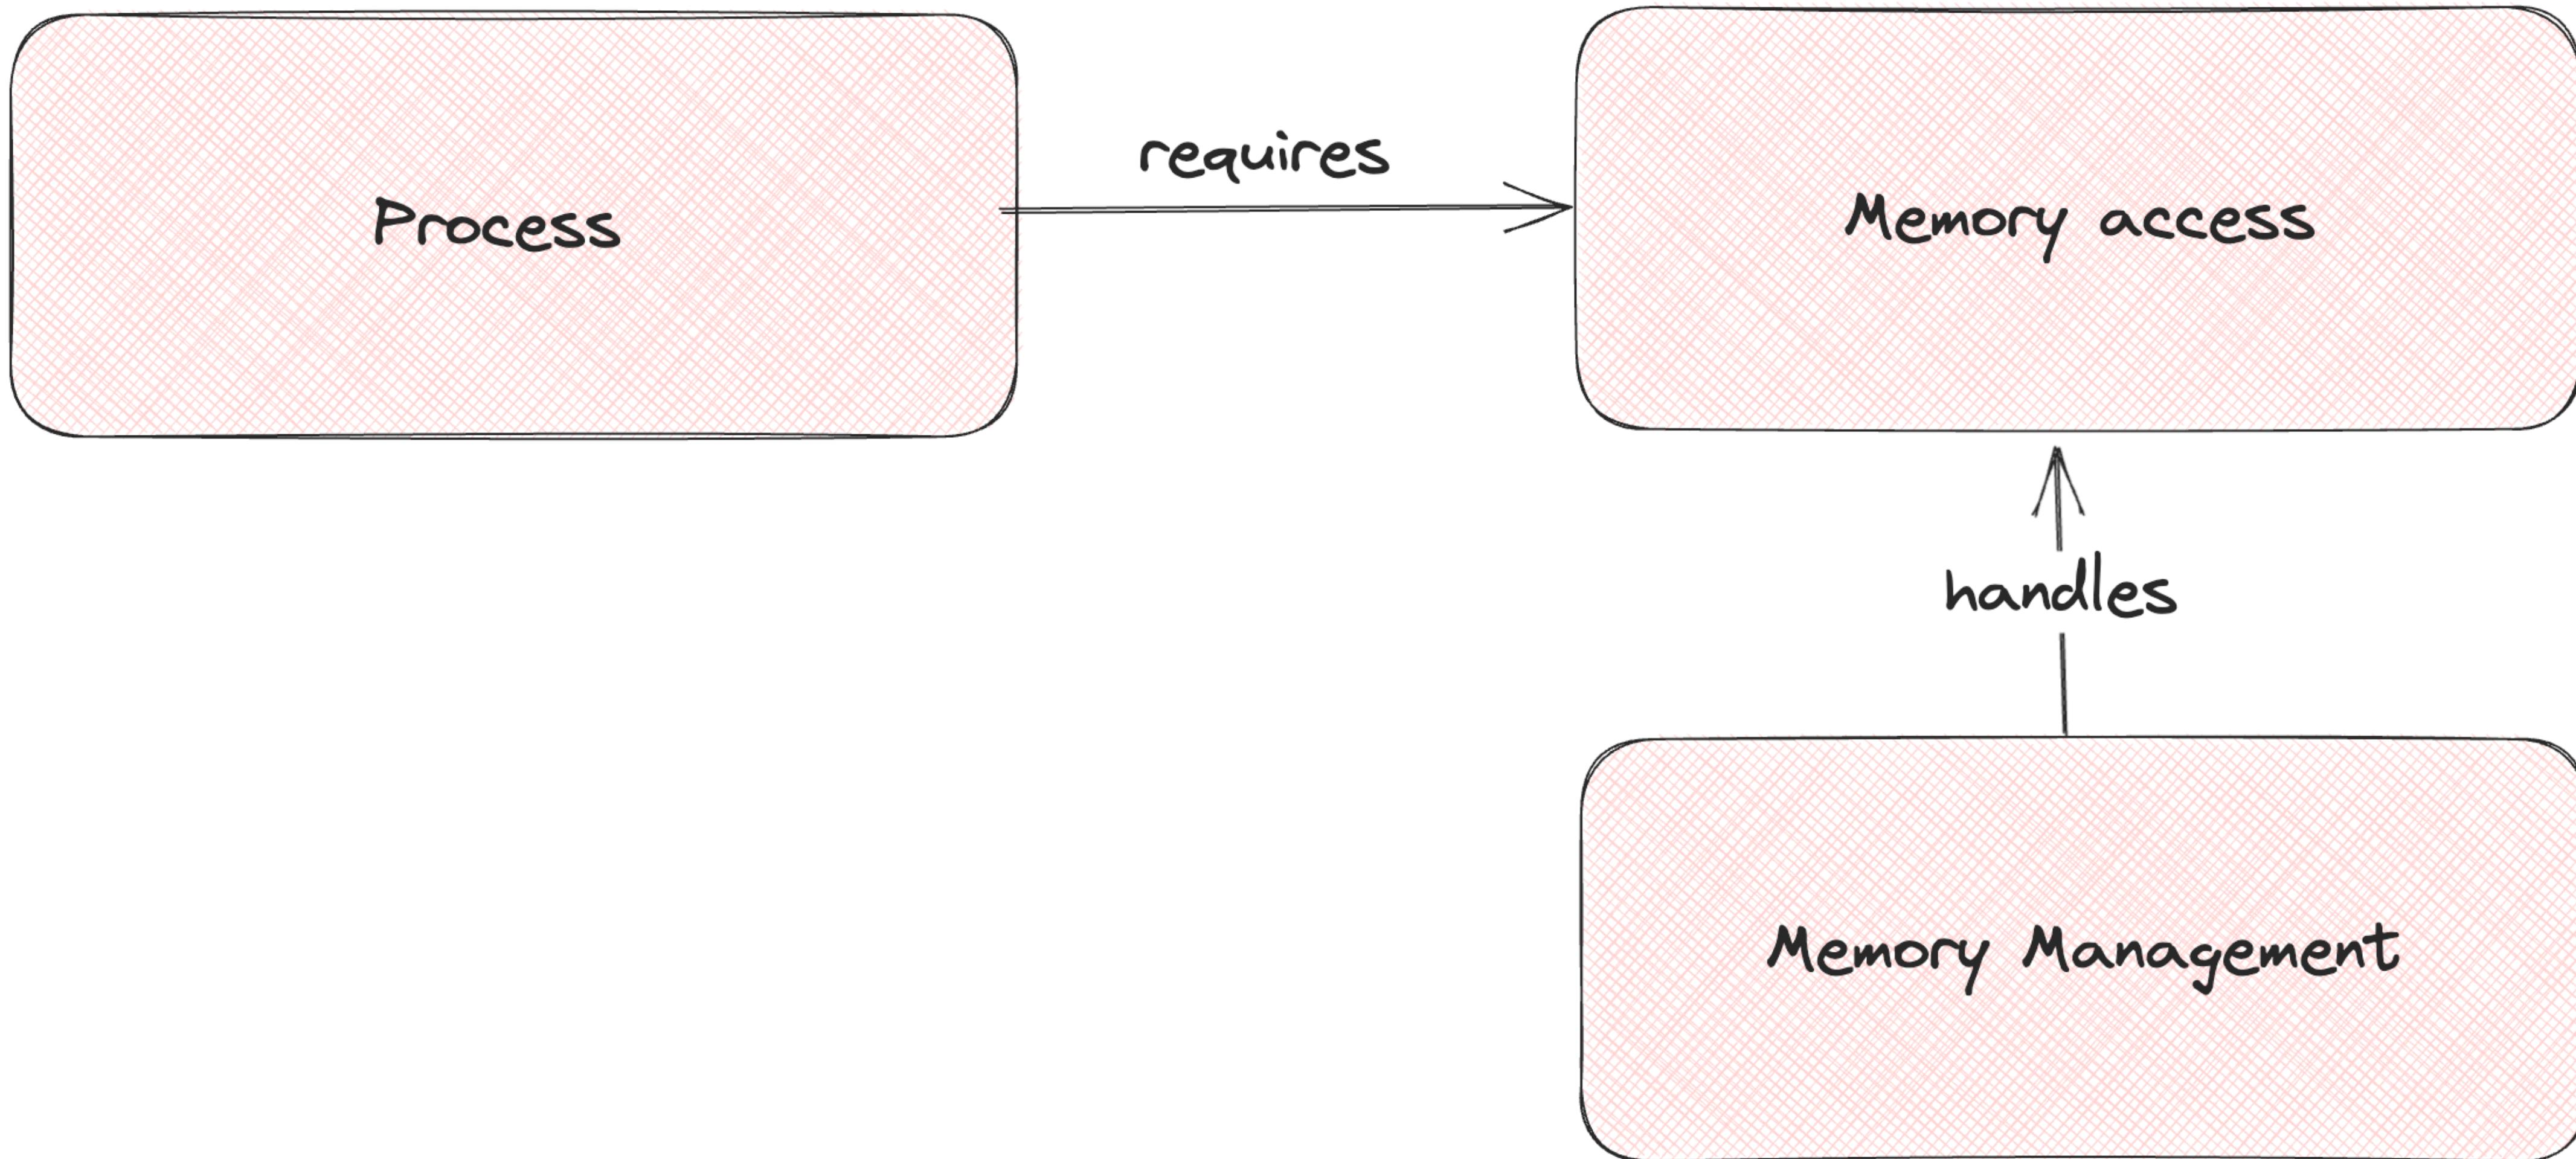
TCP	UDP
Connection Oriented	Not Connection Oriented
Reliability (order is maintained and retransmission)	Unreliable (At L4)
Higher overhead - reliability, error checking, etc	Low overhead
Flow control (based on network)	No implicit flow control
Error detection - retransmit erroneous packets	Has some error checking - Erroneous packets are discarded without notification
Congestion Control	No Congestion Control
Use cases: HTTP/HTTPS, File transfer, Mail	Use cases: Streaming data, VoIP, DNS queries, ..



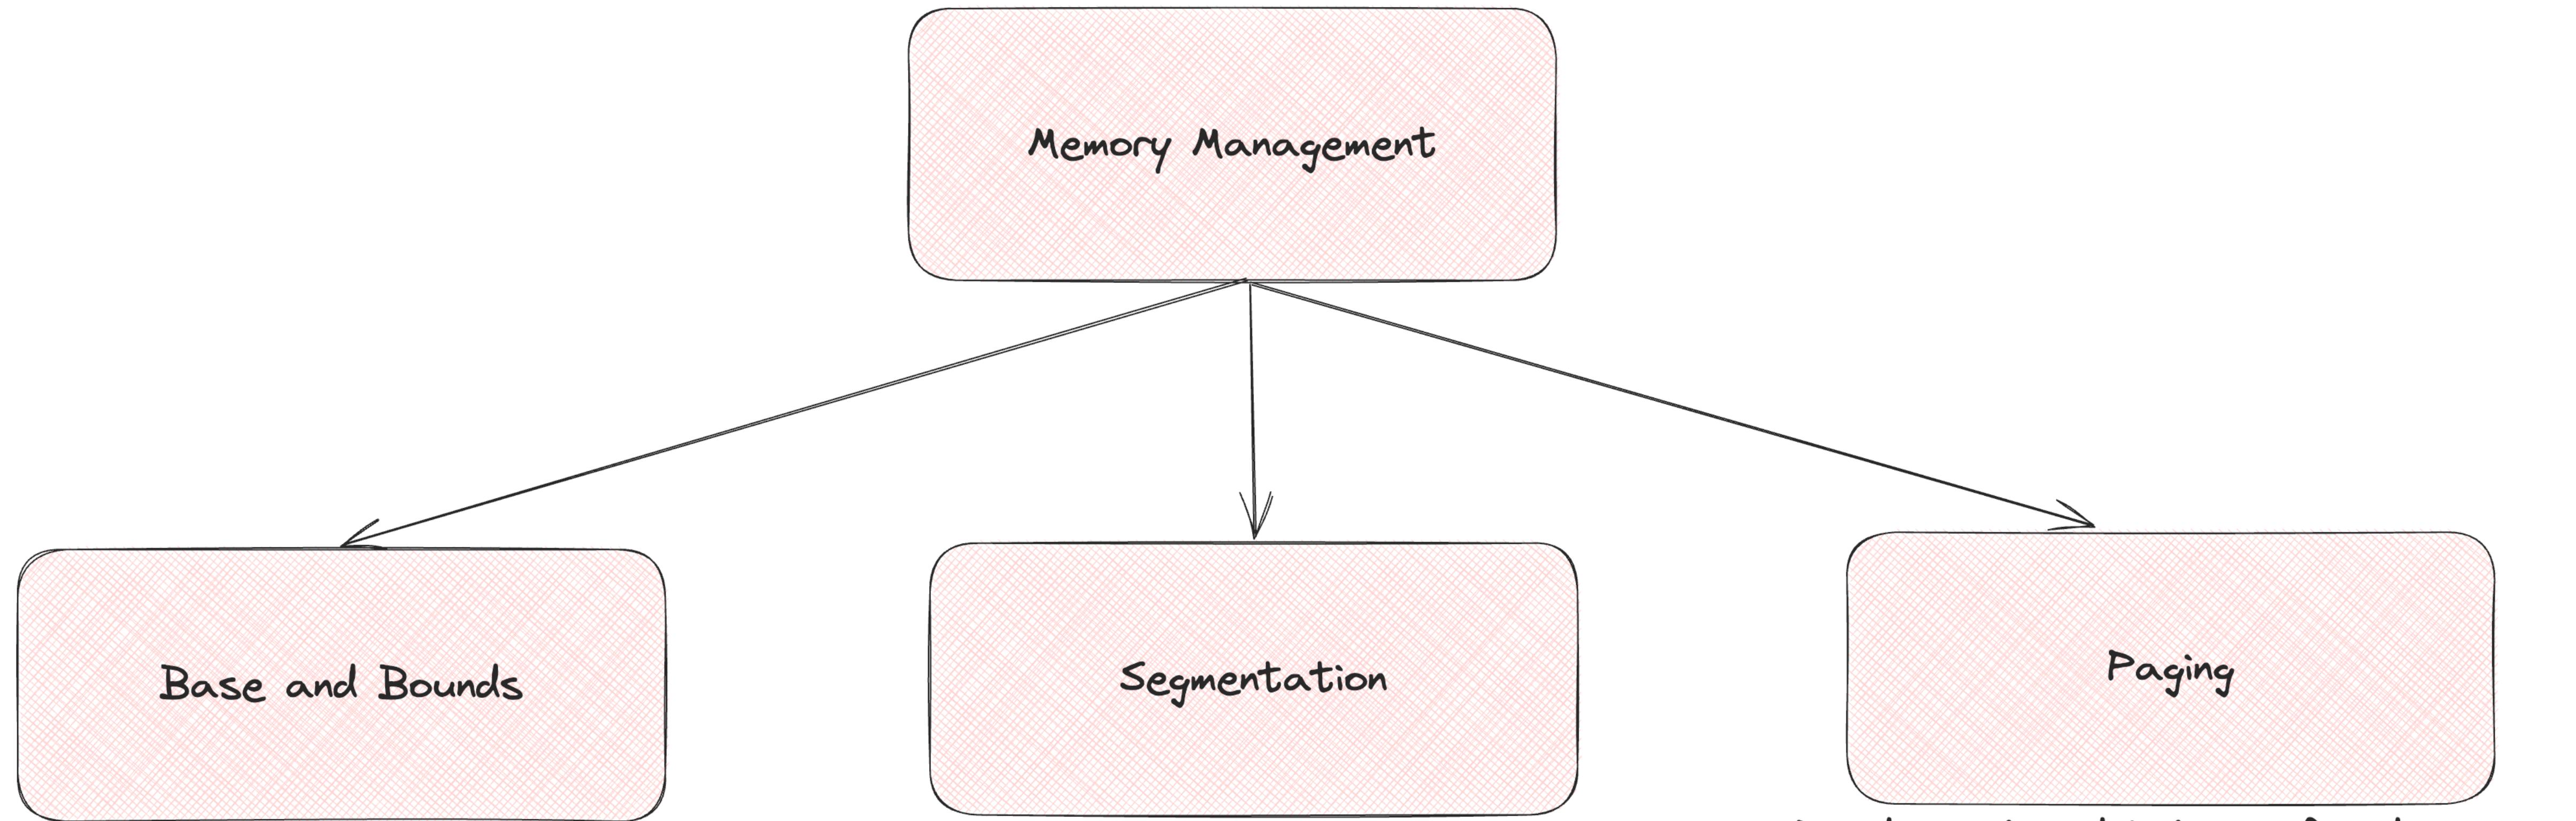
TCP and how it works!



Process requires Memory: Memory Virtualization



Memory Management



Base register
and bounds register

(just add and check
if less than bounds for
translation)

Generalized Base and bounds

(for each segment there is a
base and bounds that needs to
be assigned)

Paging

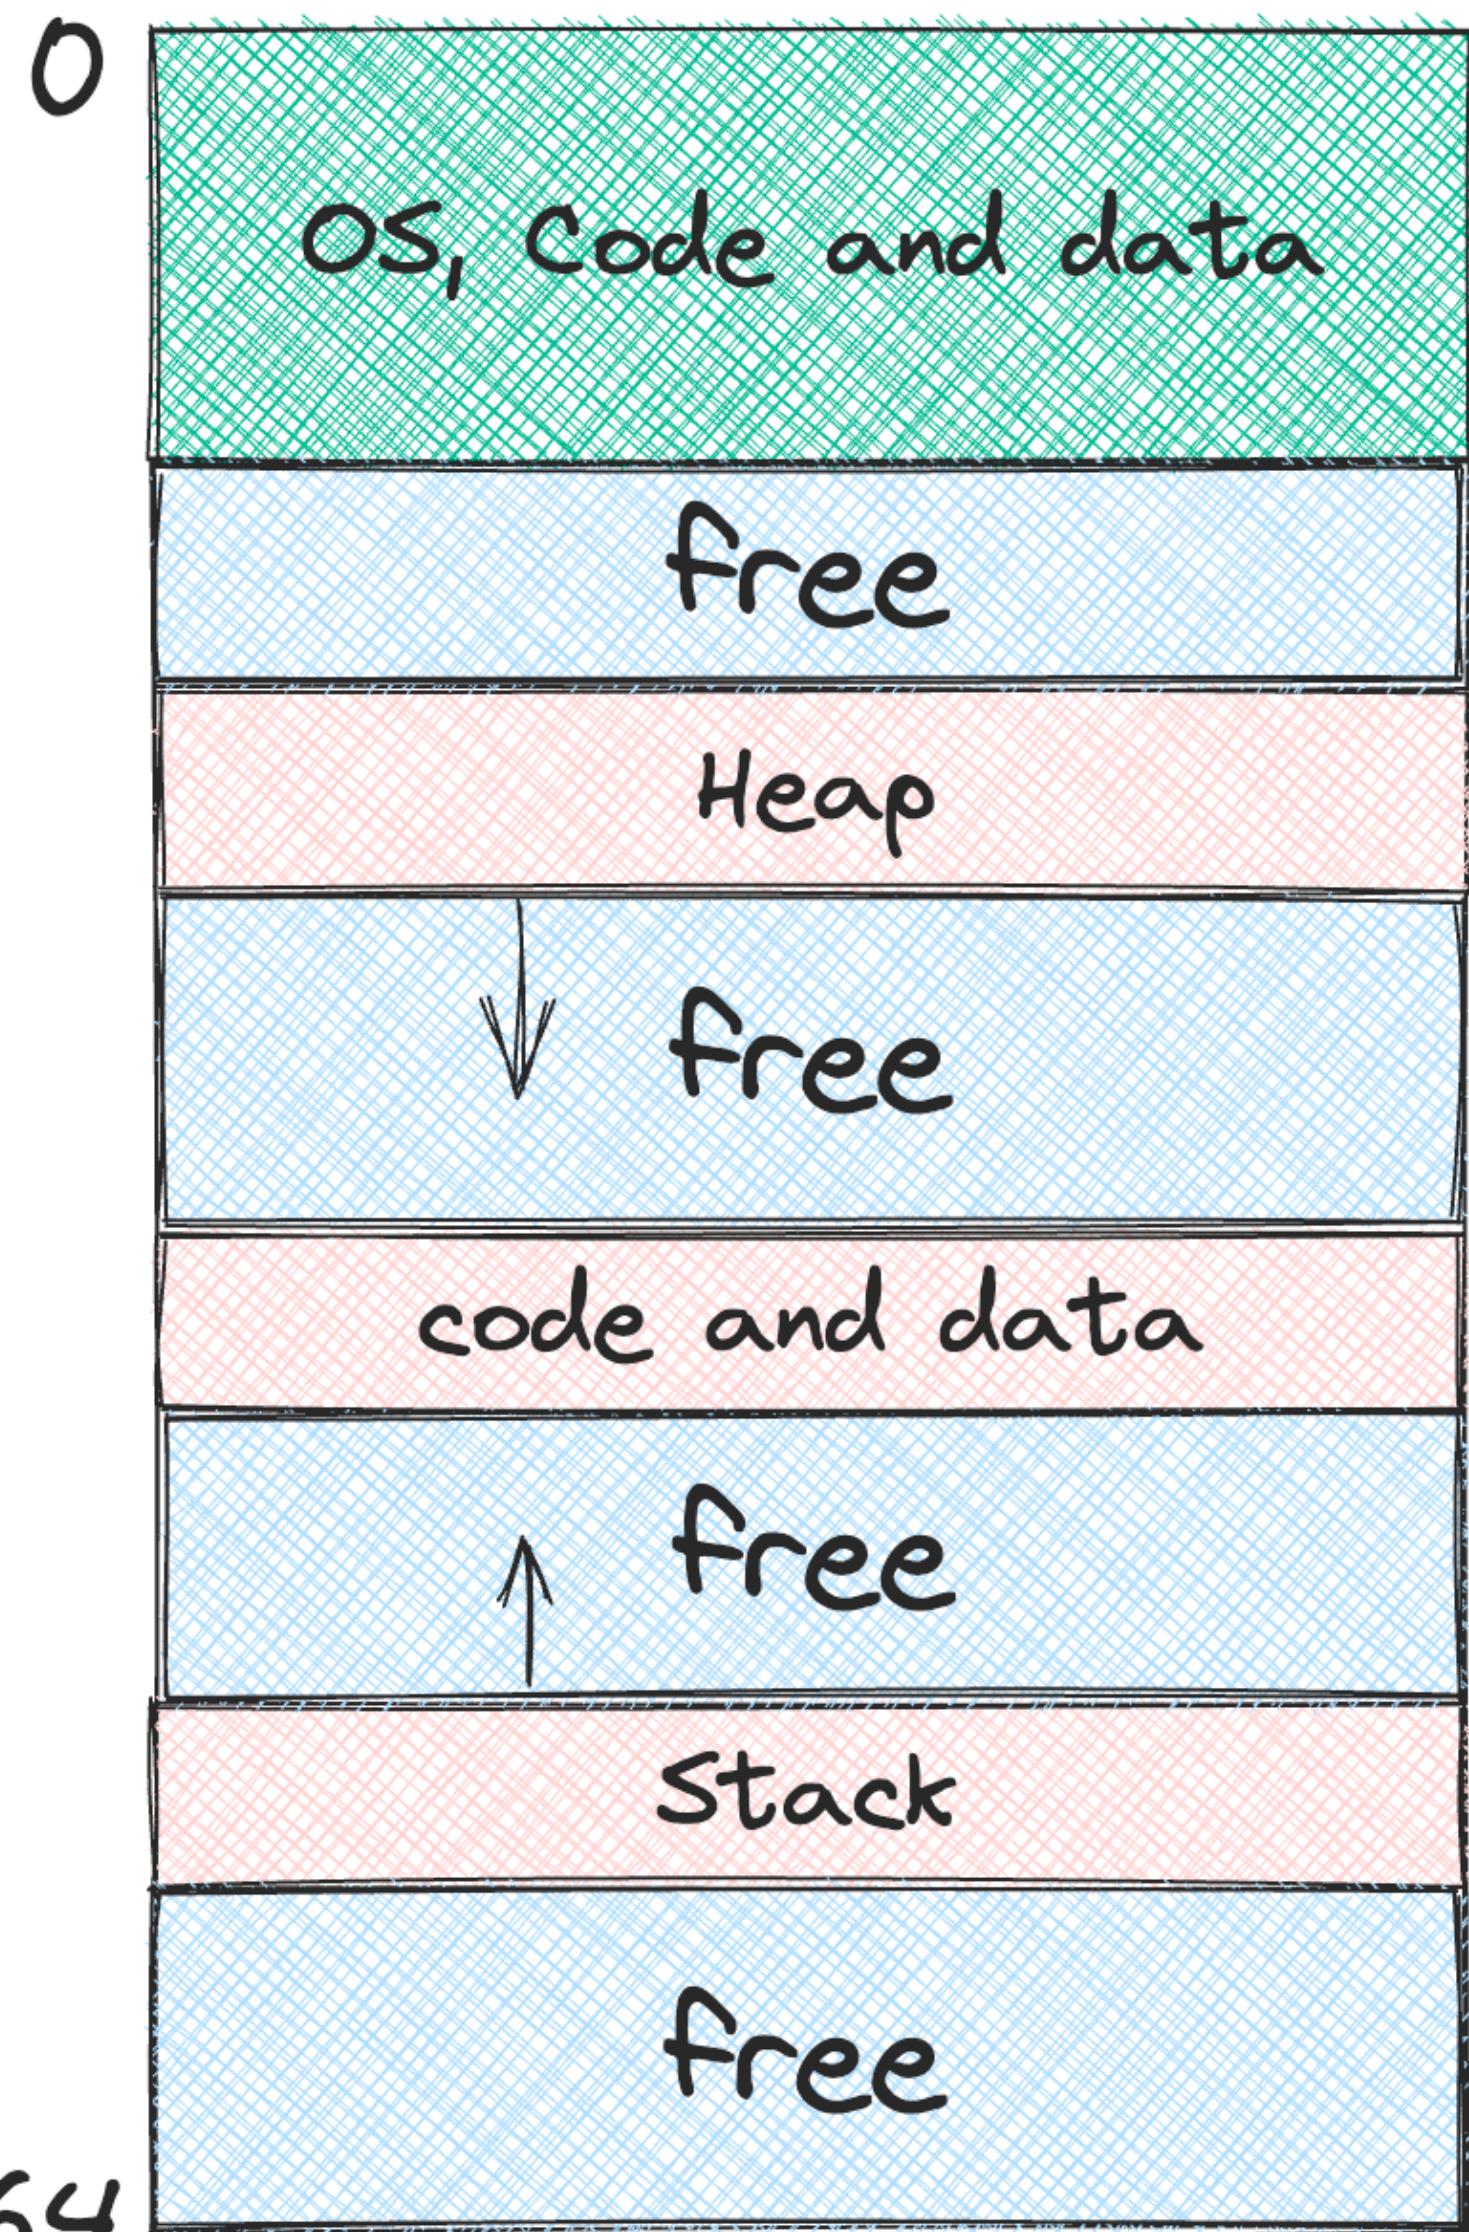
Dividing VA and PA into fixed chunks
(Translation with the help of a page
table)



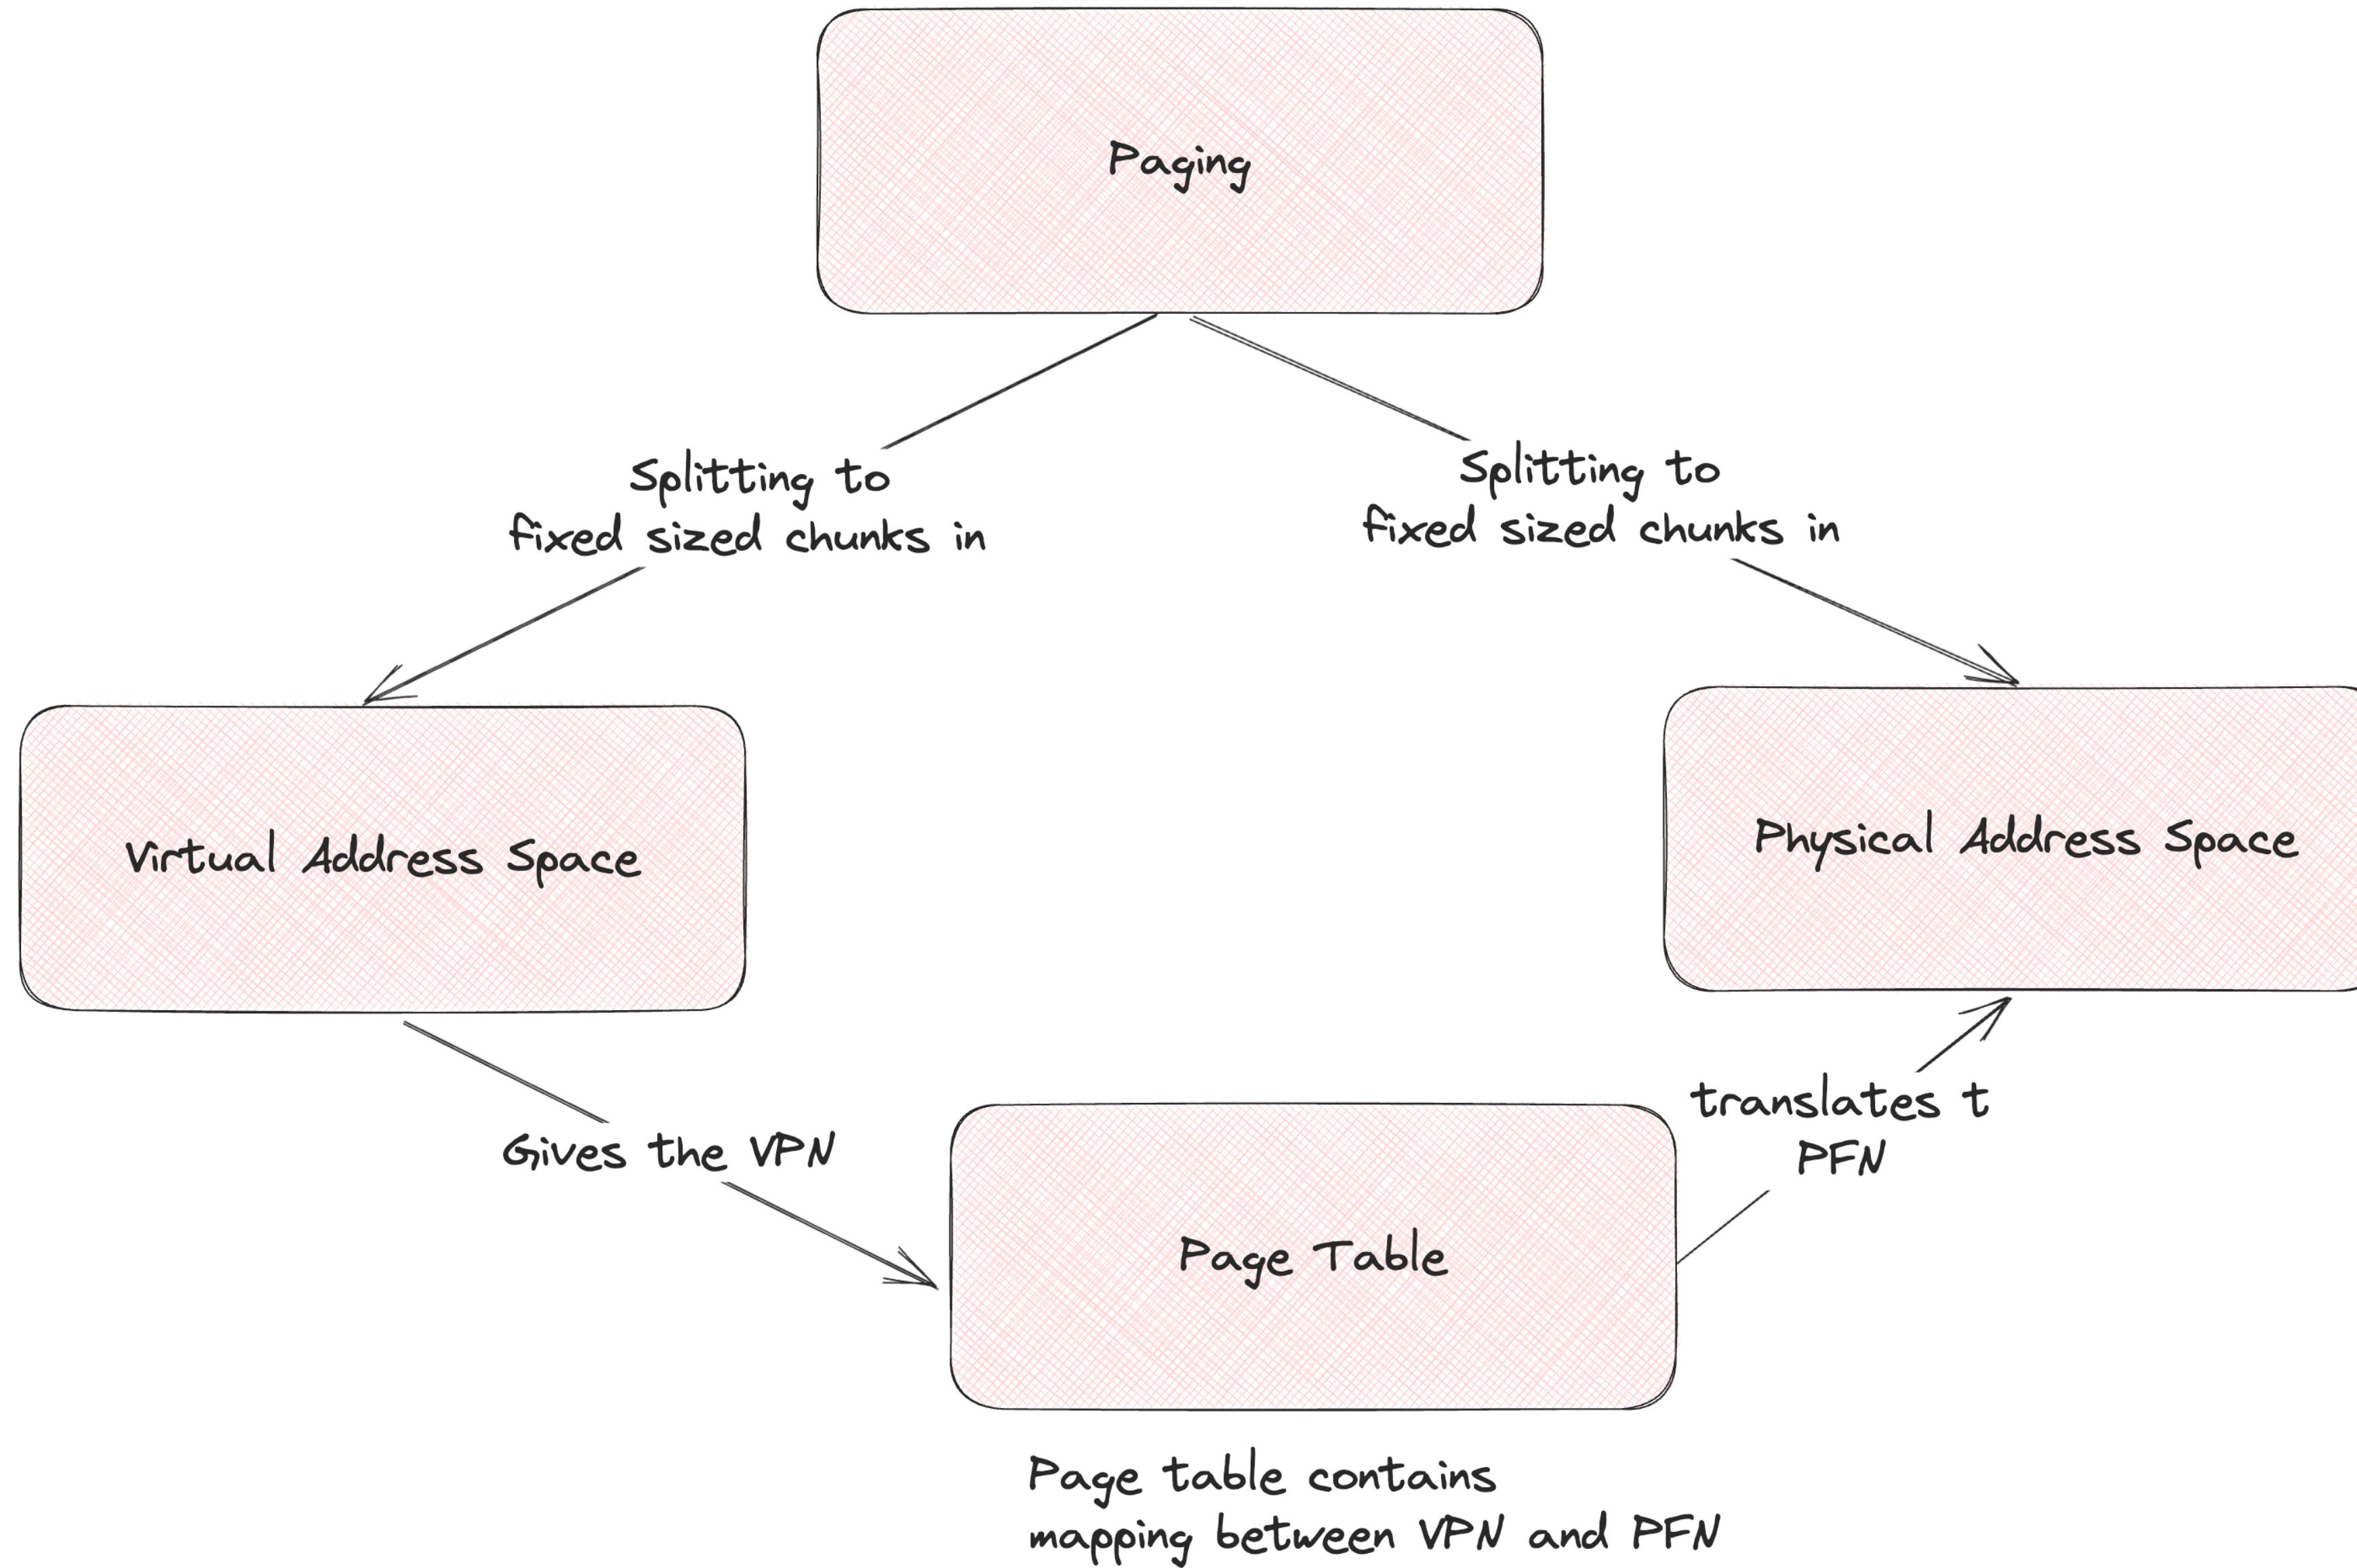
Segmentation

Generalized Base and Bounds

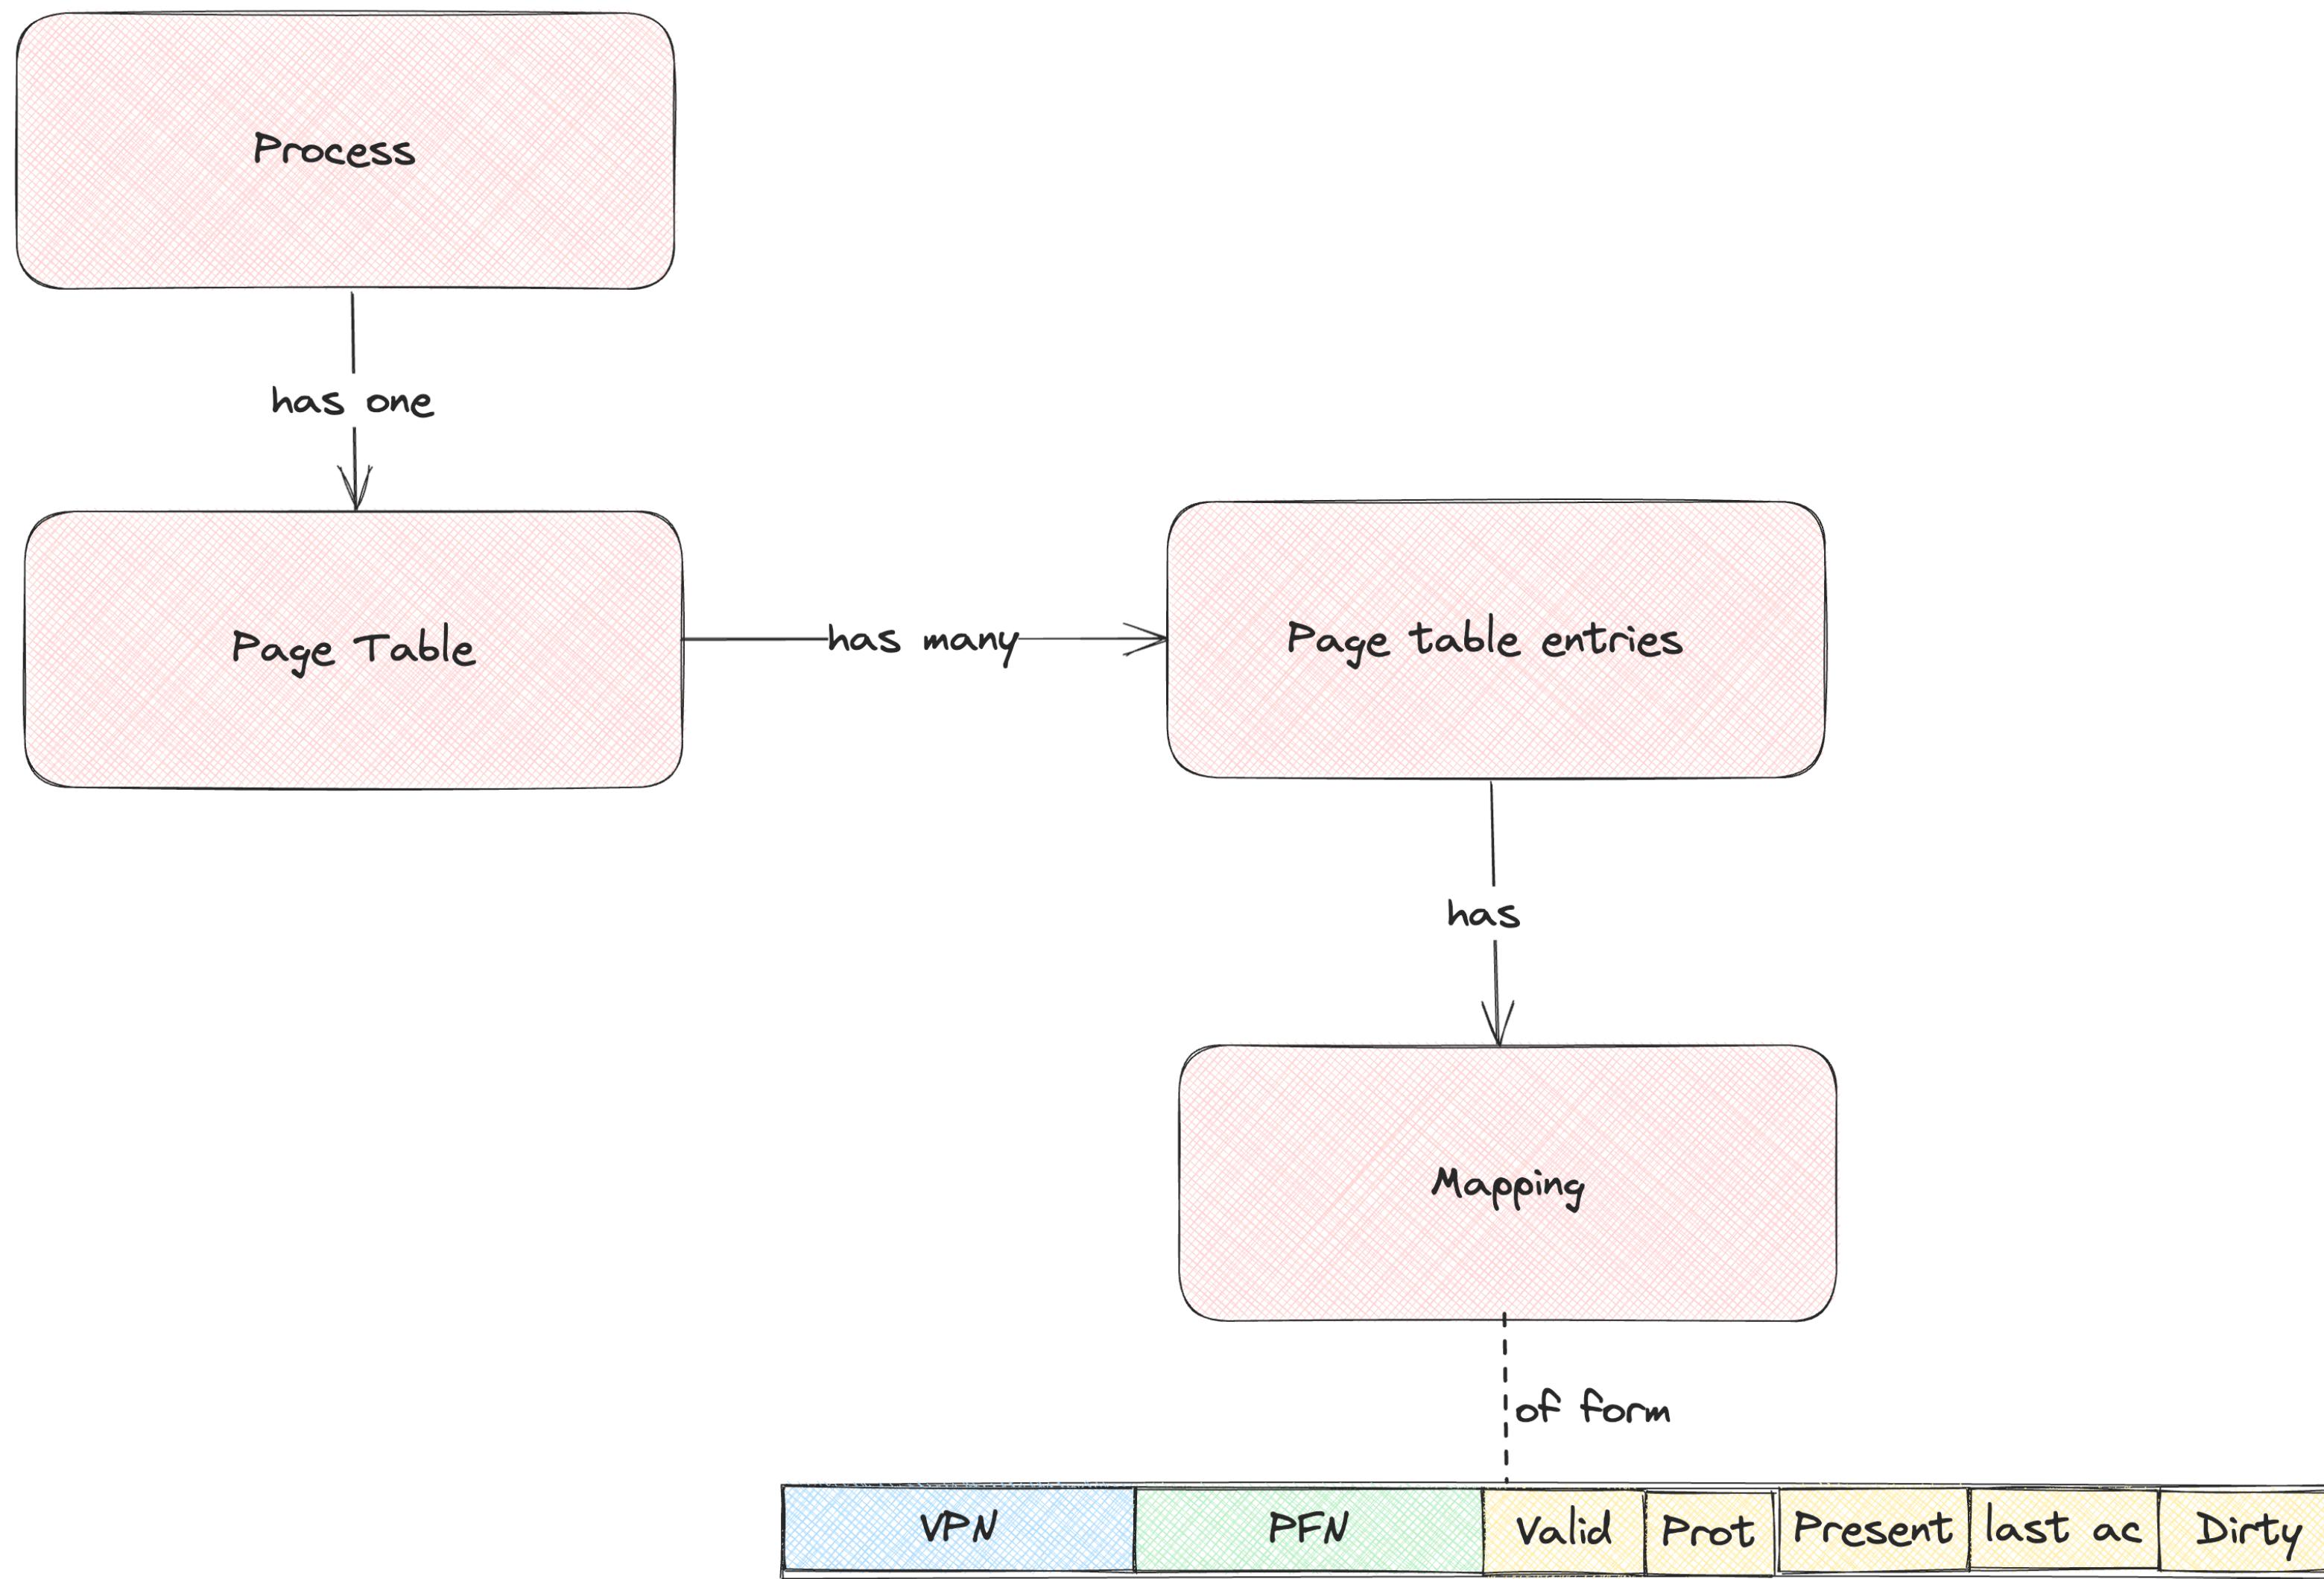
- Only used memory is allocated in physical memory
 - Allows allocating large address space
 - Sparse address space
- Different segments per process - code, stack, heap
- **For translation:** use first bits to identify segments and perform translation
- Results in **External fragmentation**



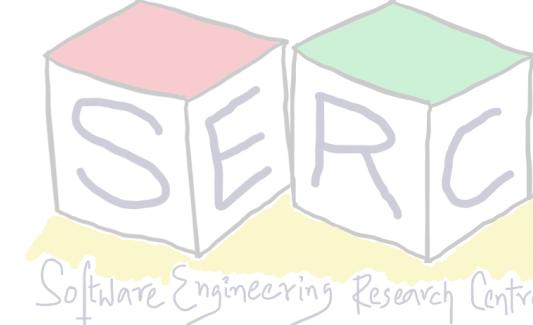
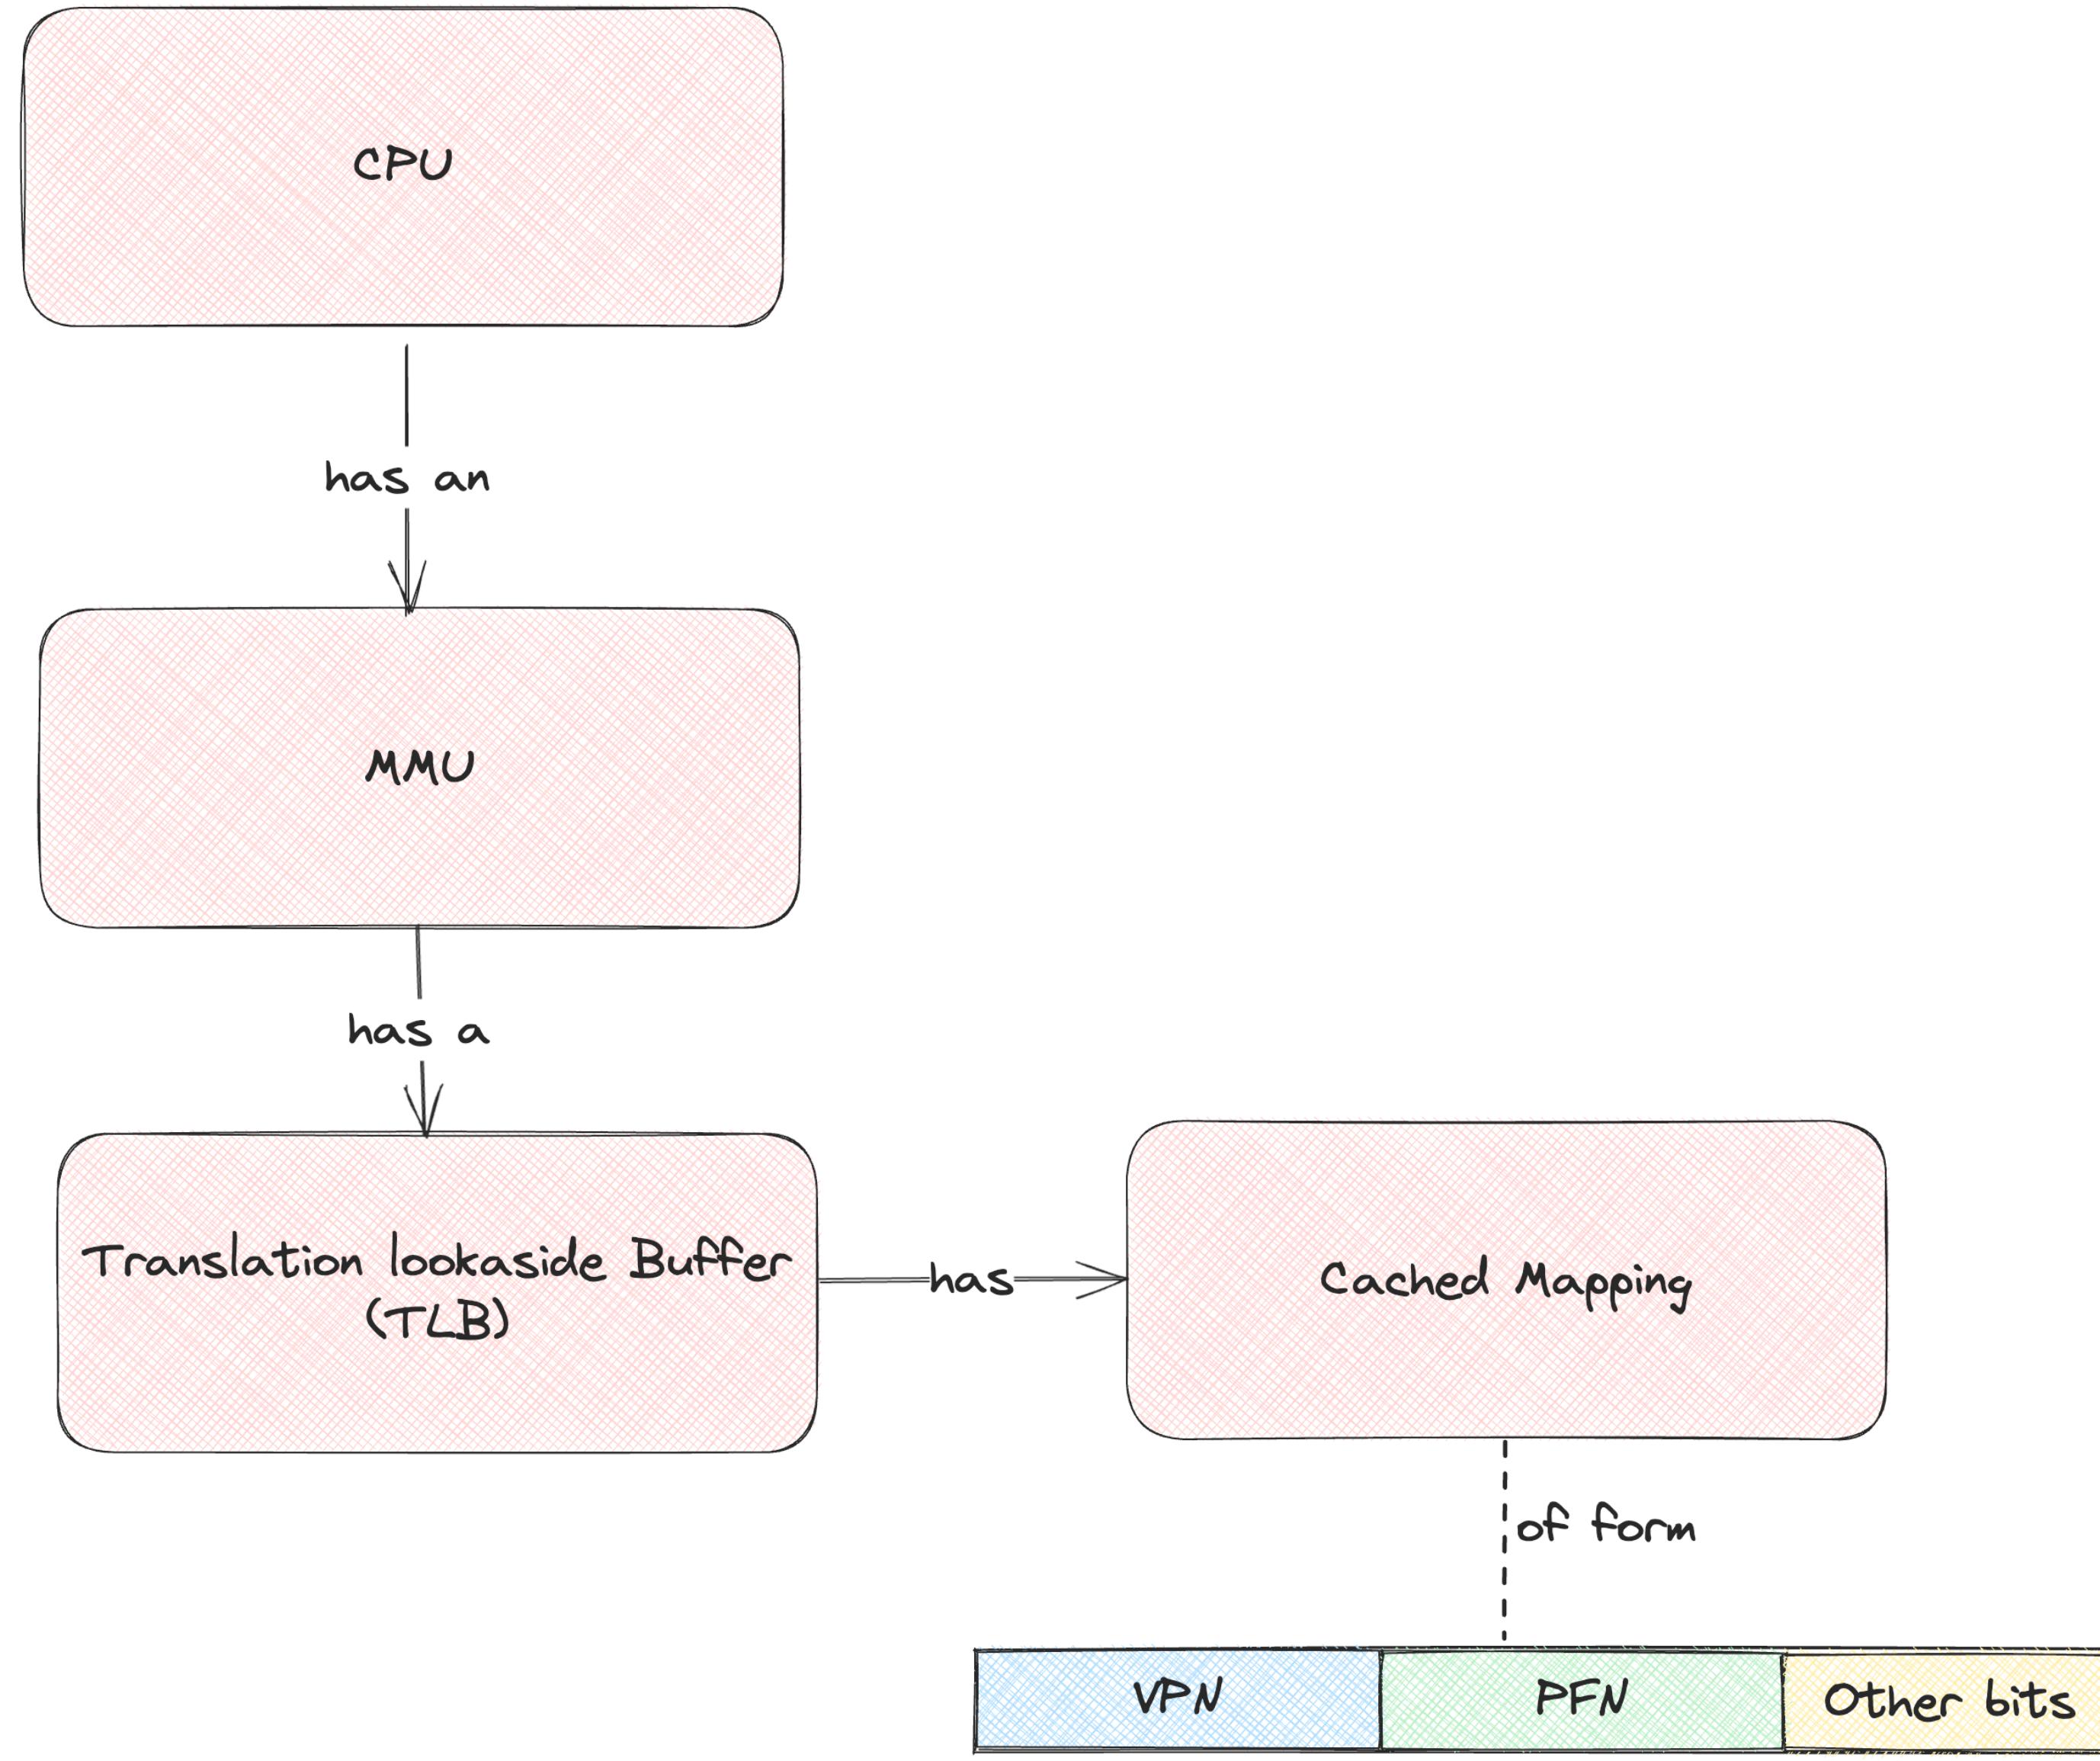
Paging

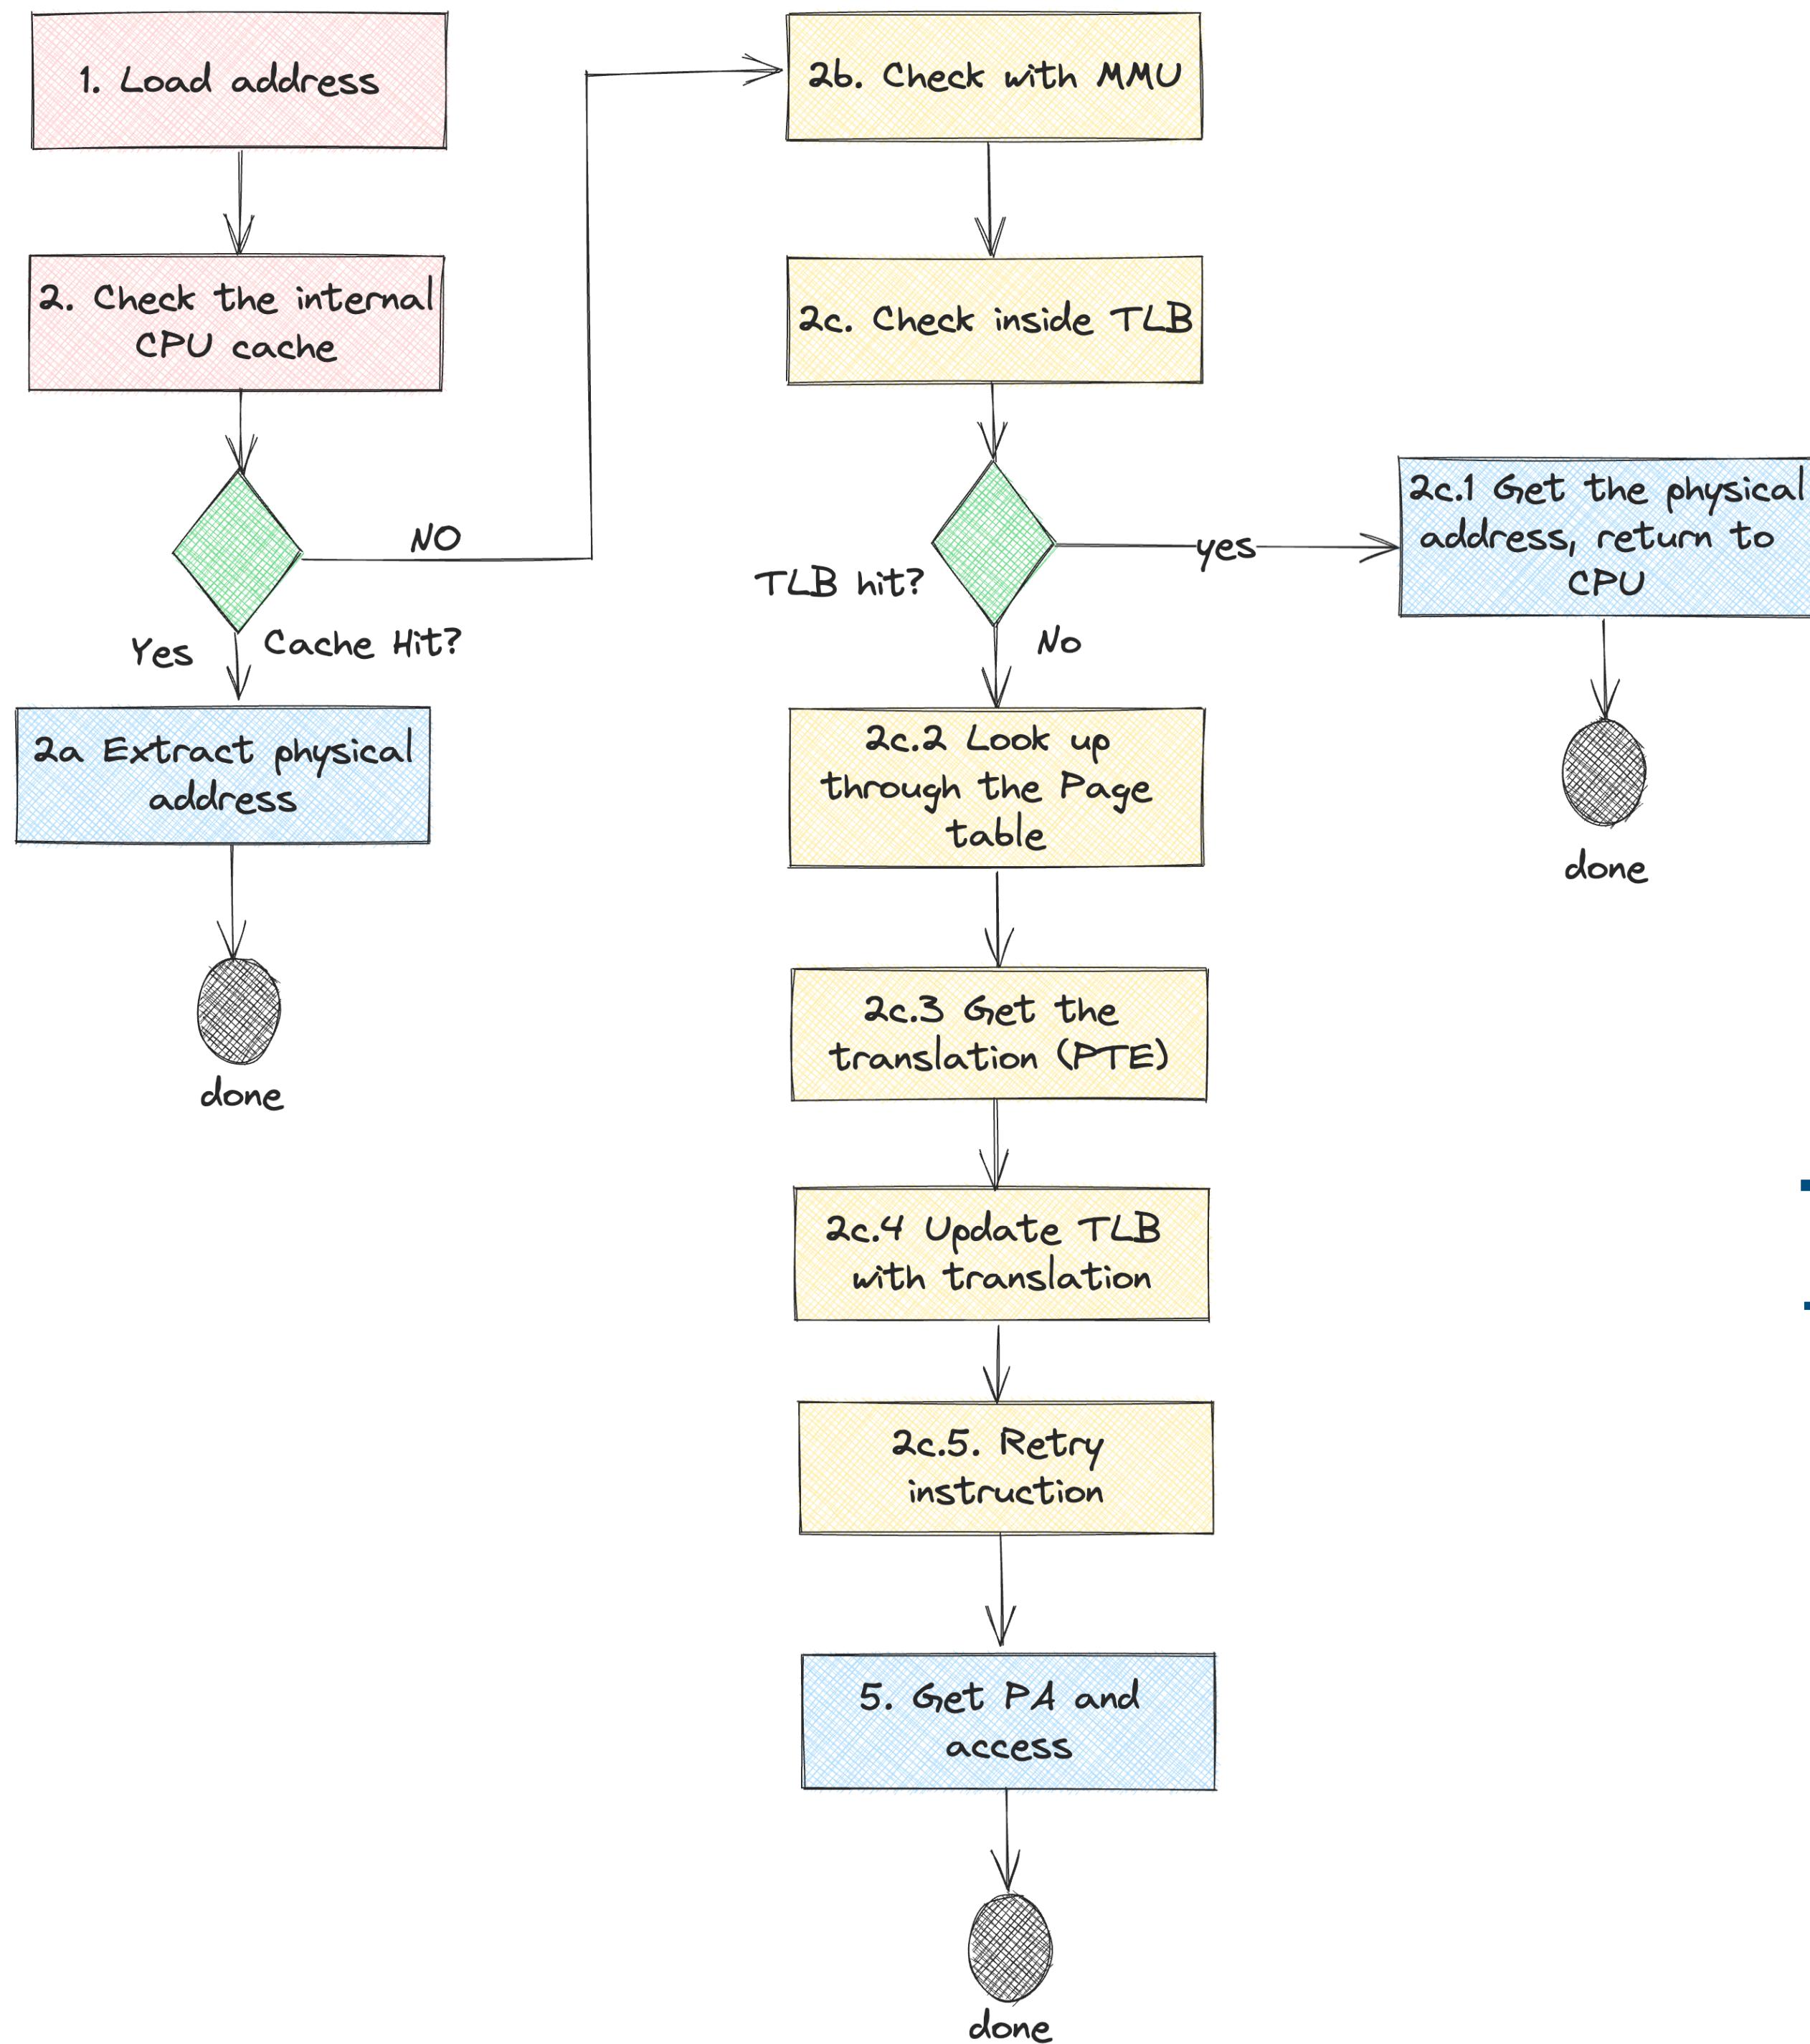


Page Tables

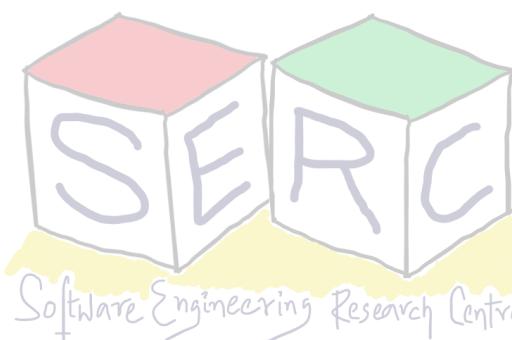


Paging - TLB



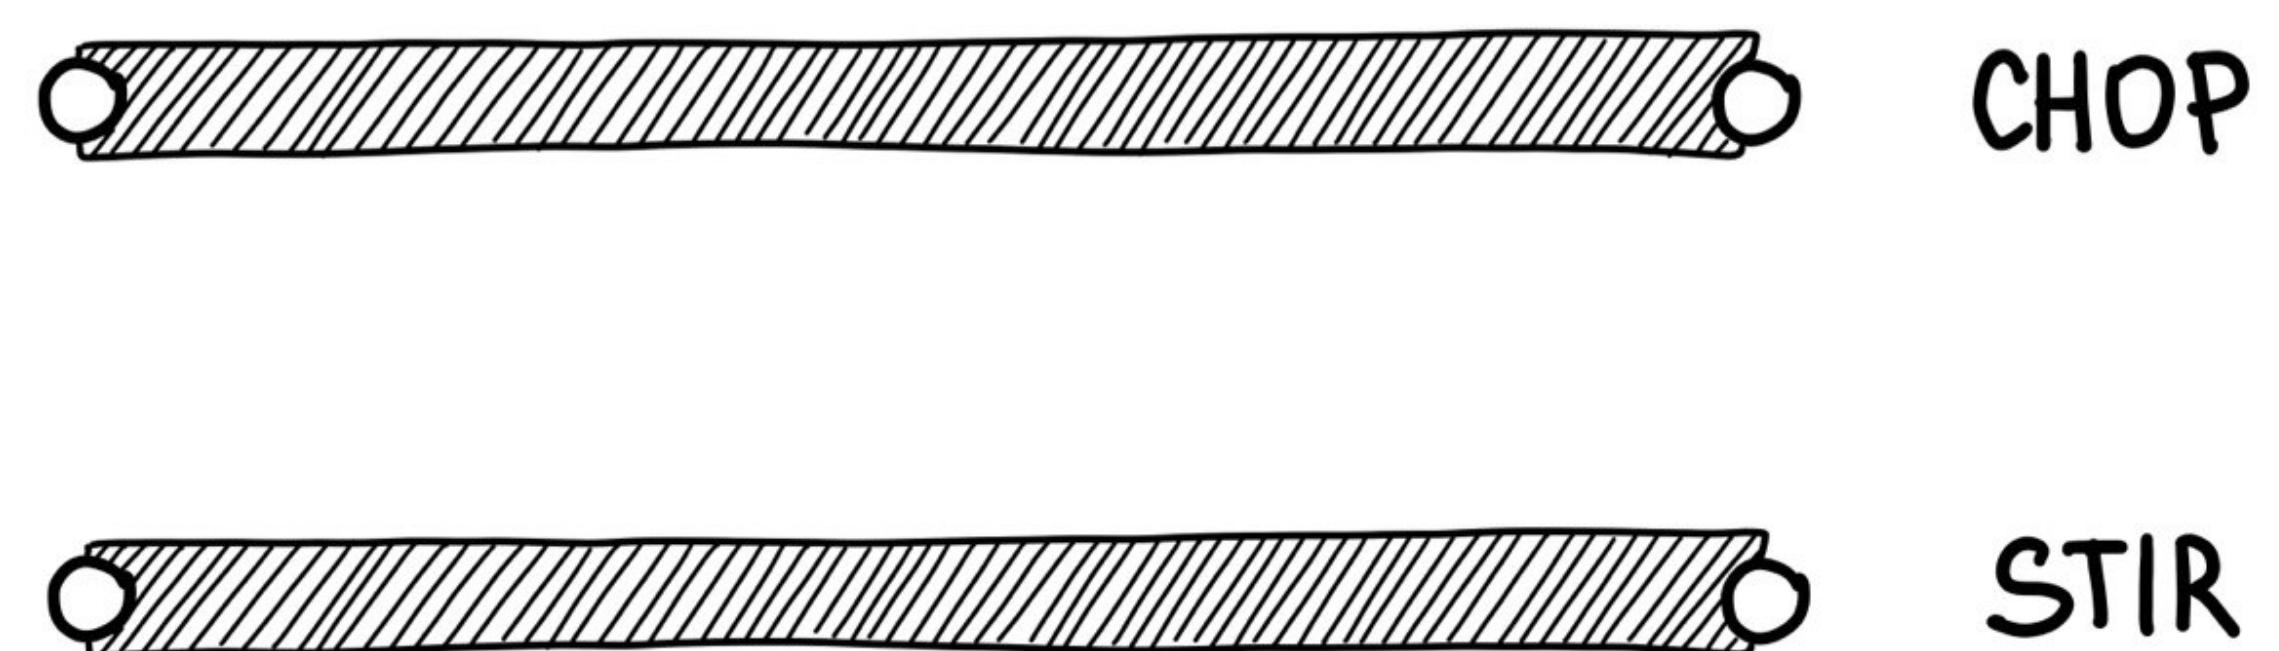
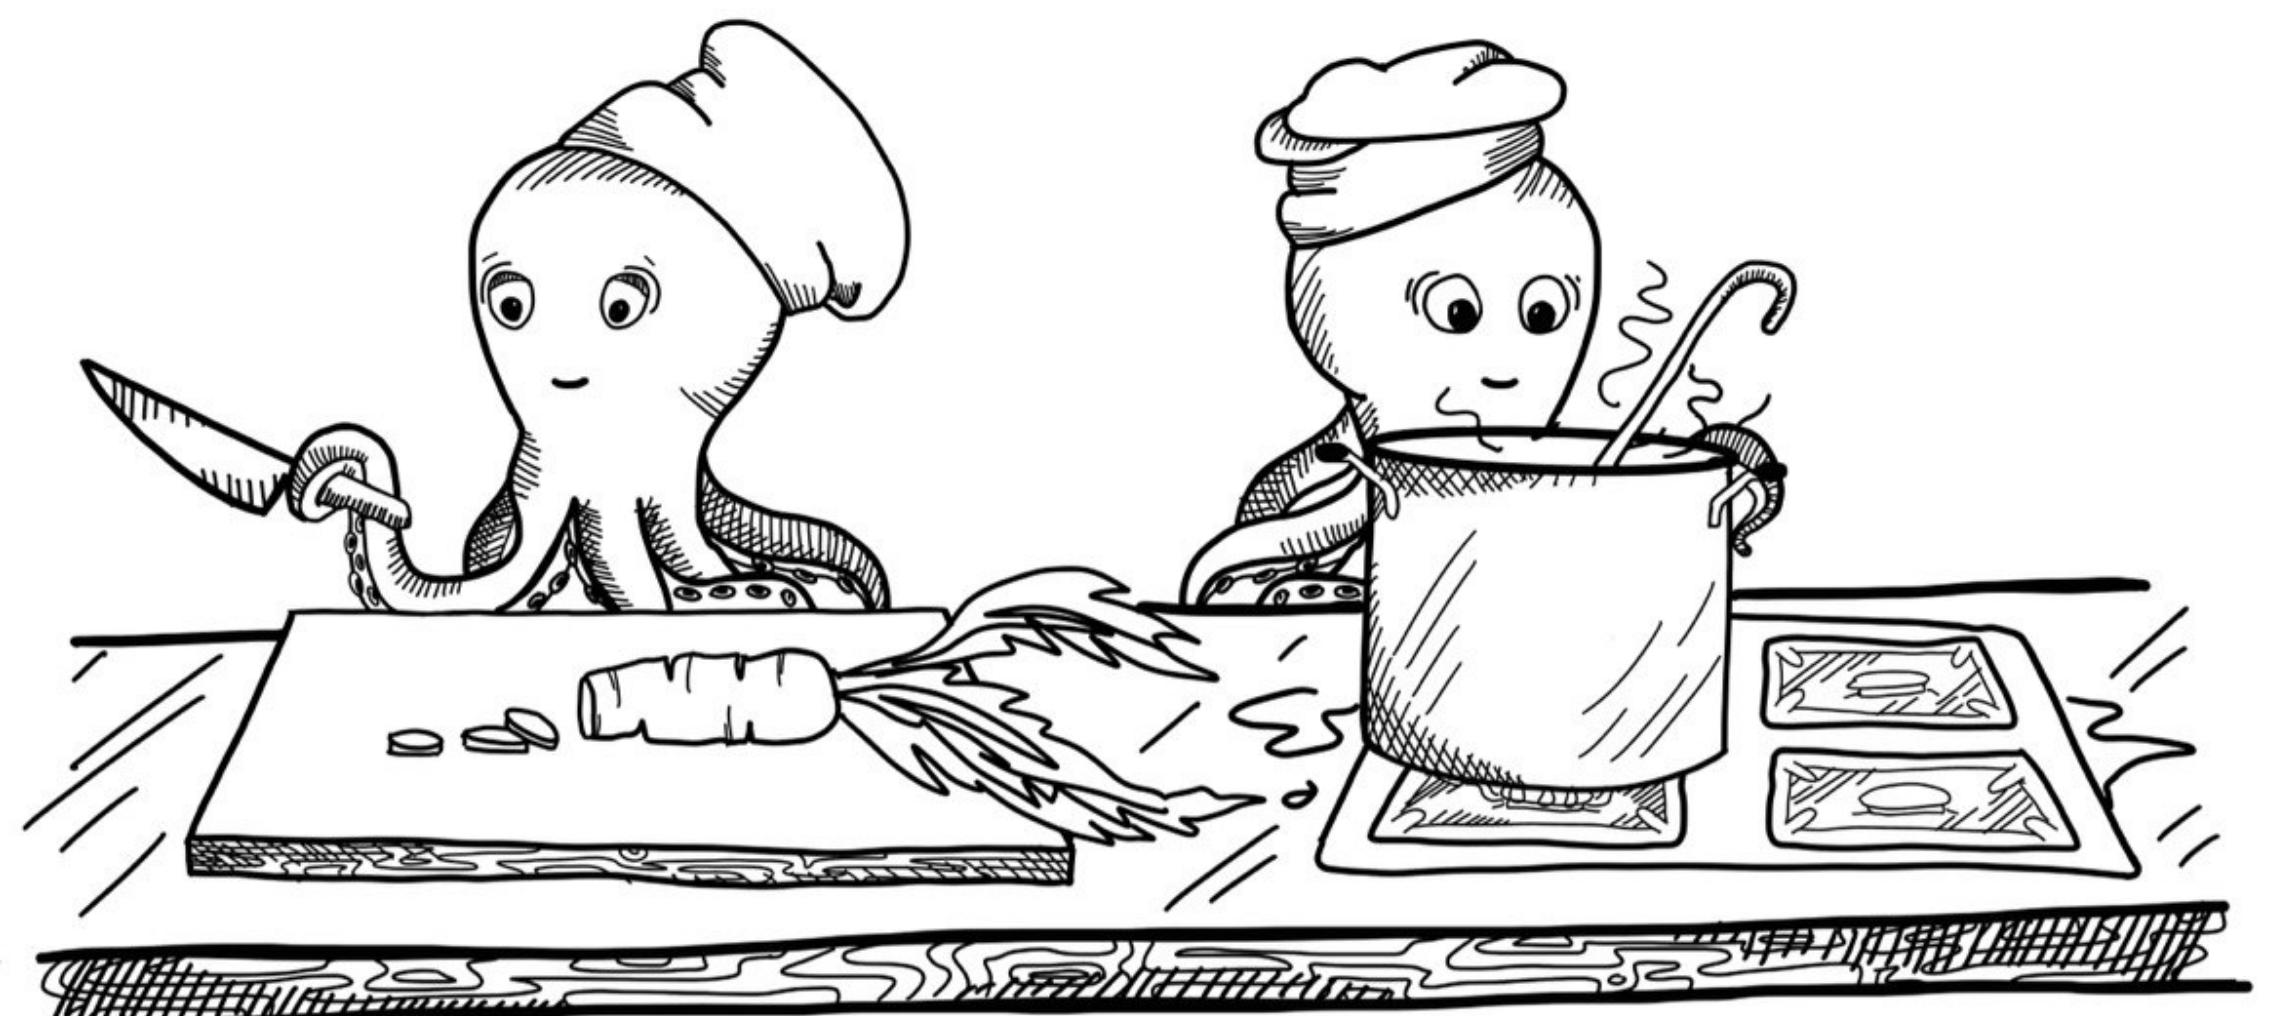
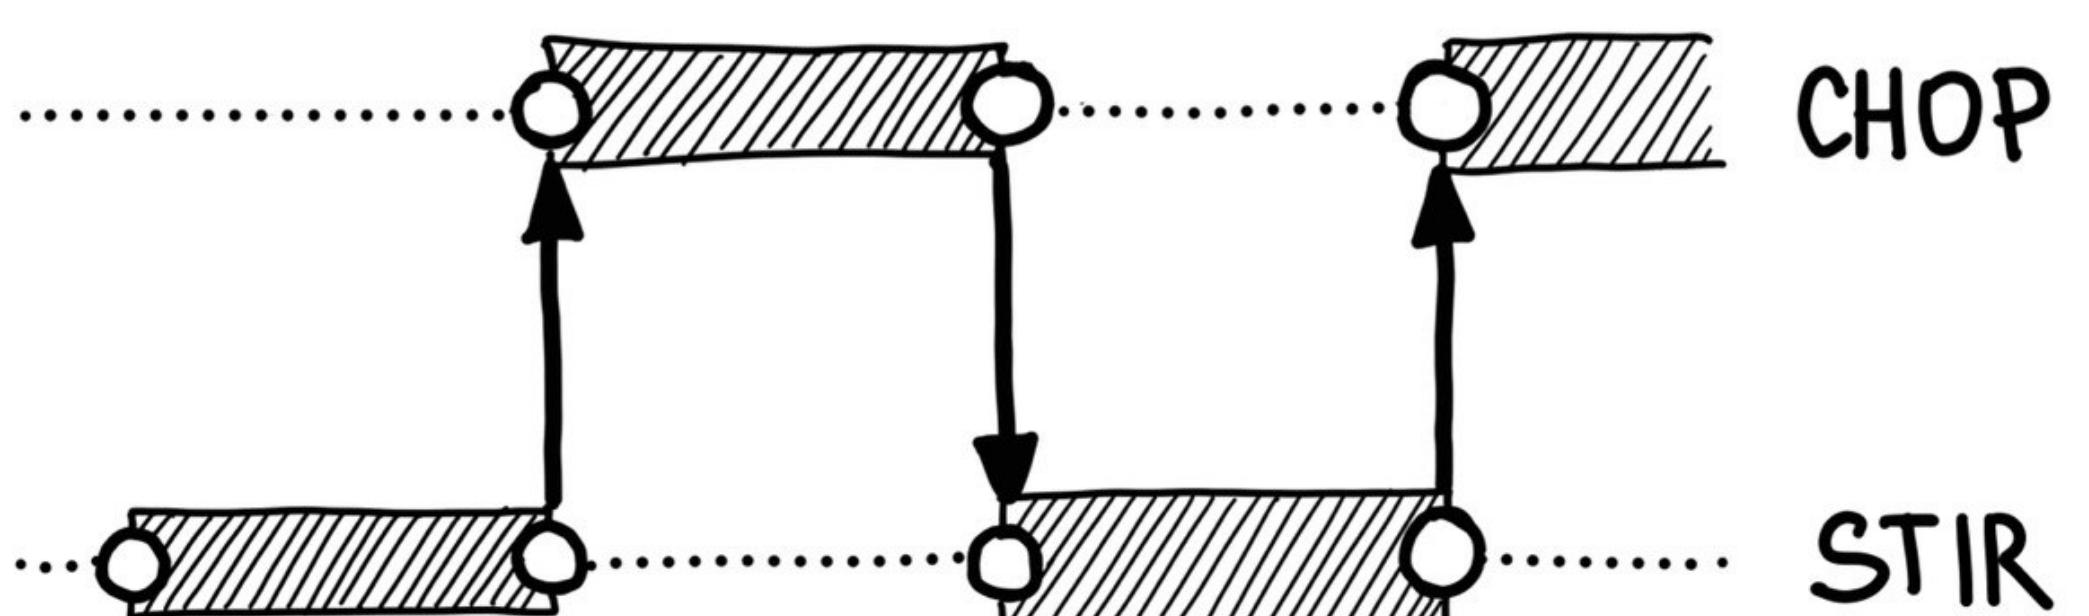
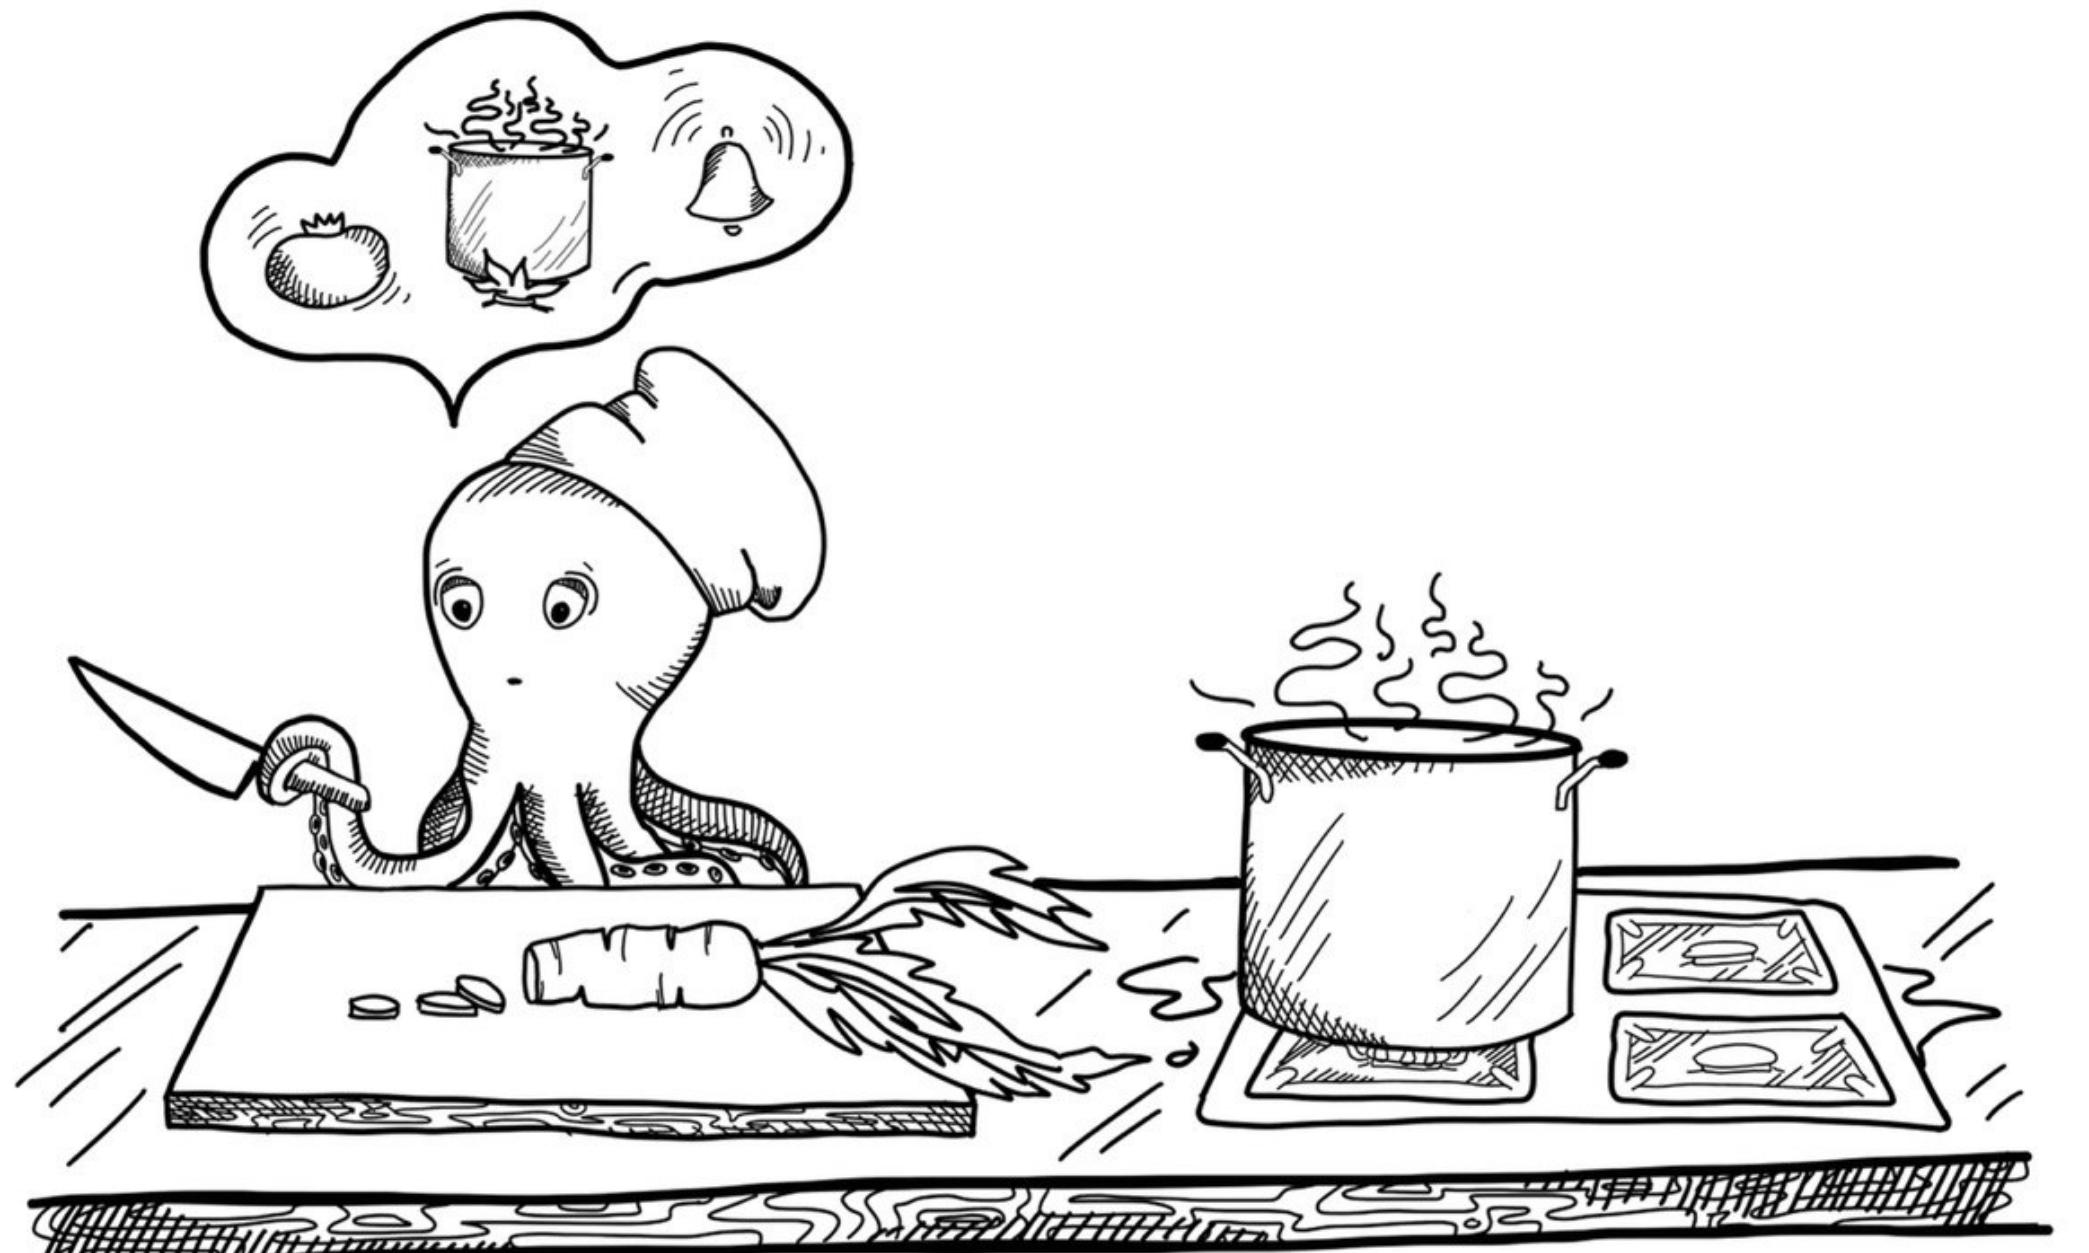


The overall address translation Process

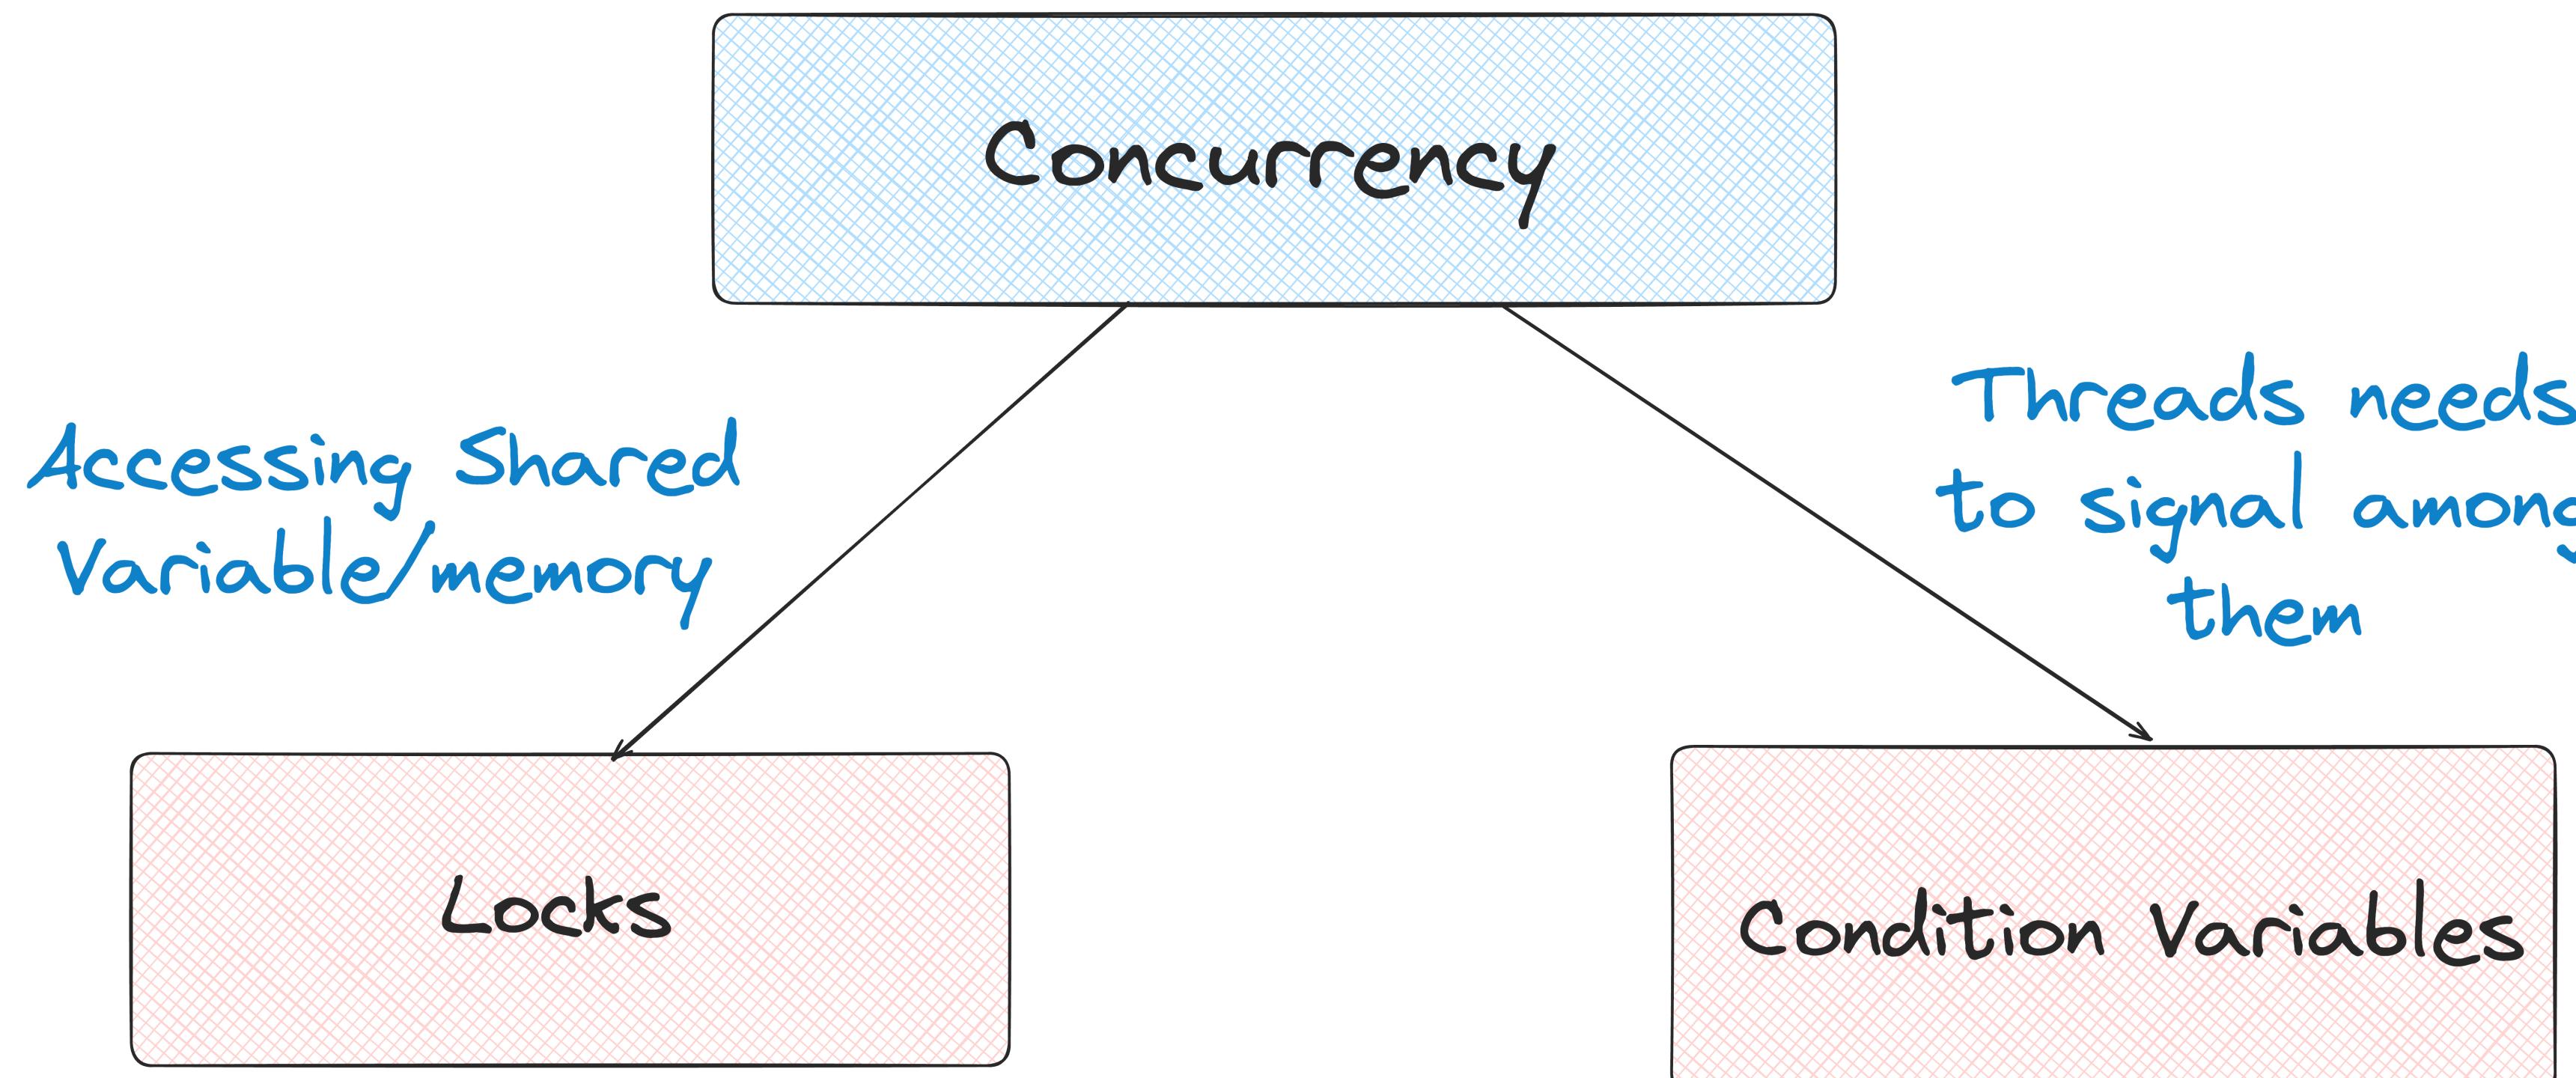


Concurrency and Parallelism

What is what?



Locks and Condition Variables



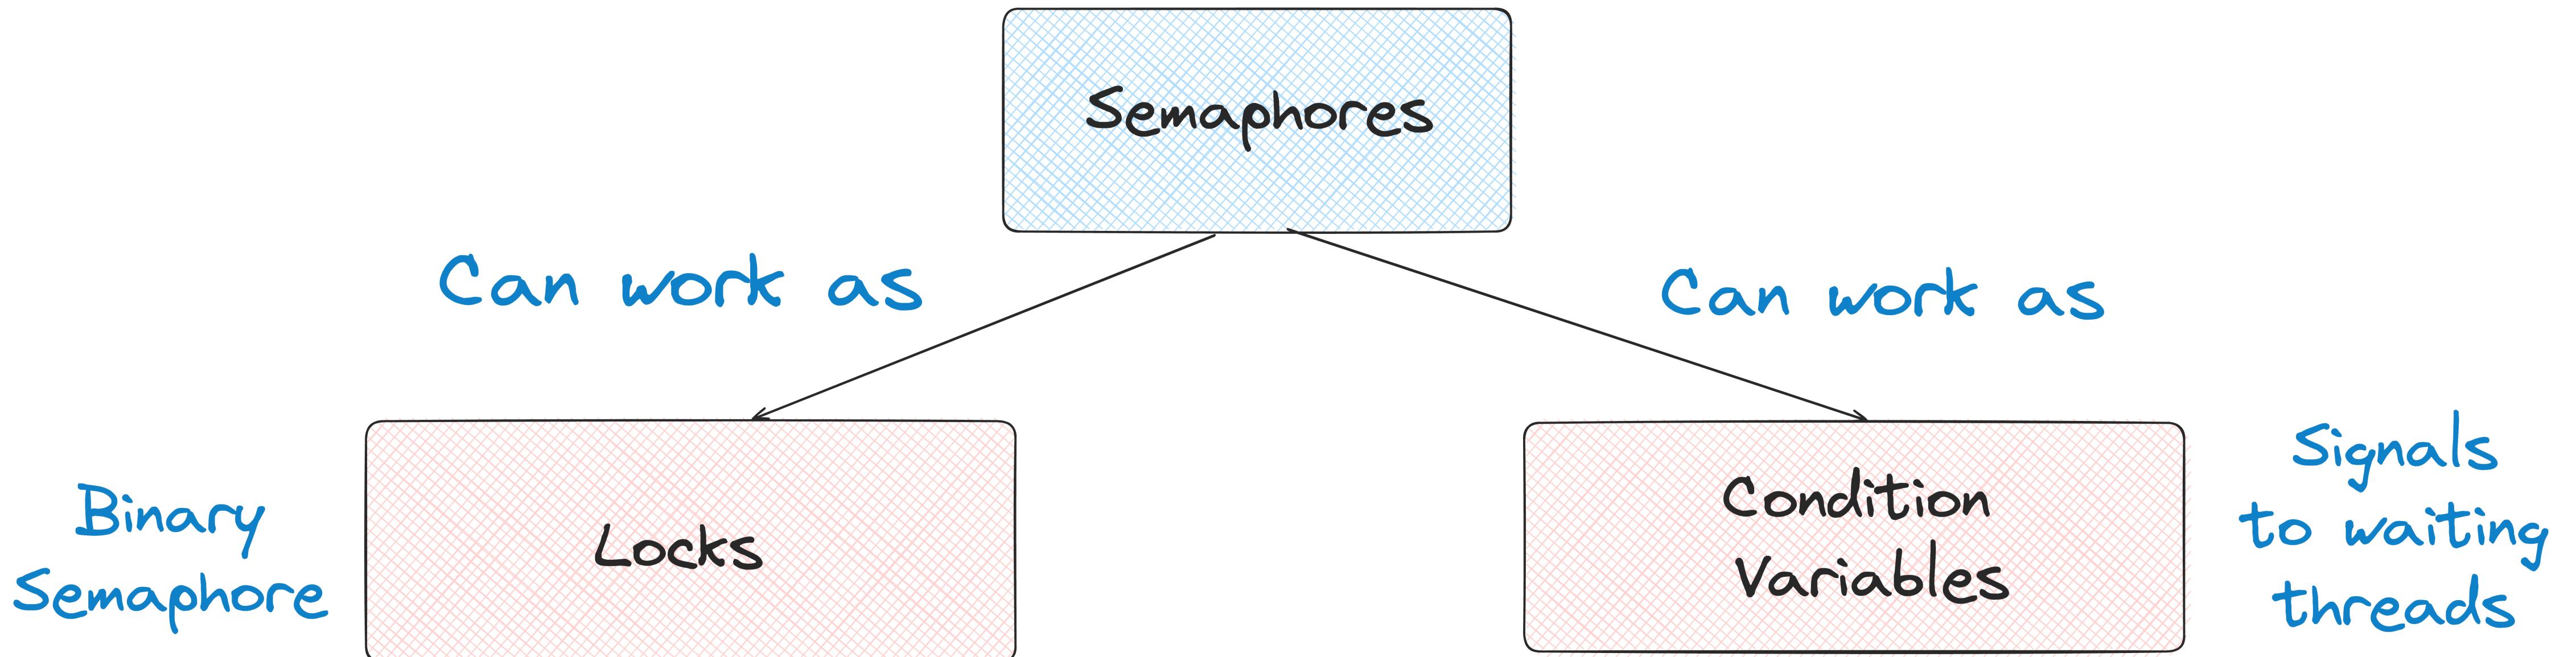
Hardware Primitives
Software locks

Use Signal and Wait
Adds threads to queues
when waiting

Combination of Locks and Condition variables may be required
to accomplish different tasks



Semaphores



Two key operations: wait () and post()
Initialization of semaphore holds the key



Producer Consumer Problem Using Semaphores

- Let us start with 2 semaphores: empty and wait, Buffer with MAX = 1

● ● ● Get and Put for large sized buffer

```
int buffer[MAX];
int fill = 0;
int use = 0;
int count = 0;

void put (int value)
{
    buffer[fill] = value;
    fill = (fill + 1)%MAX;
    count++;
}

int get()
{
    int tmp = buffer[use];
    use = (use + 1)%MAX;
    count--;
    return tmp;
}
```



● ● ● Producer-Consumer with buffer

```
sem_t empty;
sem_t full;

void *producer(void *arg)
{
    int i;
    int maxLoops = (int)arg;
    for (i=0;i<maxLoops;i++)
    {
        sem_wait(&empty);
        put (i);
        sem_post(&full);
    }
}

void *consumer(void *arg)
{
    int i;
    int maxLoops = (int)arg;
    for (i=0;i<maxLoops;i++)
    {
        sem_wait(&full);
        int tmp = get();
        sem_post(&empty);
        printf("%d\n", tmp);
    }
}
```

Producer Consumer Problem Using Semaphores

The Solution

Producer

```
sem_wait (&empty);
sem_wait (&mutex);
put(i);
sem_post (&mutex);
sem_post (&full);
```

Consumer

```
sem_wait (&full);
sem_wait (&mutex);
get();
sem_post (&mutex);
sem_post (&empty);
```

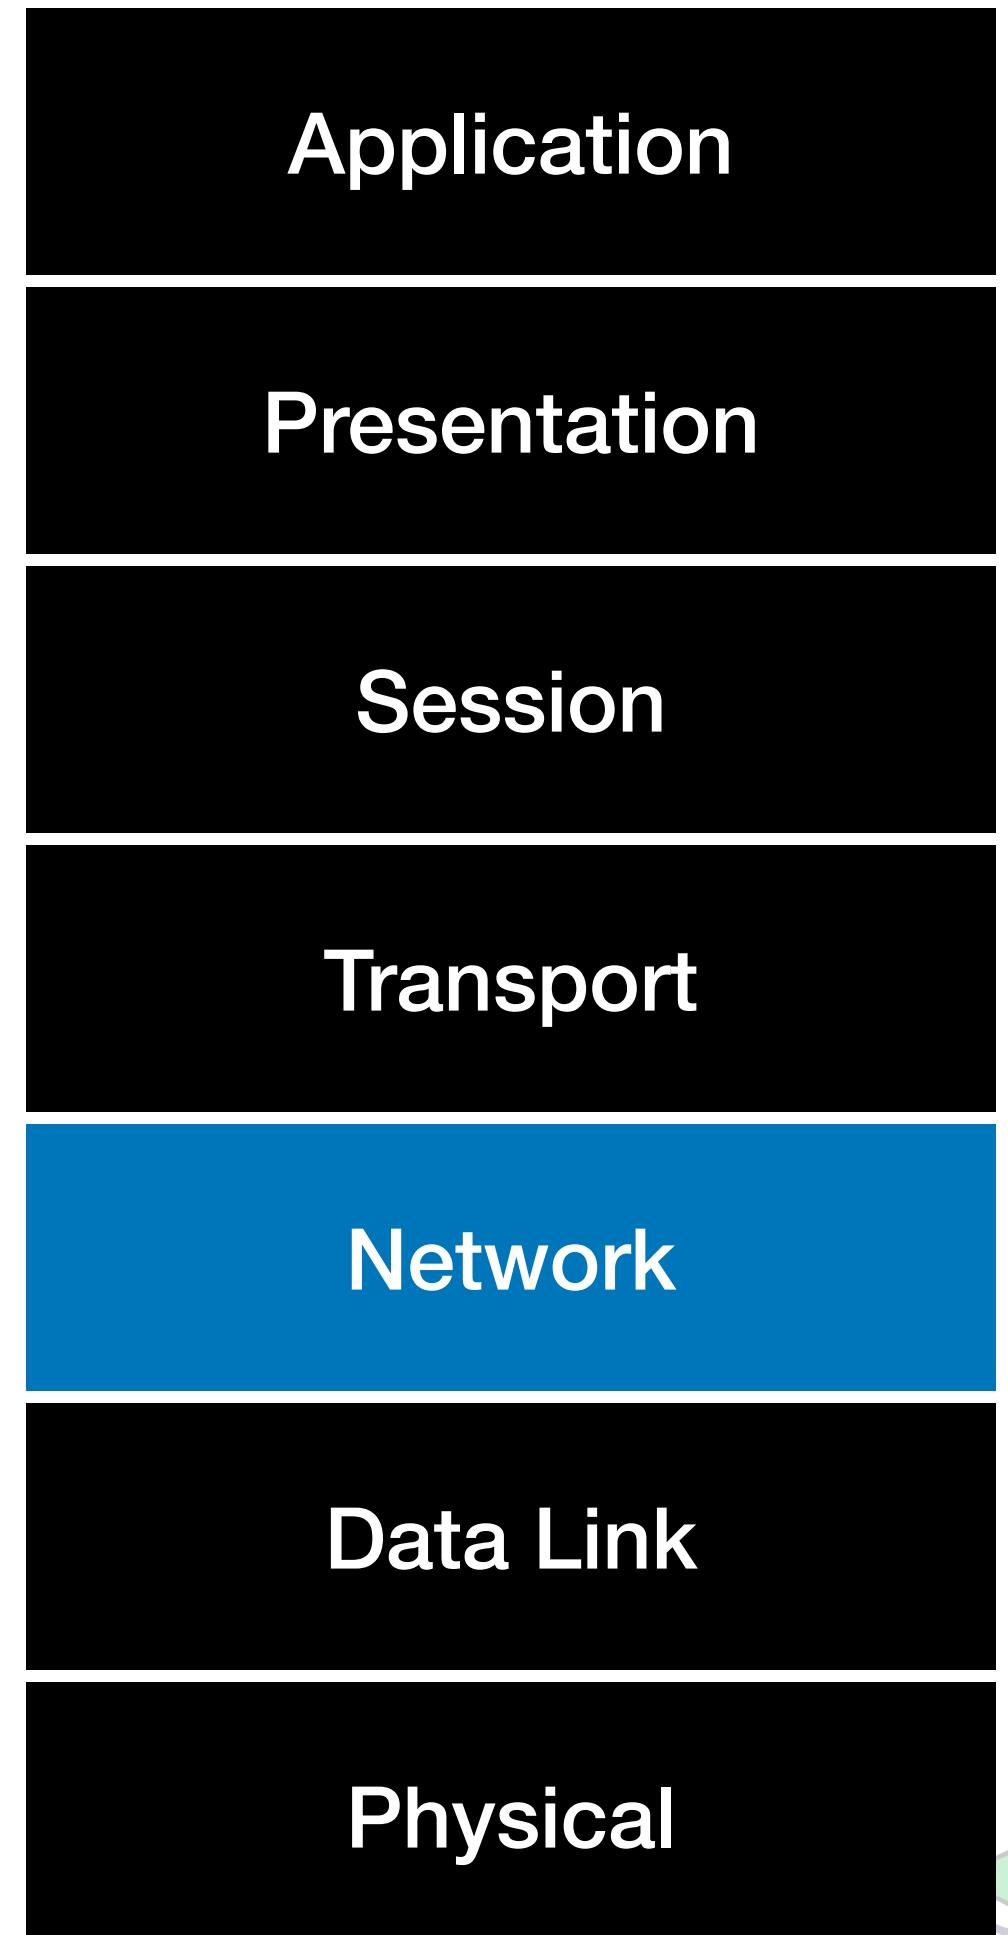
- Add mutex lock around put and get - Avoid deadlocks!
- Let producer and consumer get the signal and then lock when entering CS



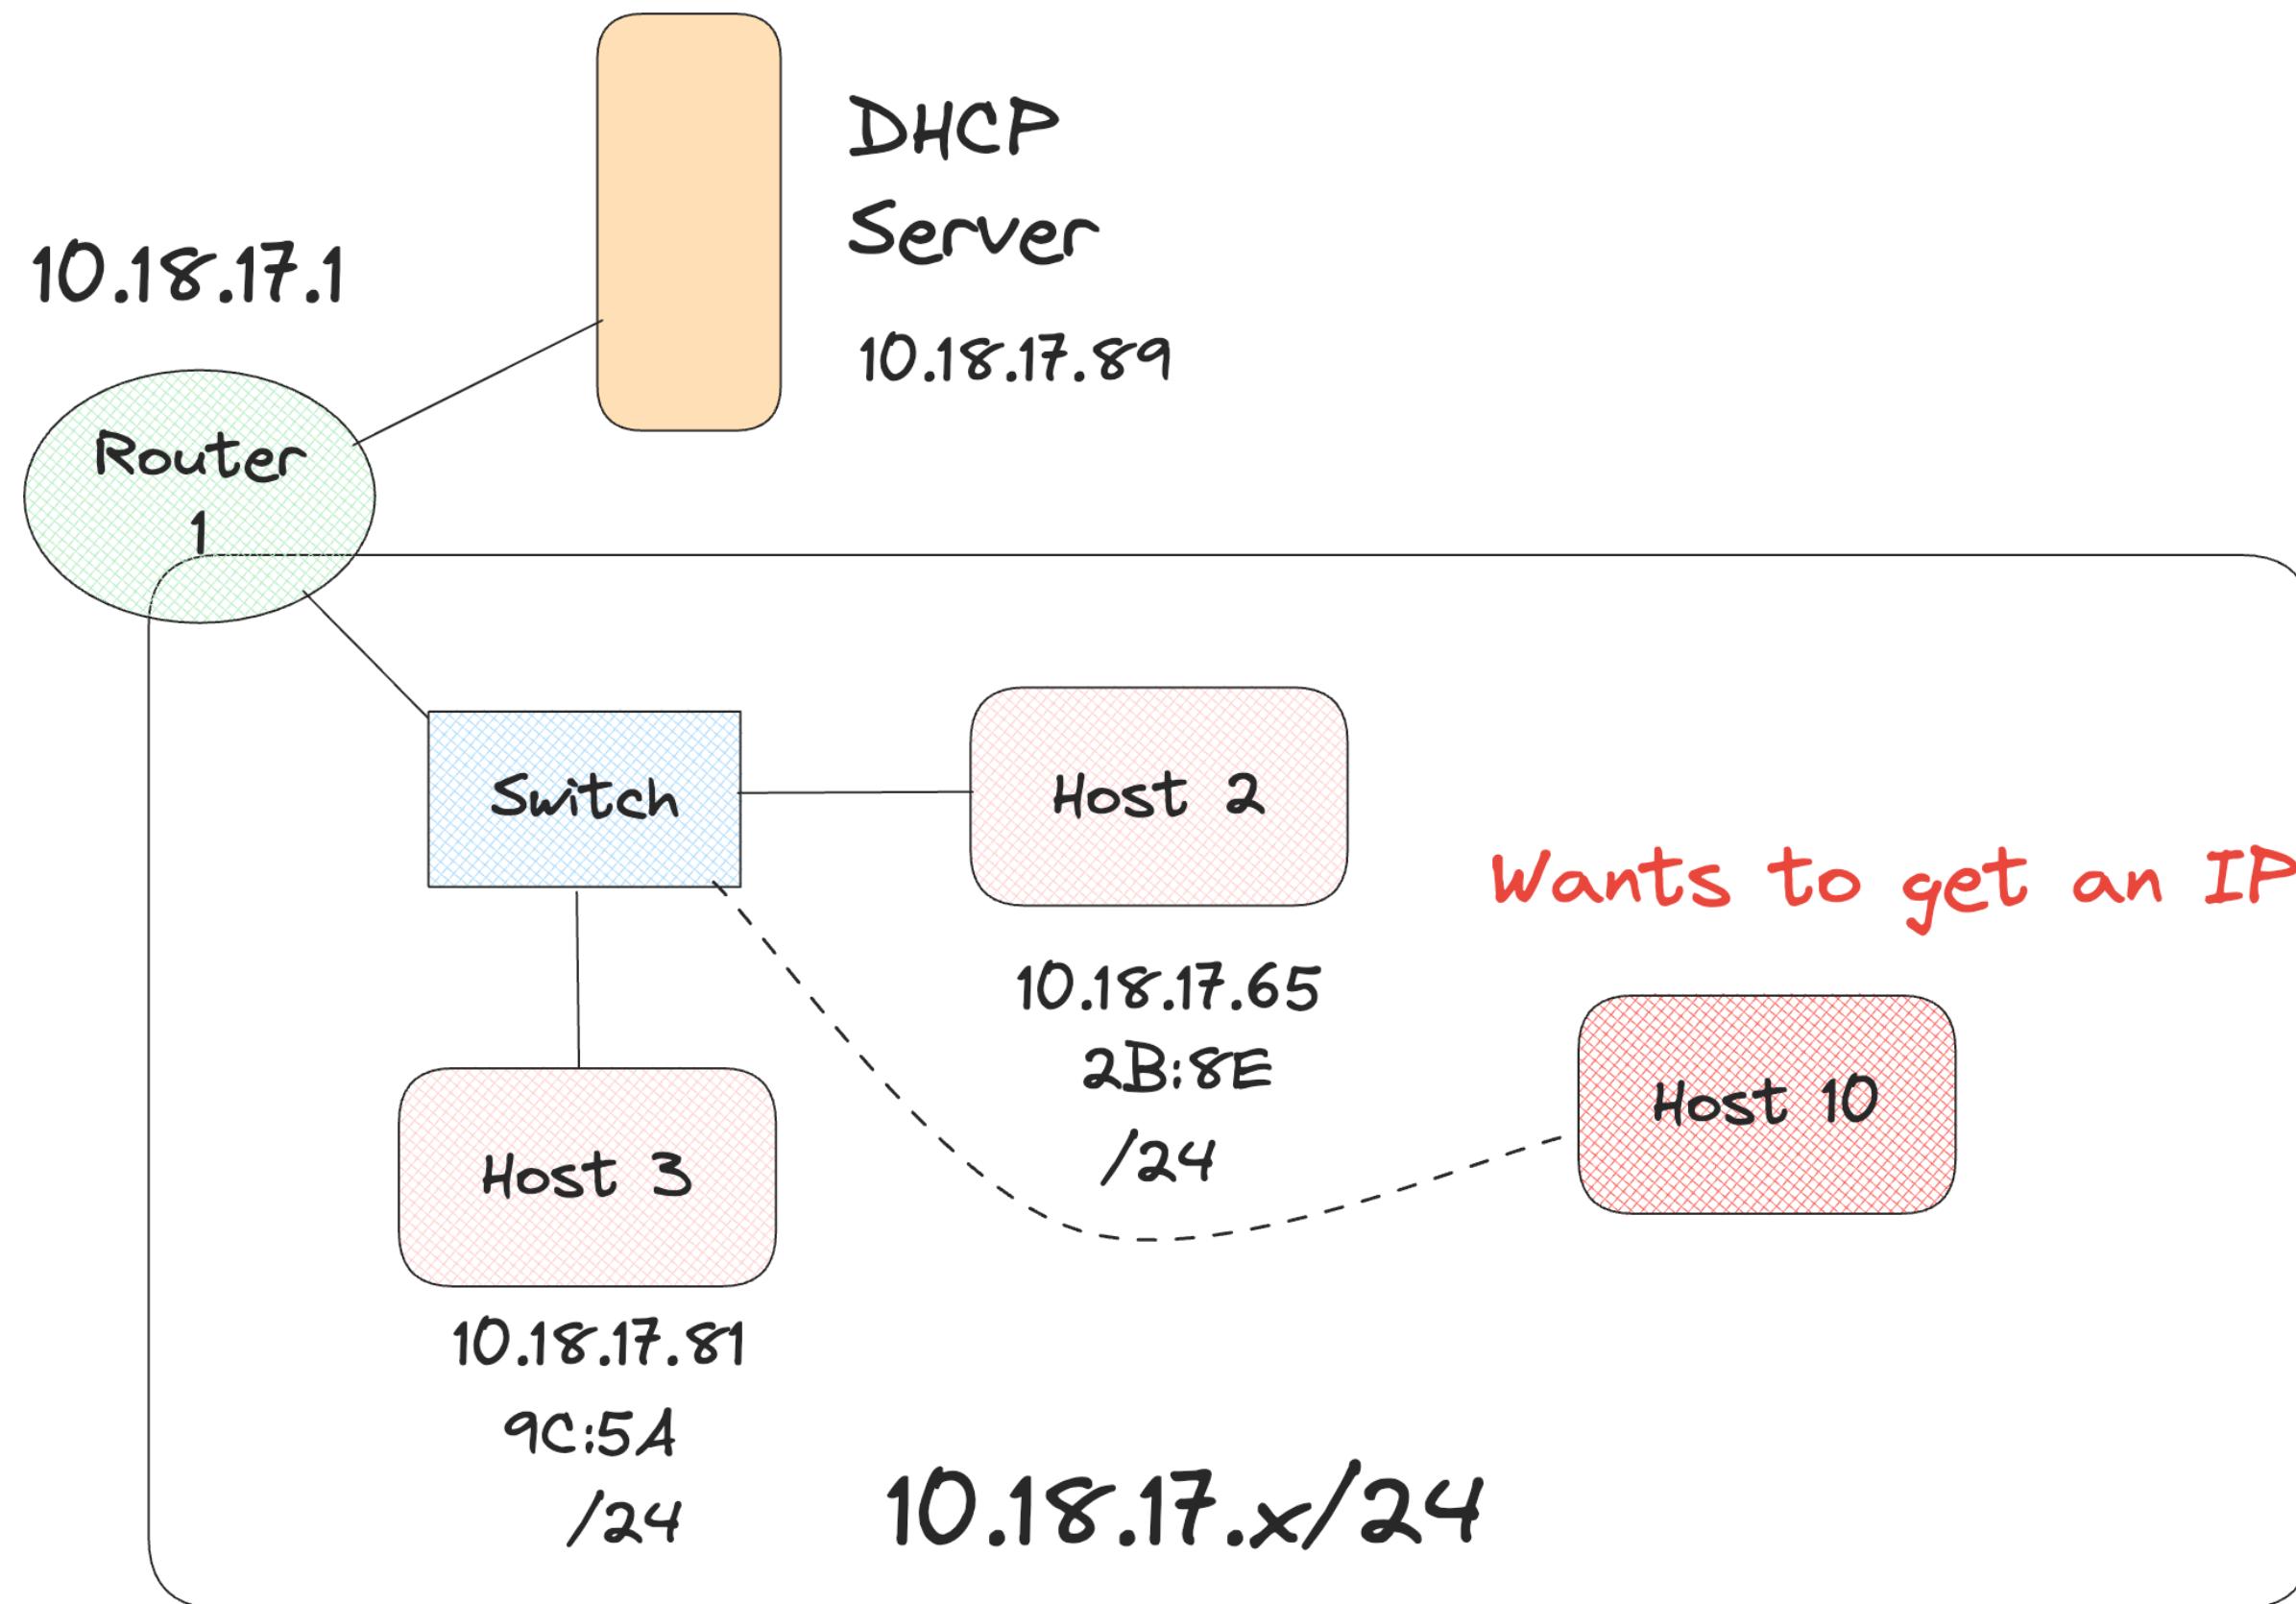
Network Layer (L3)

End-to-end Communication

- Manages routing through different routes in a large network
- Uses an addressing scheme - IP addressing
 - 32 bits represented as 4 octets (IPv4)
- Performs functionalities such as Logical addressing (IP), Path selection and packet forwarding
- L3 technologies: routers, even hosts are L3, L3 switches



Getting IP Address - DHCP



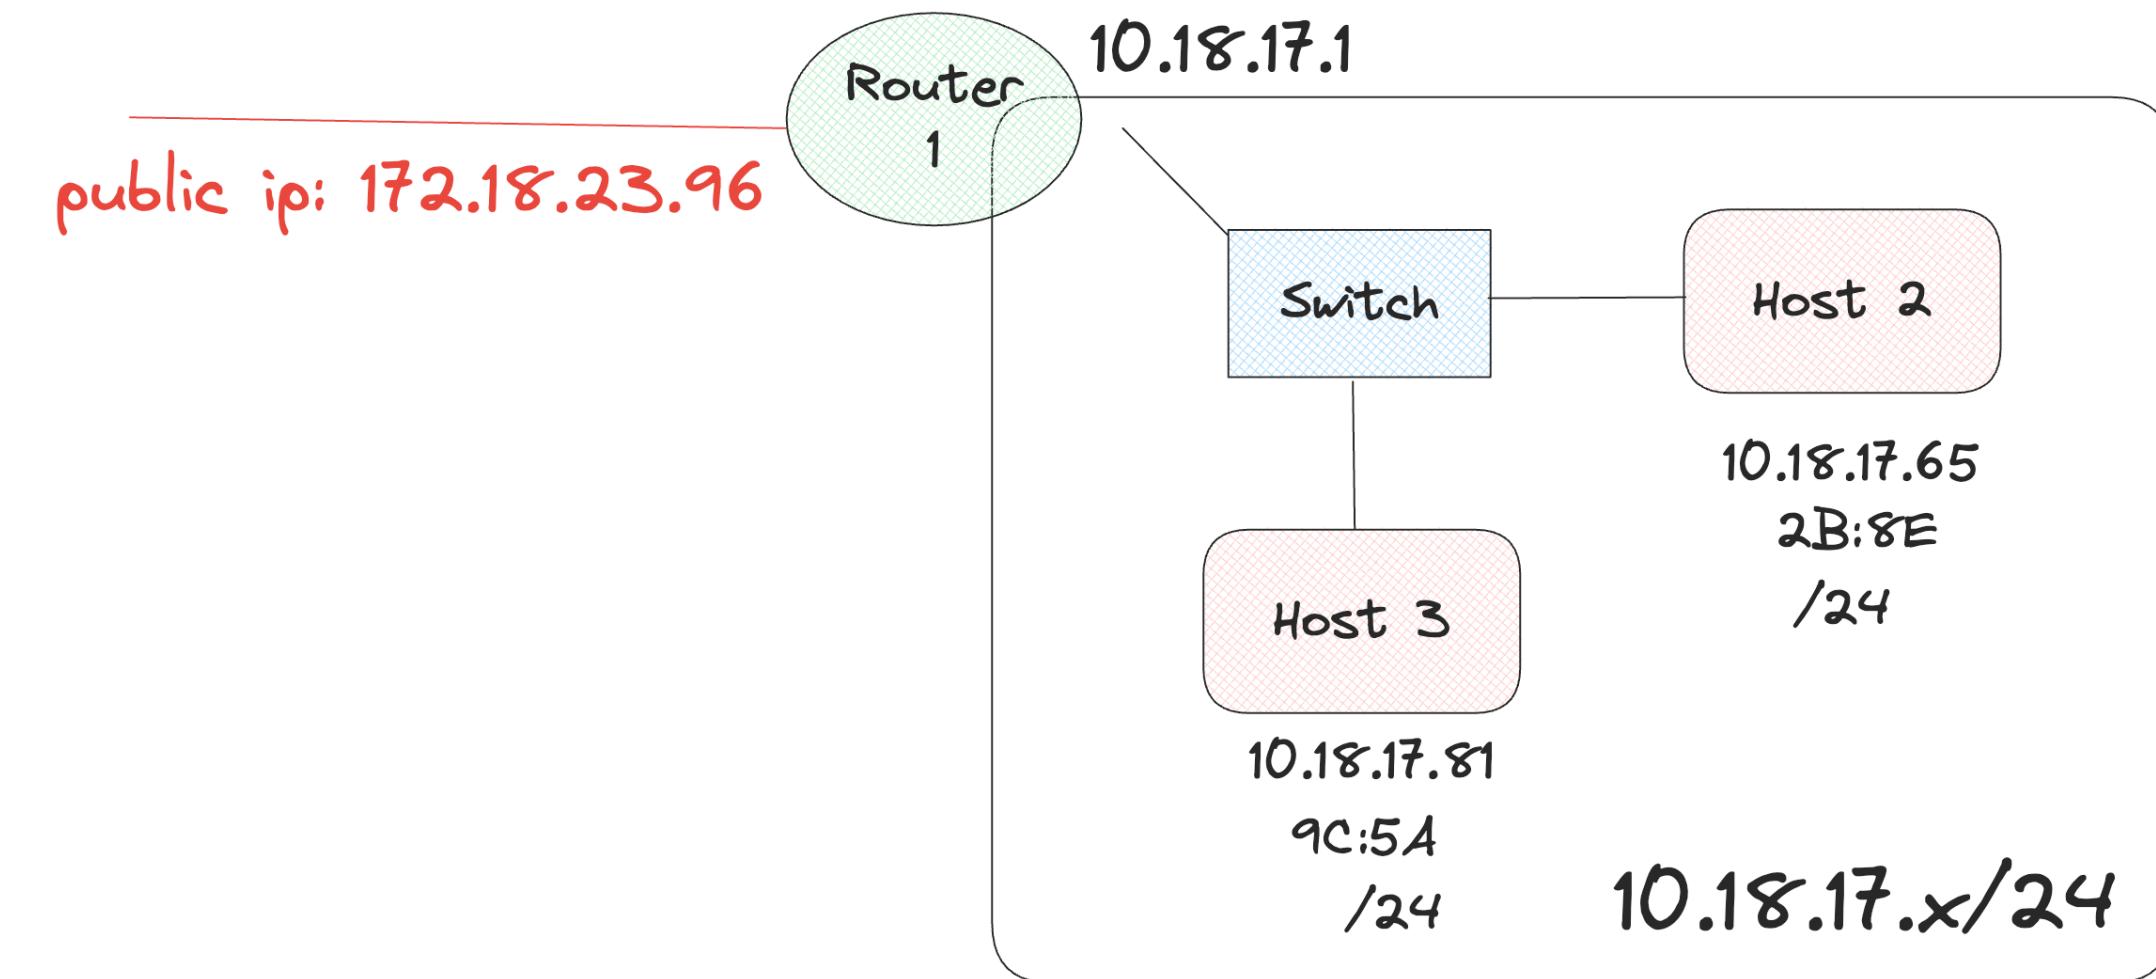
- Host 10 is the client here
- It sends out a broadcast DHCP request to every node in the network to get DHCP server
- Every device in the network will get the request
- DHCP runs over UDP
- Client uses port 68 and server port (listens on port 67)



Network Address Translation (NAT)

NAT Translation Table

WAN side address	LAN side address
172.18.23.96 5501	10.18.17.81 3801
....



- All devices in the network share just one IPV4 address as far as the outside world is concerned
- NAT allows a router (similar device) to translate private IP addresses to its own public IP address
- When devices from network wants to communicate with outside network:
 - NAT modifies the source IP to make it appear that communication is from the larger public IP
 - A translation table is used for managing the translations
- **Multiple types:** Static NAT, Dynamic NAT, Port Address Translation or NAT Overload

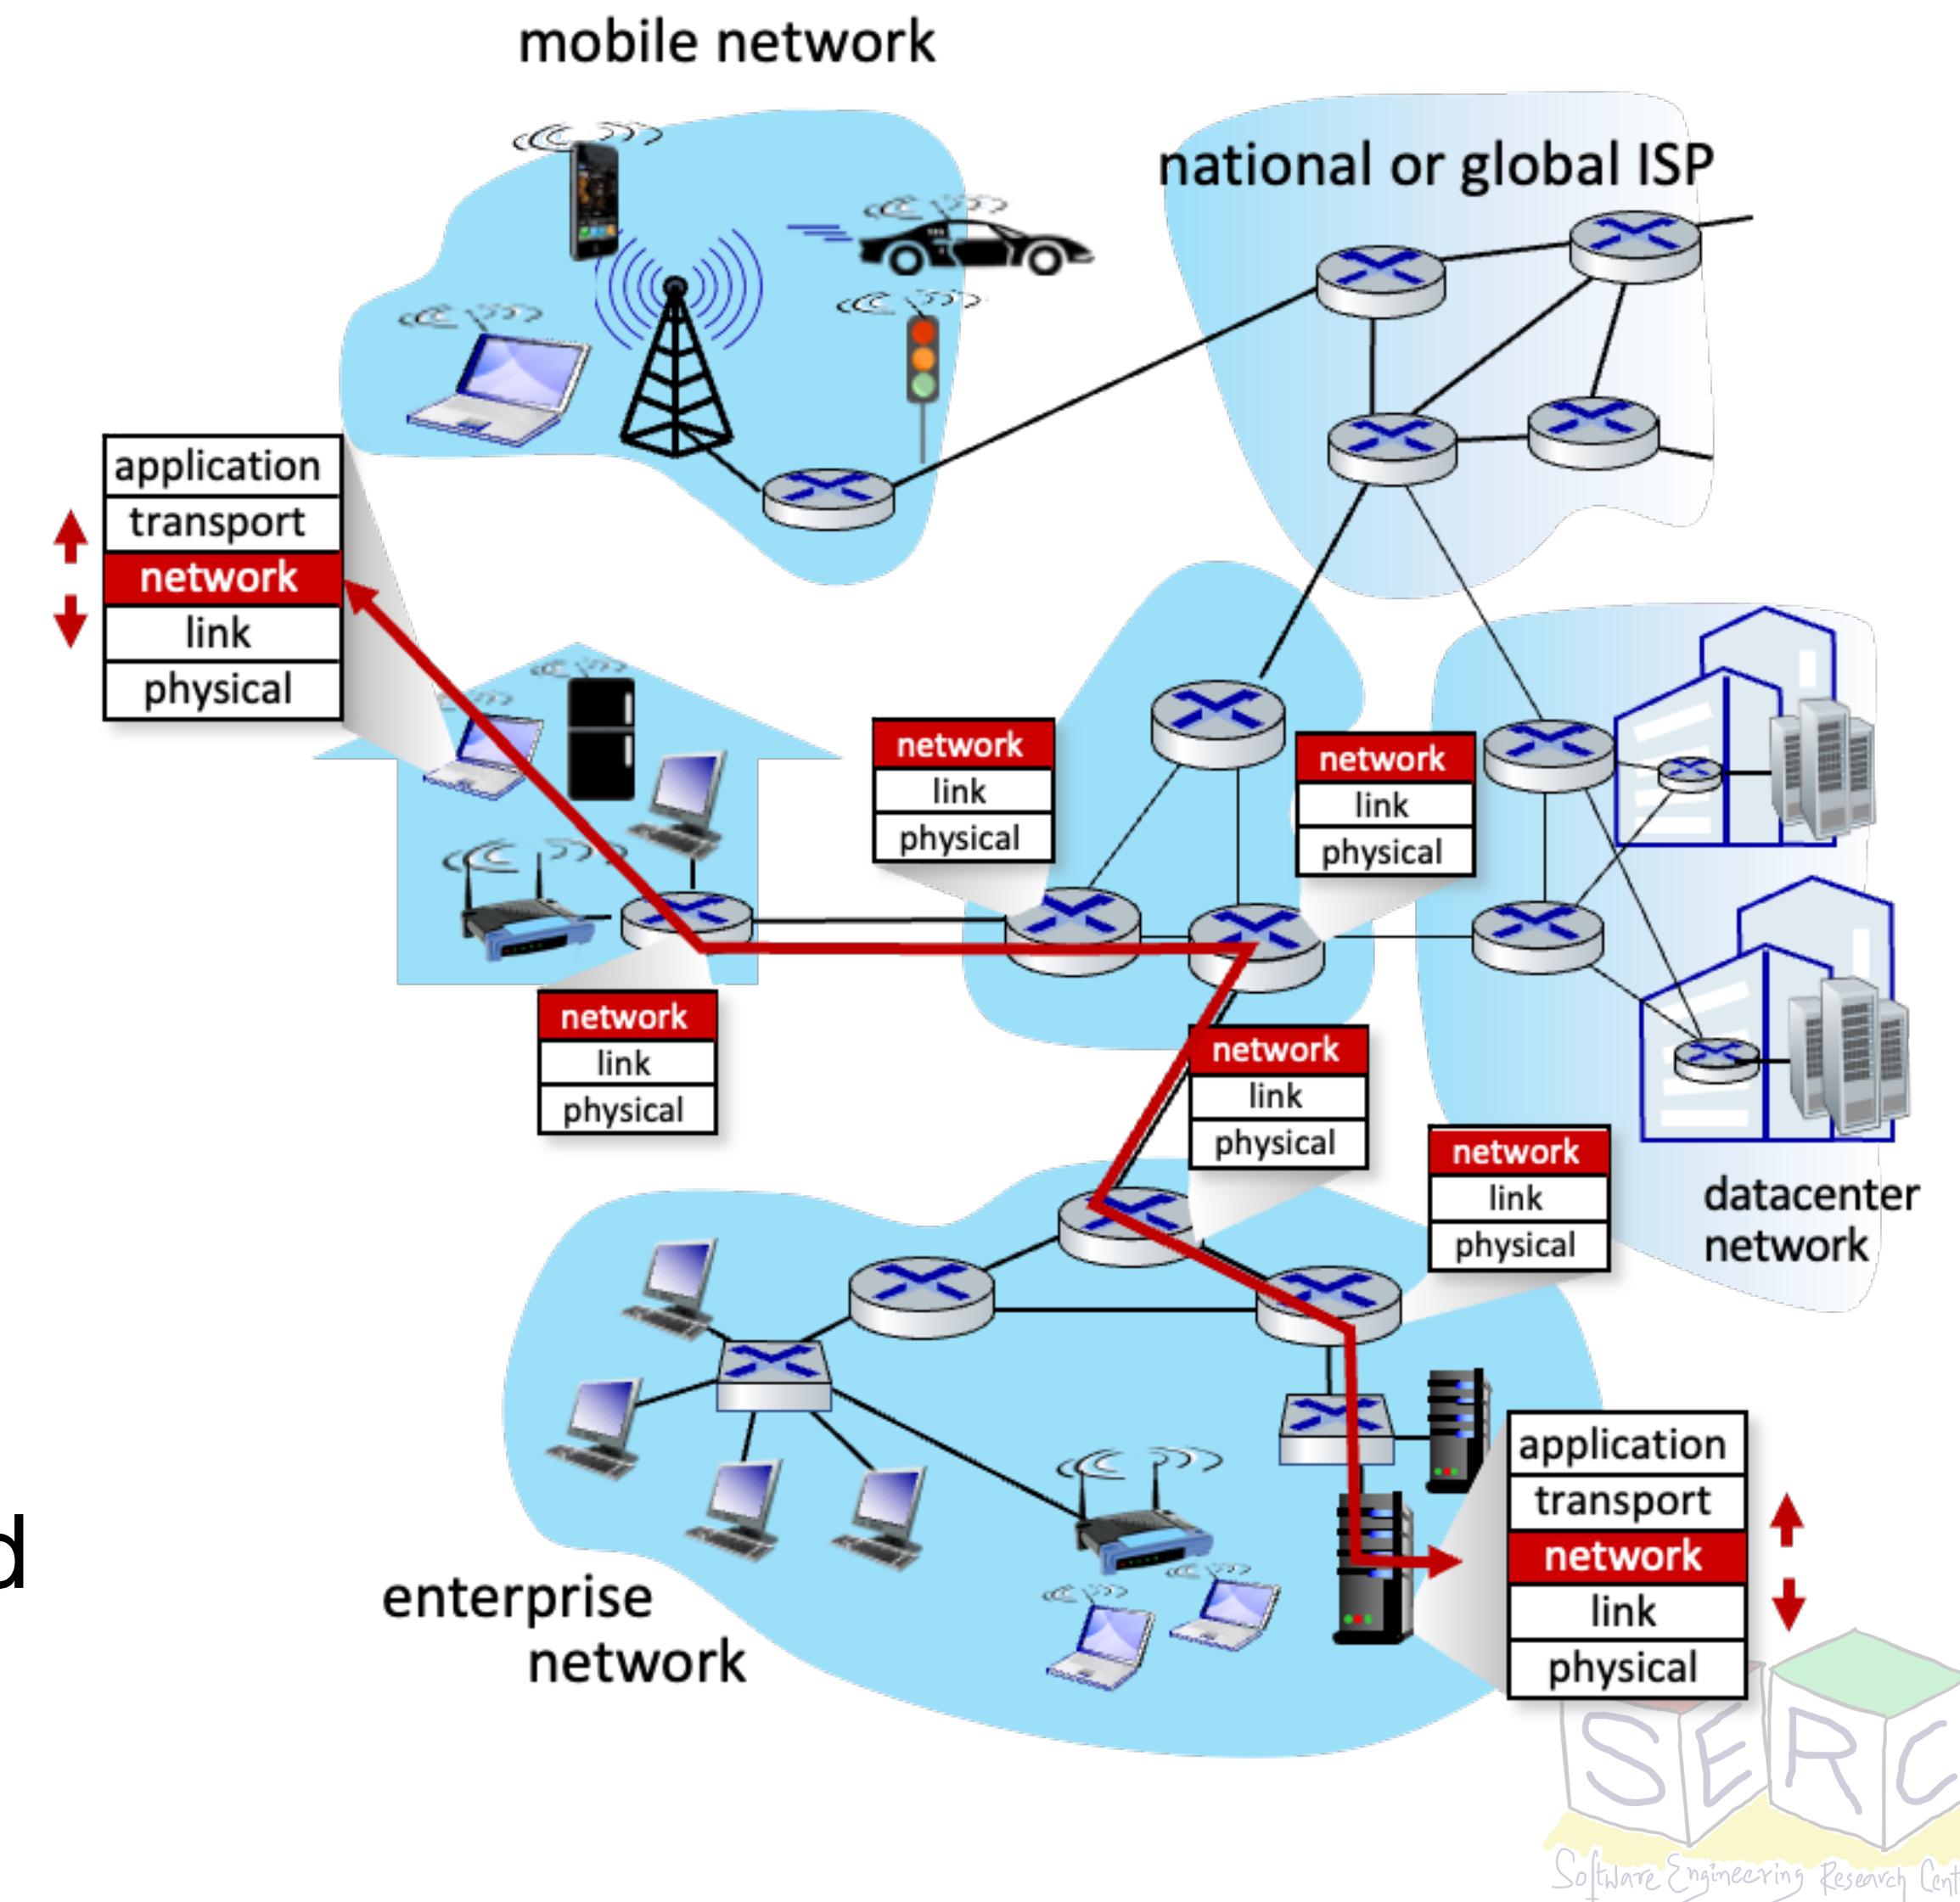
Network Layer - Functionalities

Addressing

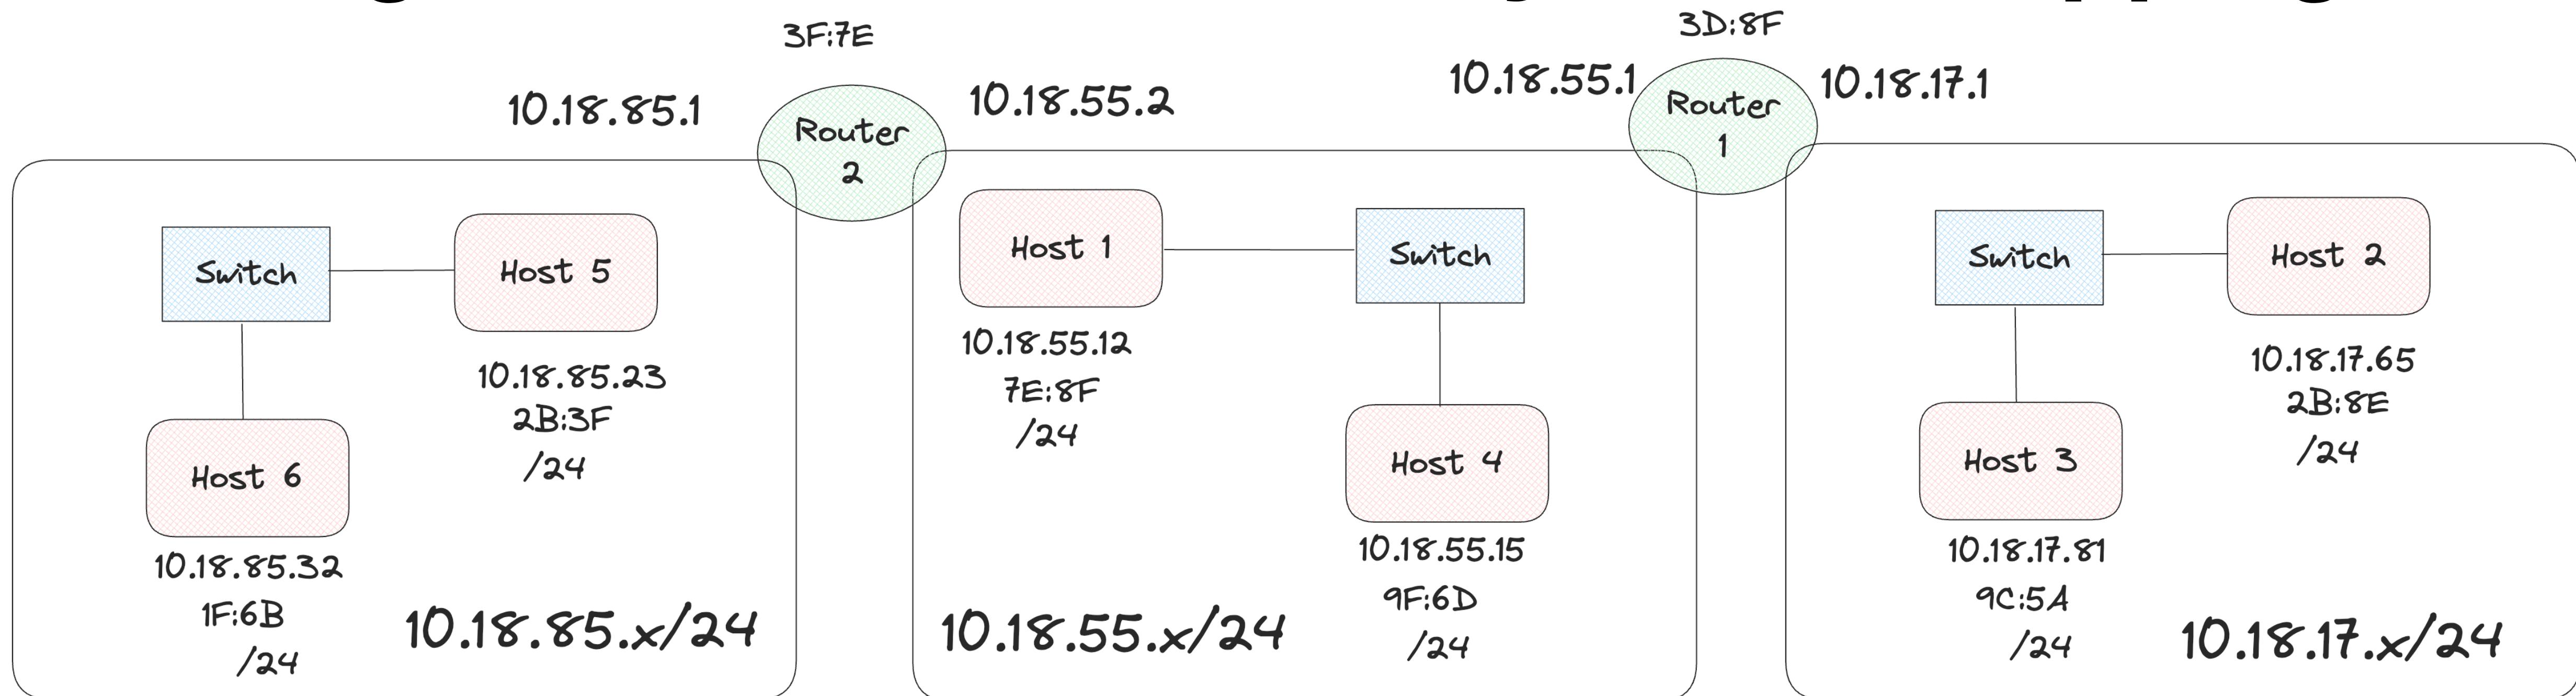
- Devices in network are assigned logical address for unique identification - IP
- Network layer uses IP to forward packets to the intended destinations

Route Determination

- Identifies best path for packets to reach to destination
- This process is dynamic and changes based on network conditions



Routing Tables - Static and Dynamic Mappings



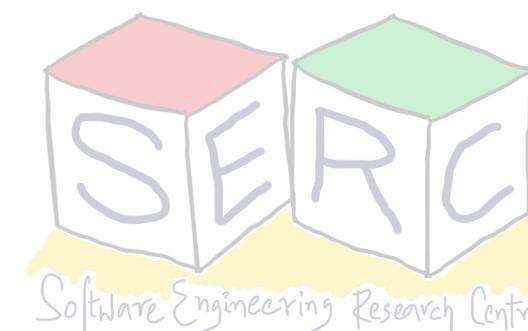
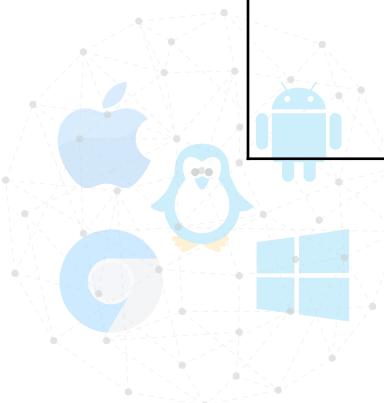
Type	Destination	Interface
DC	10.18.85.x/24	Left
DC	10.18.55.x/24	Right
Static	10.18.17.x/24	10.18.55.1

Router 2 routing table

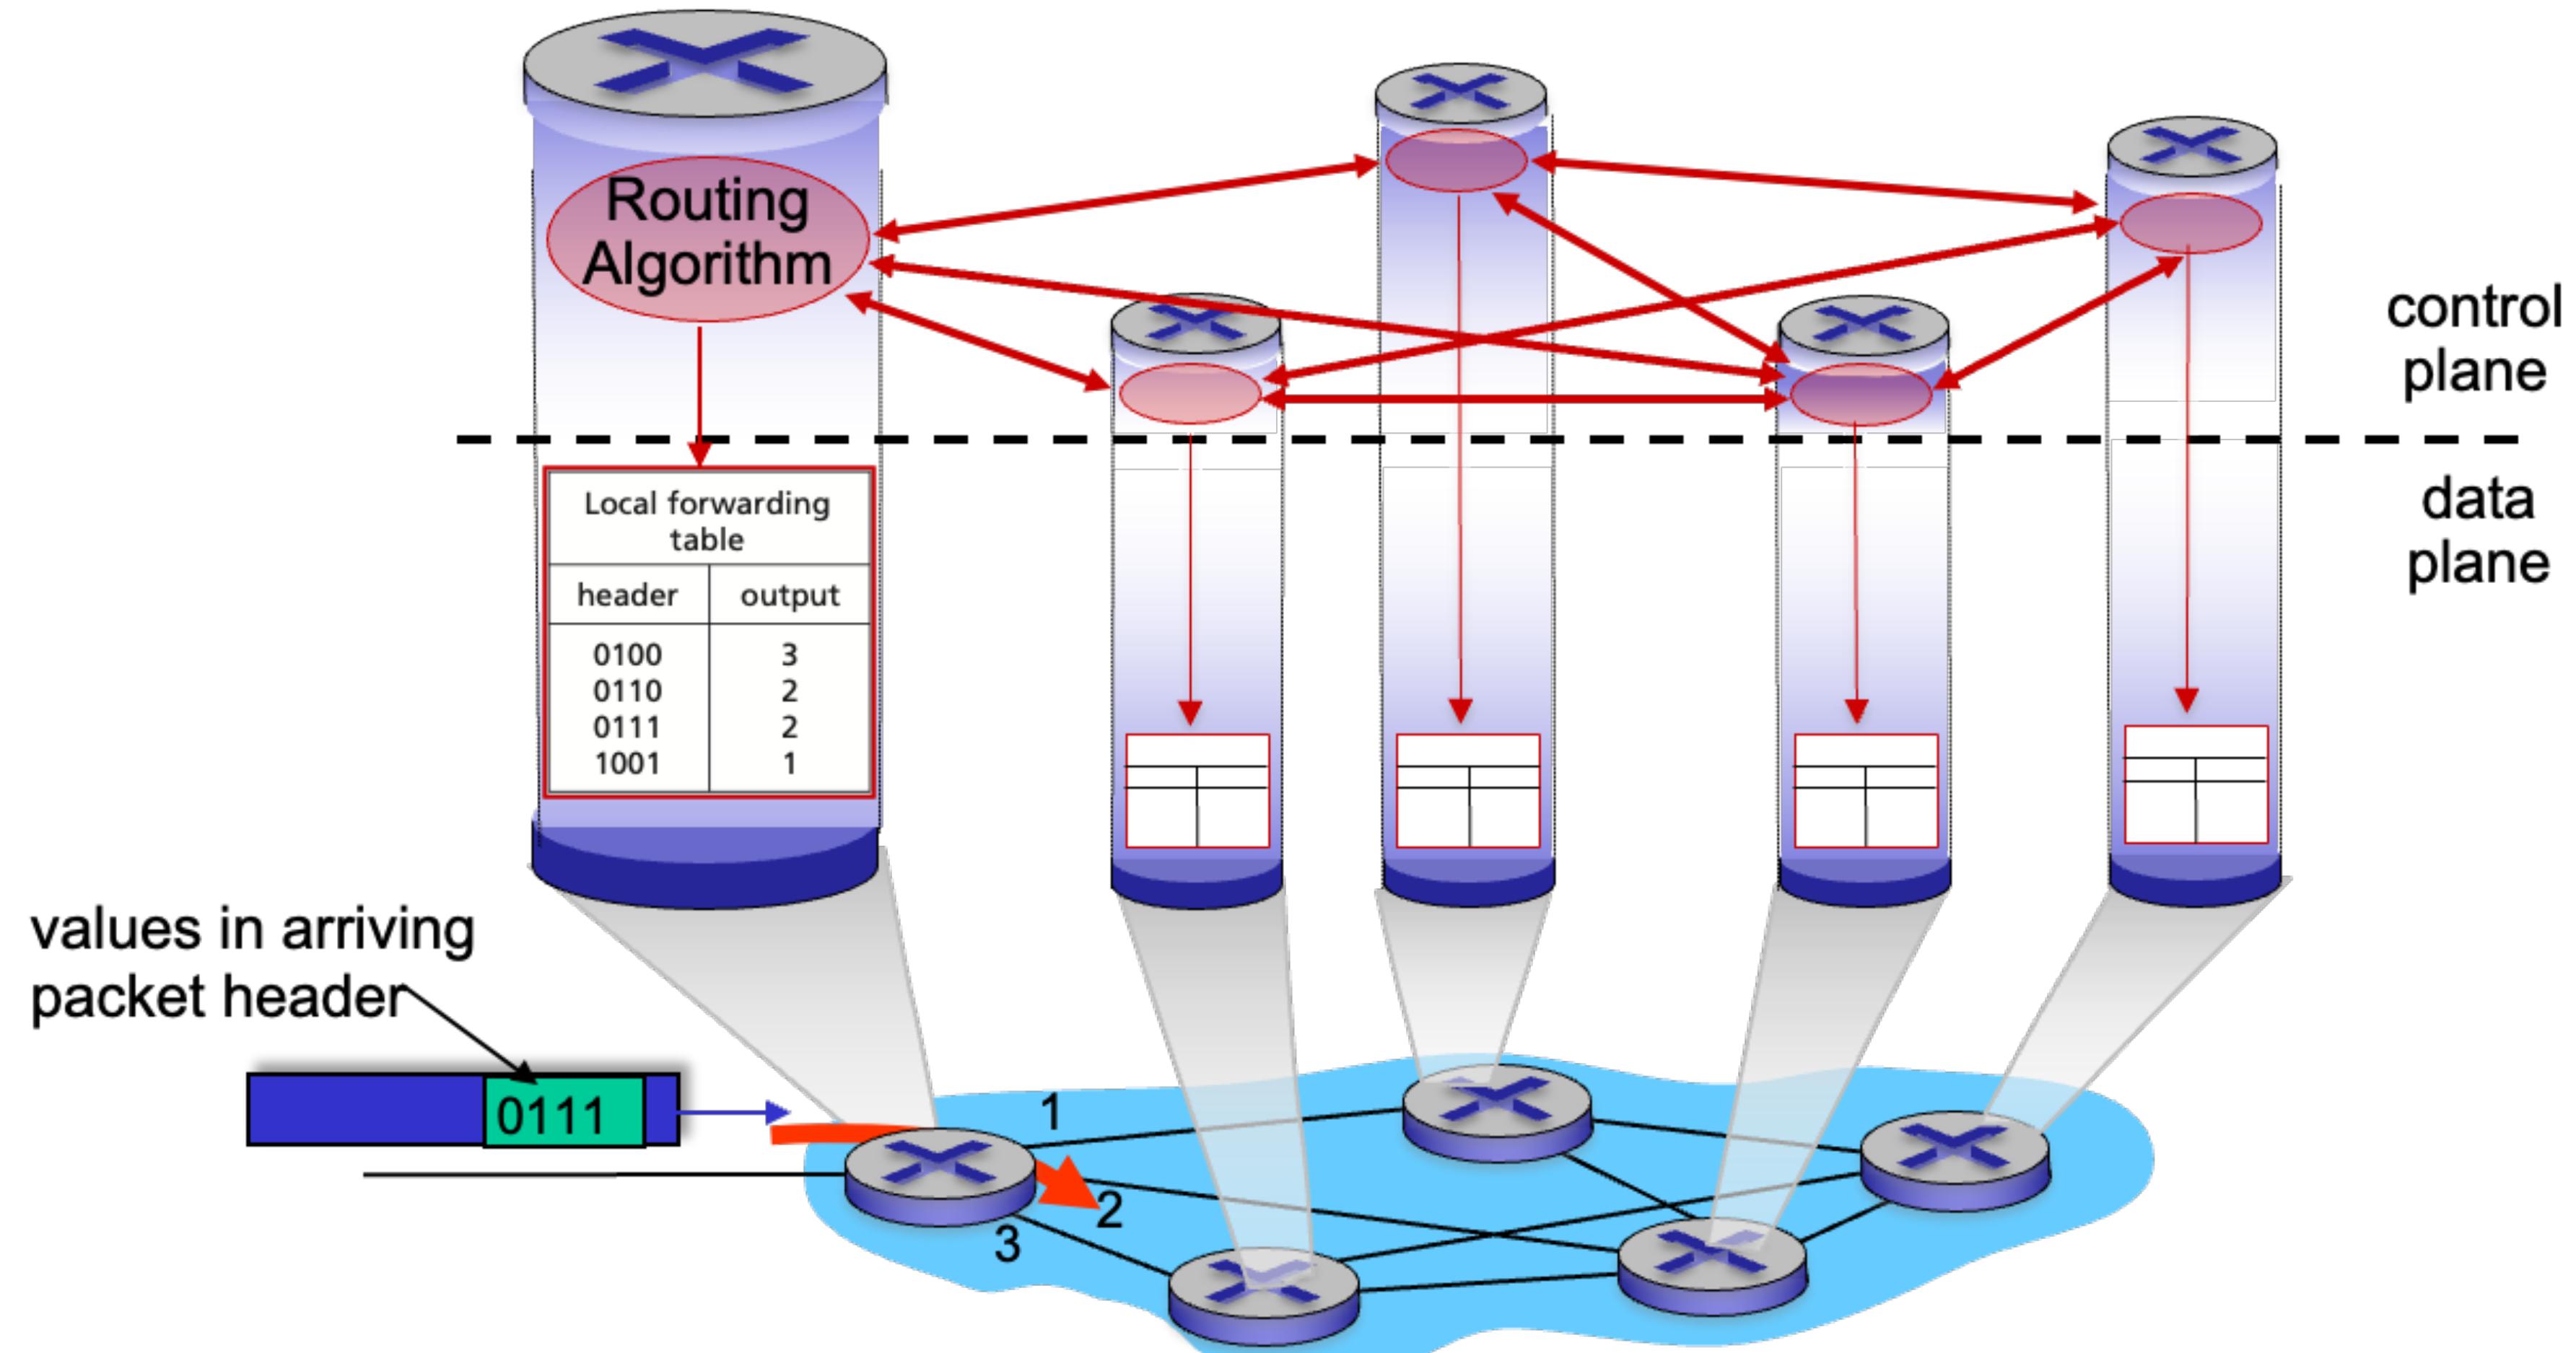
Type	Destination	Interface
DC	10.18.55.x/24	Left
DC	10.18.17.x/24	Right
Static	10.18.85.x/24	10.18.55.2

Router 1 routing table

Routers can also
Learn about the address



Traditional Control Plane Approach

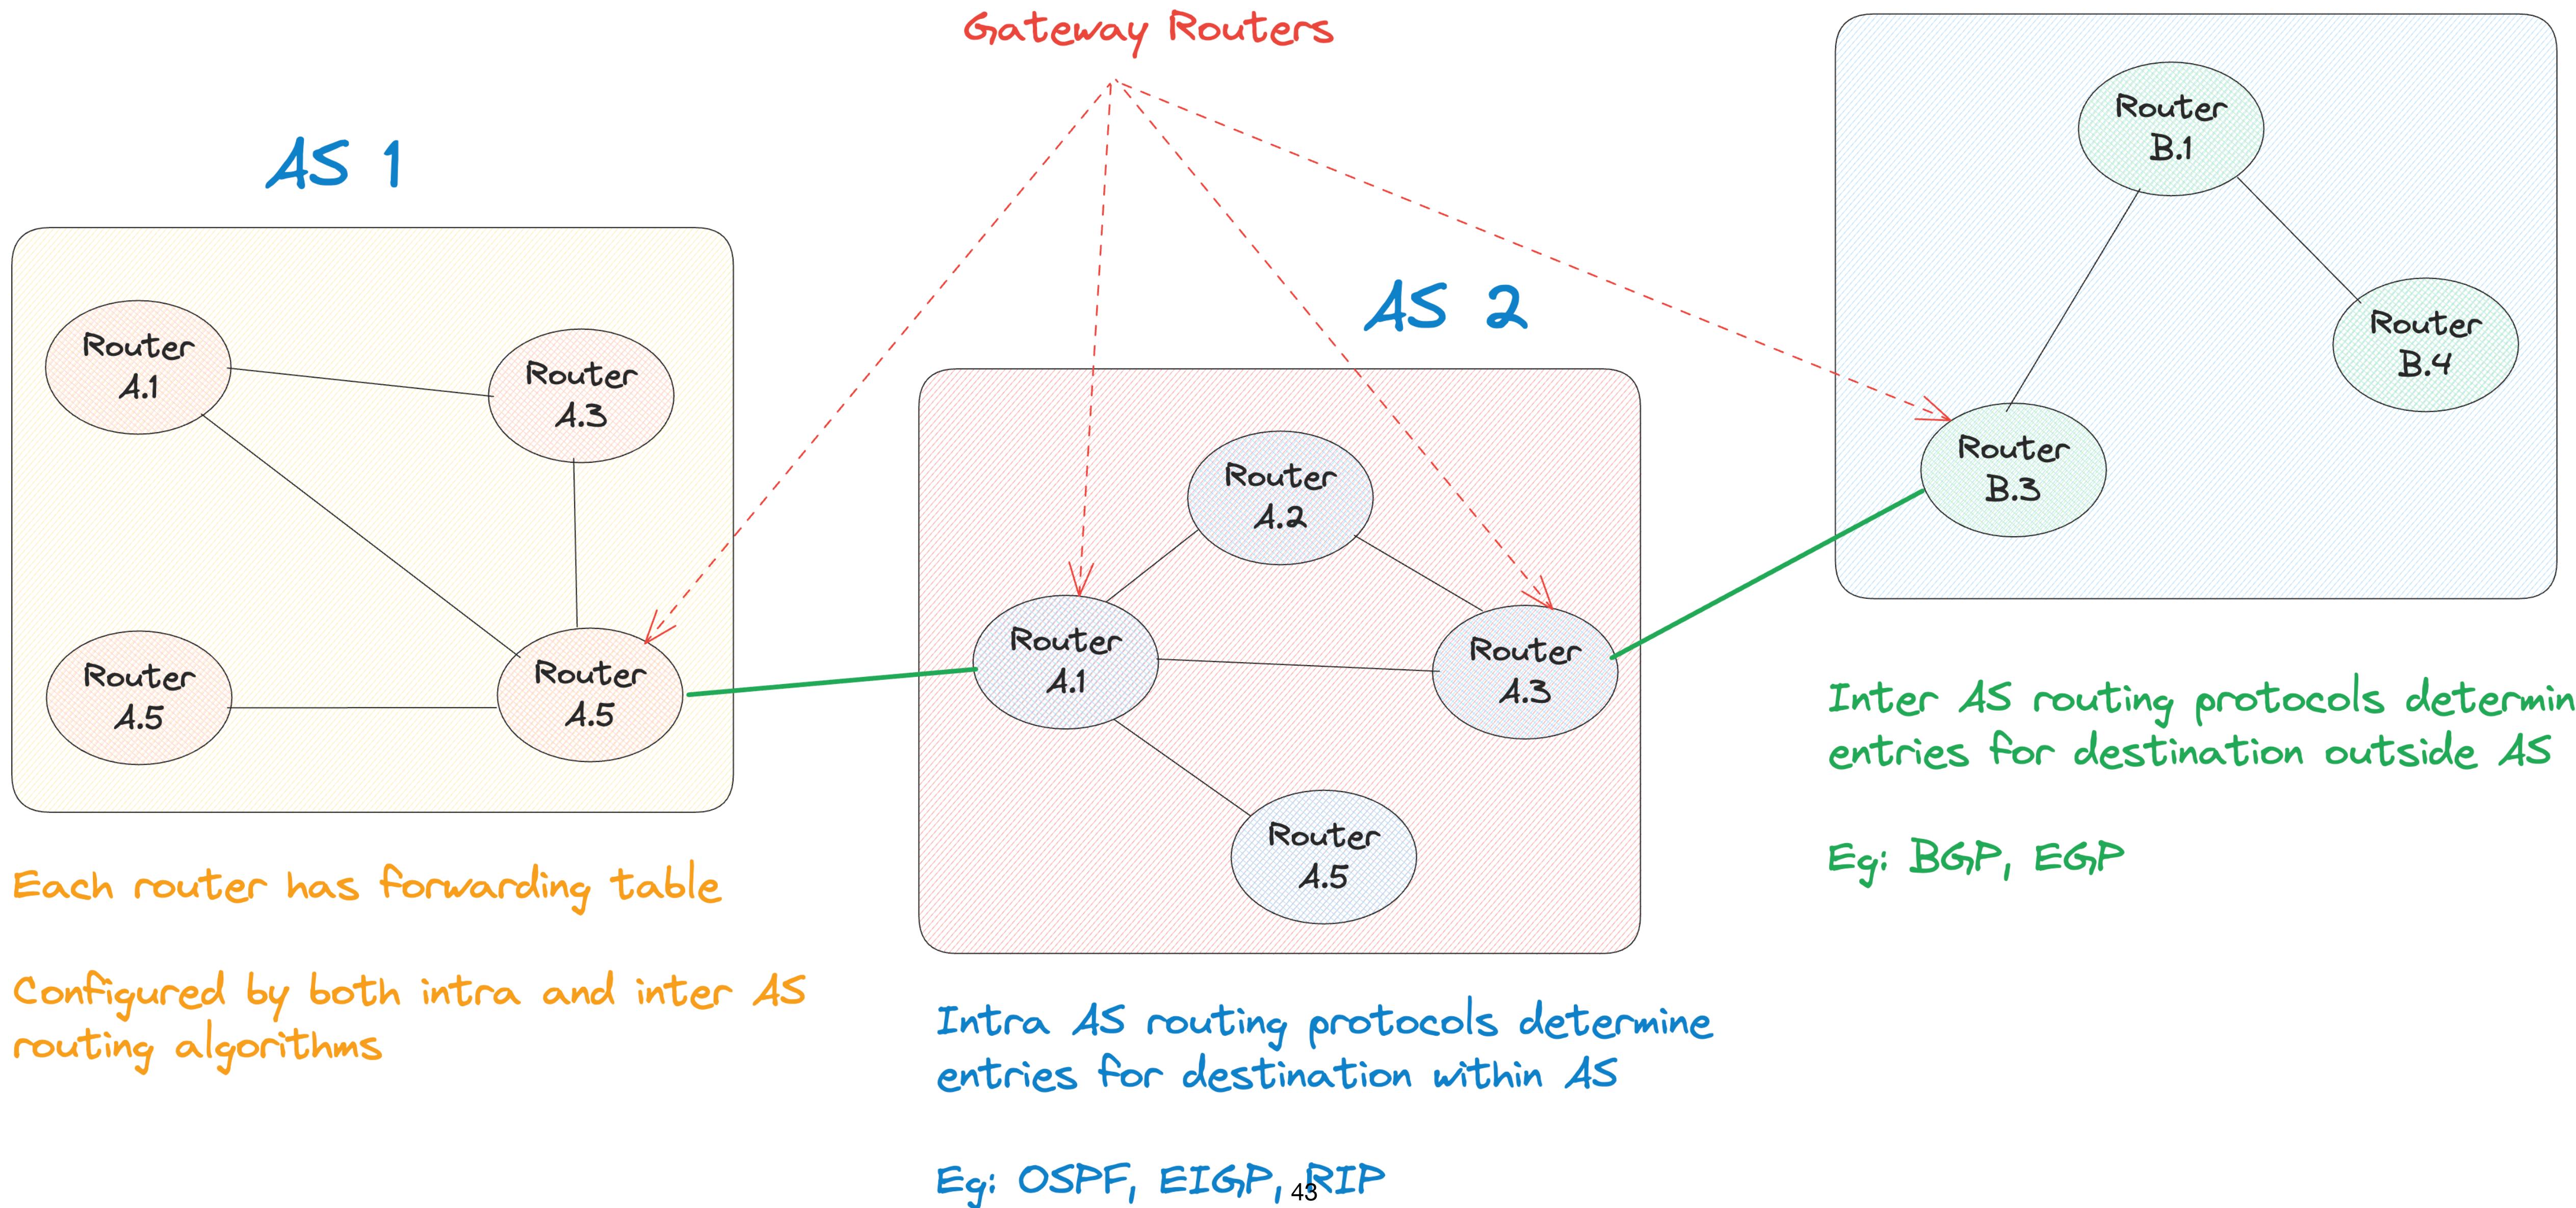


Internet approach to scalable routing

- Aggregate routers into regions known as “**Autonomous Systems**” (**AS**) a.k.a “domains”
 - Total of around 70,000 AS’s have been assigned not all are active
- There are mechanisms for handling routing within the domain and across AS
- **Intra-AS or Intra-domain**
 - All routers in AS must run the same intra-domain protocol
 - There is a **gateway router** at the edge of each AS which connects with router in another AS
- **Inter-AS or Inter-domain**
 - Routing among AS’s
 - Gateways perform inter-domain as well as intra-domain within their network

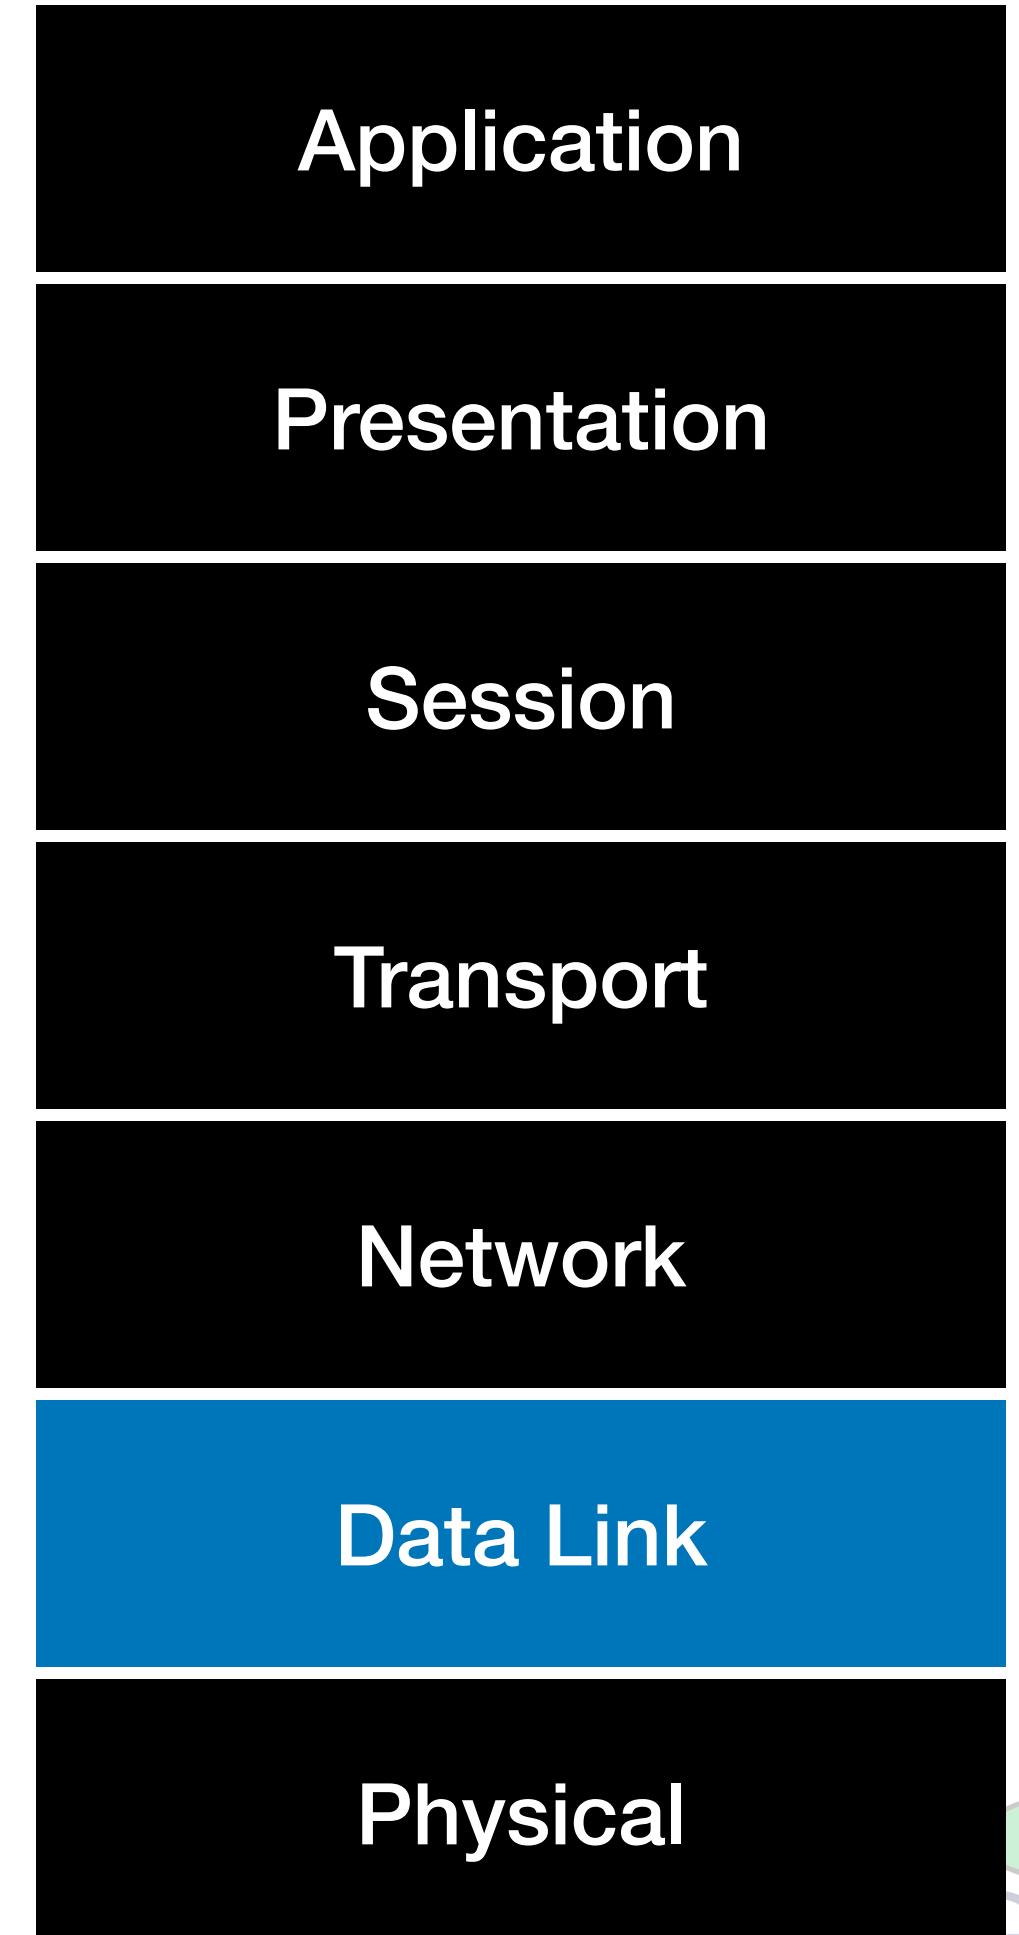


High Level Overview

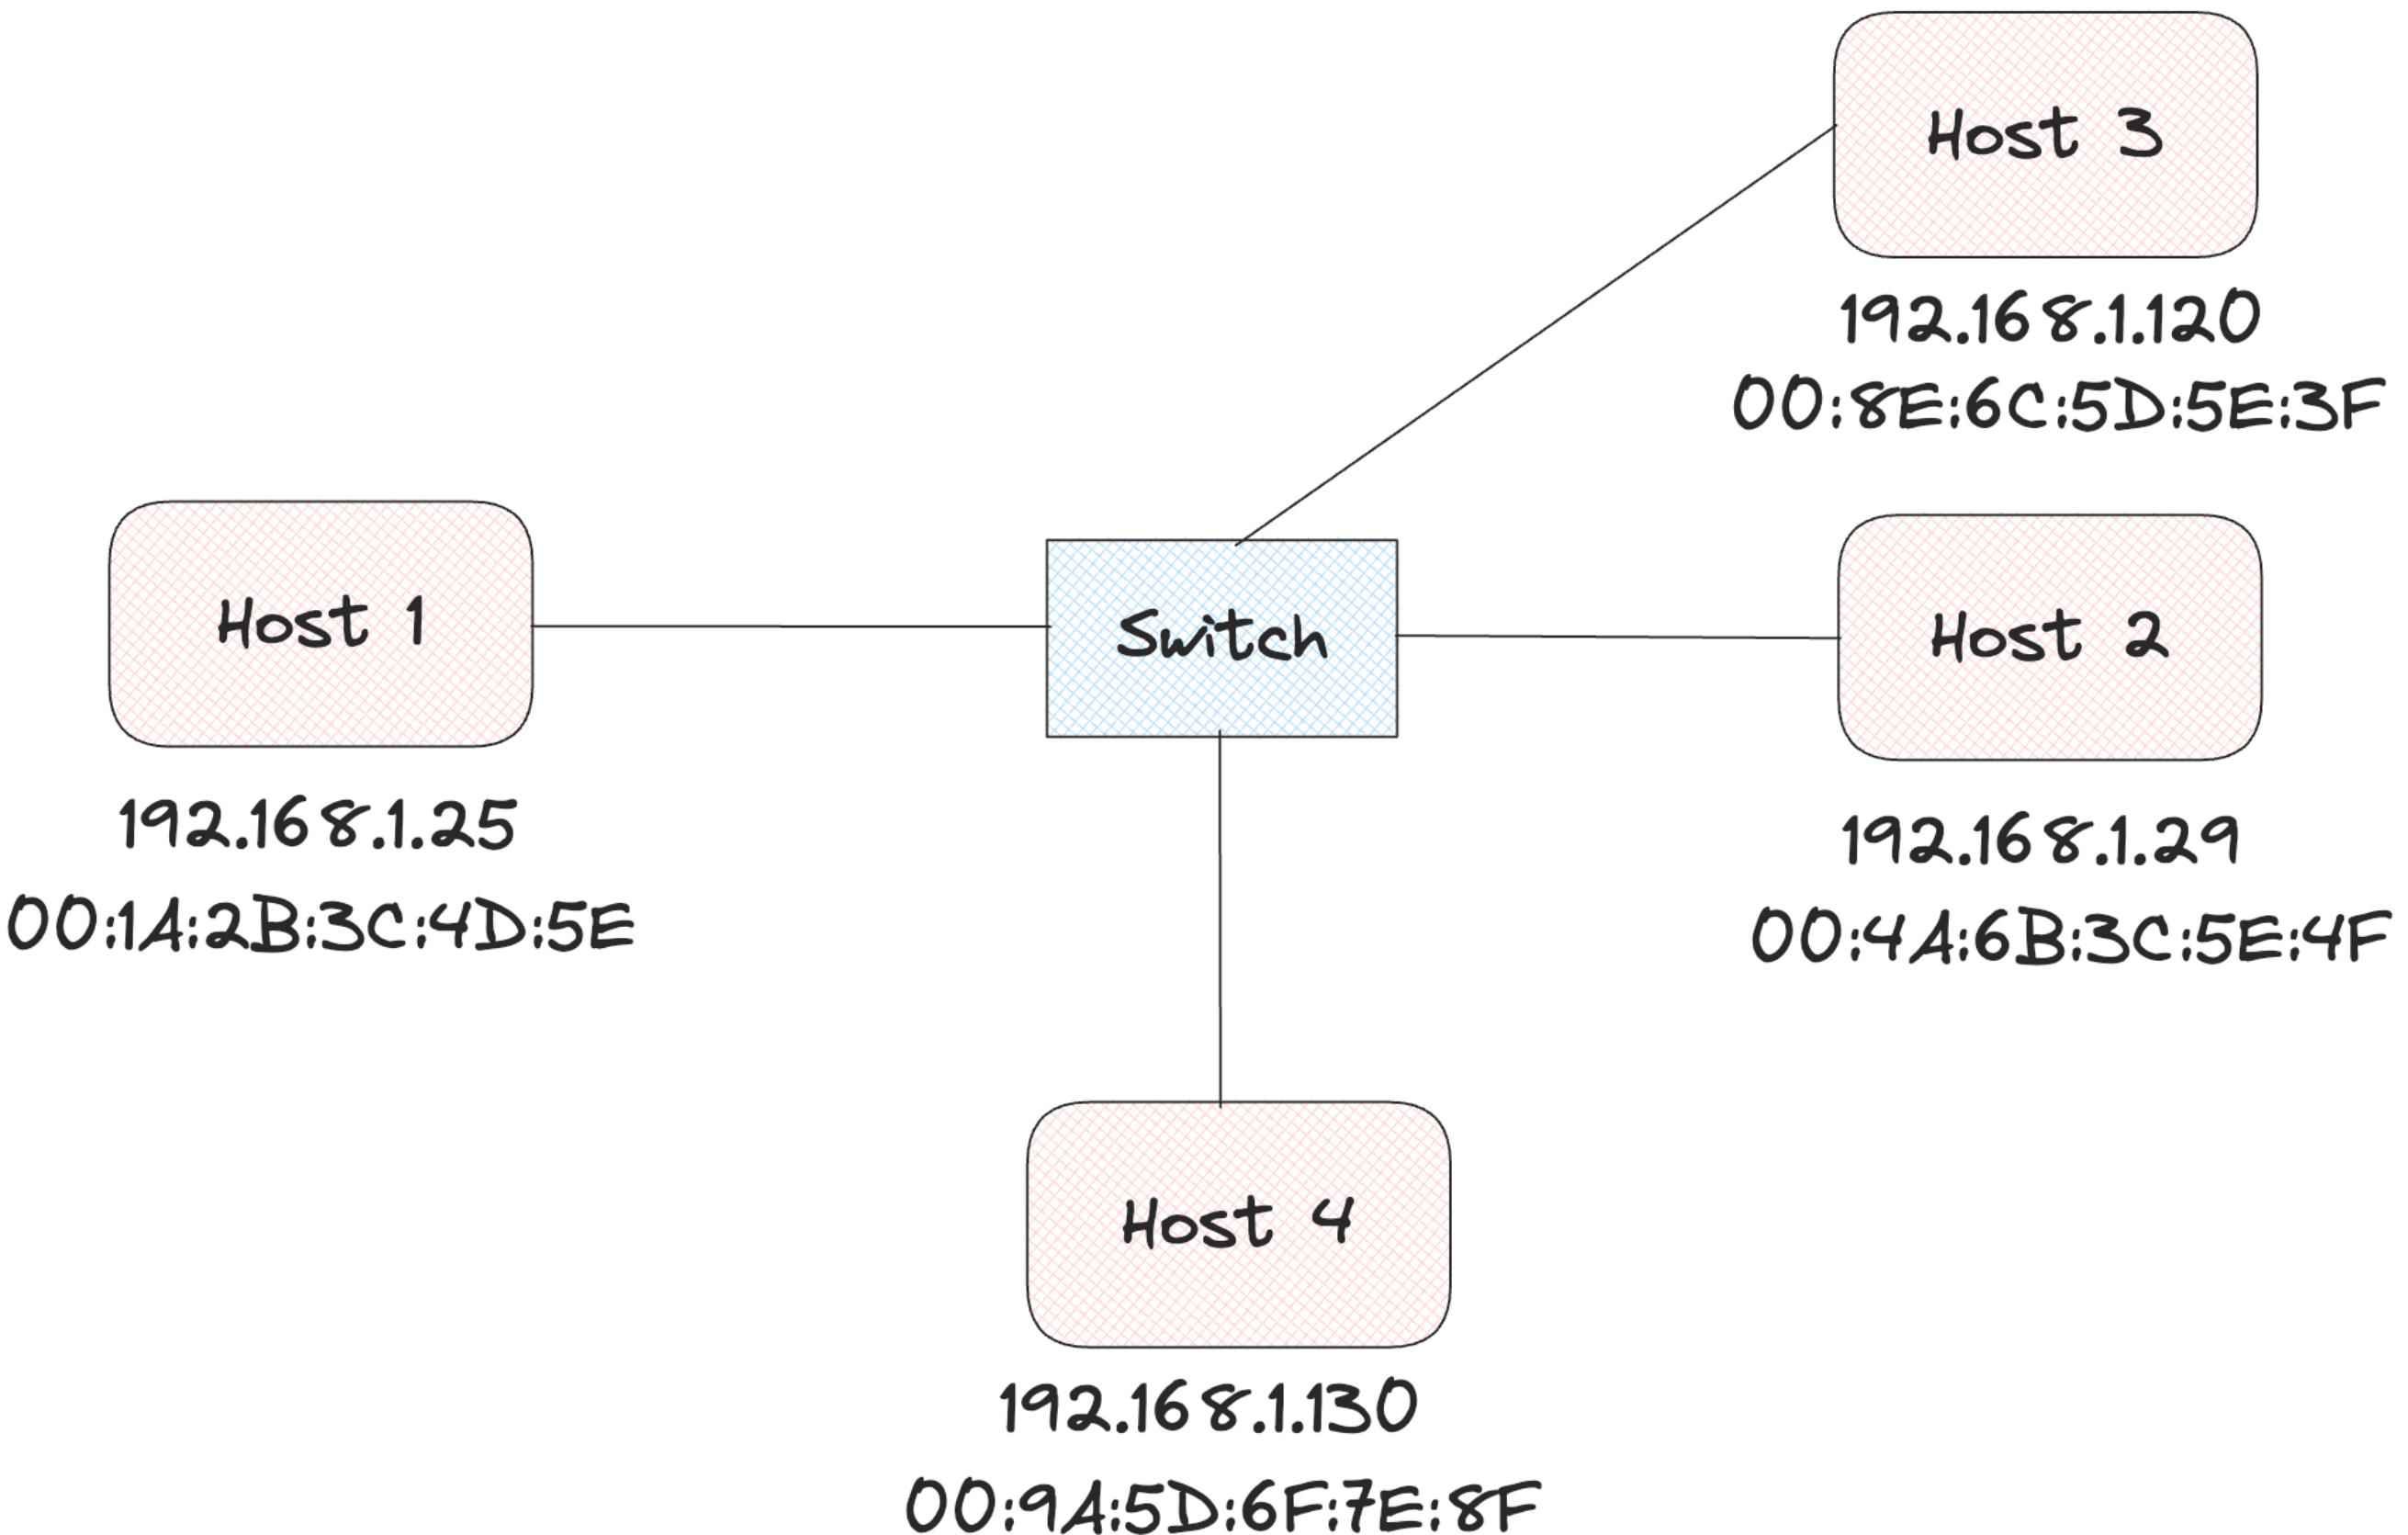


Data Link Layer (L2)

- Responsibility of transferring datagram from one node to a **physically adjacent node** over a link (no intermediate L3 routers)
- Supports **hop-to-hop** communication
- Ensures reliable connection link between two directly connected nodes (flow control, error correction and detection, etc.)
- Supported by **Media Access Control (MAC)** addressing
- Addressing scheme: MAC addressing (48 bit address, 12 hex digits, 6 bytes)
 - Eg: **00:1A:2B:3C:4D:5E**
 - First three identify manufacturer (IEEE)
 - Next three are assigned by manufacturer and should be unique



ARP - Address Resolution Protocol



- Each IP node (router, host) on the LAN has a table - **ARP Table**
- IP/MAC address mappings for some LAN nodes
 - <ip address, MAC address, TTL>
 - TTL: Time to live, time after which the mapping will be forgotten (20 mins)

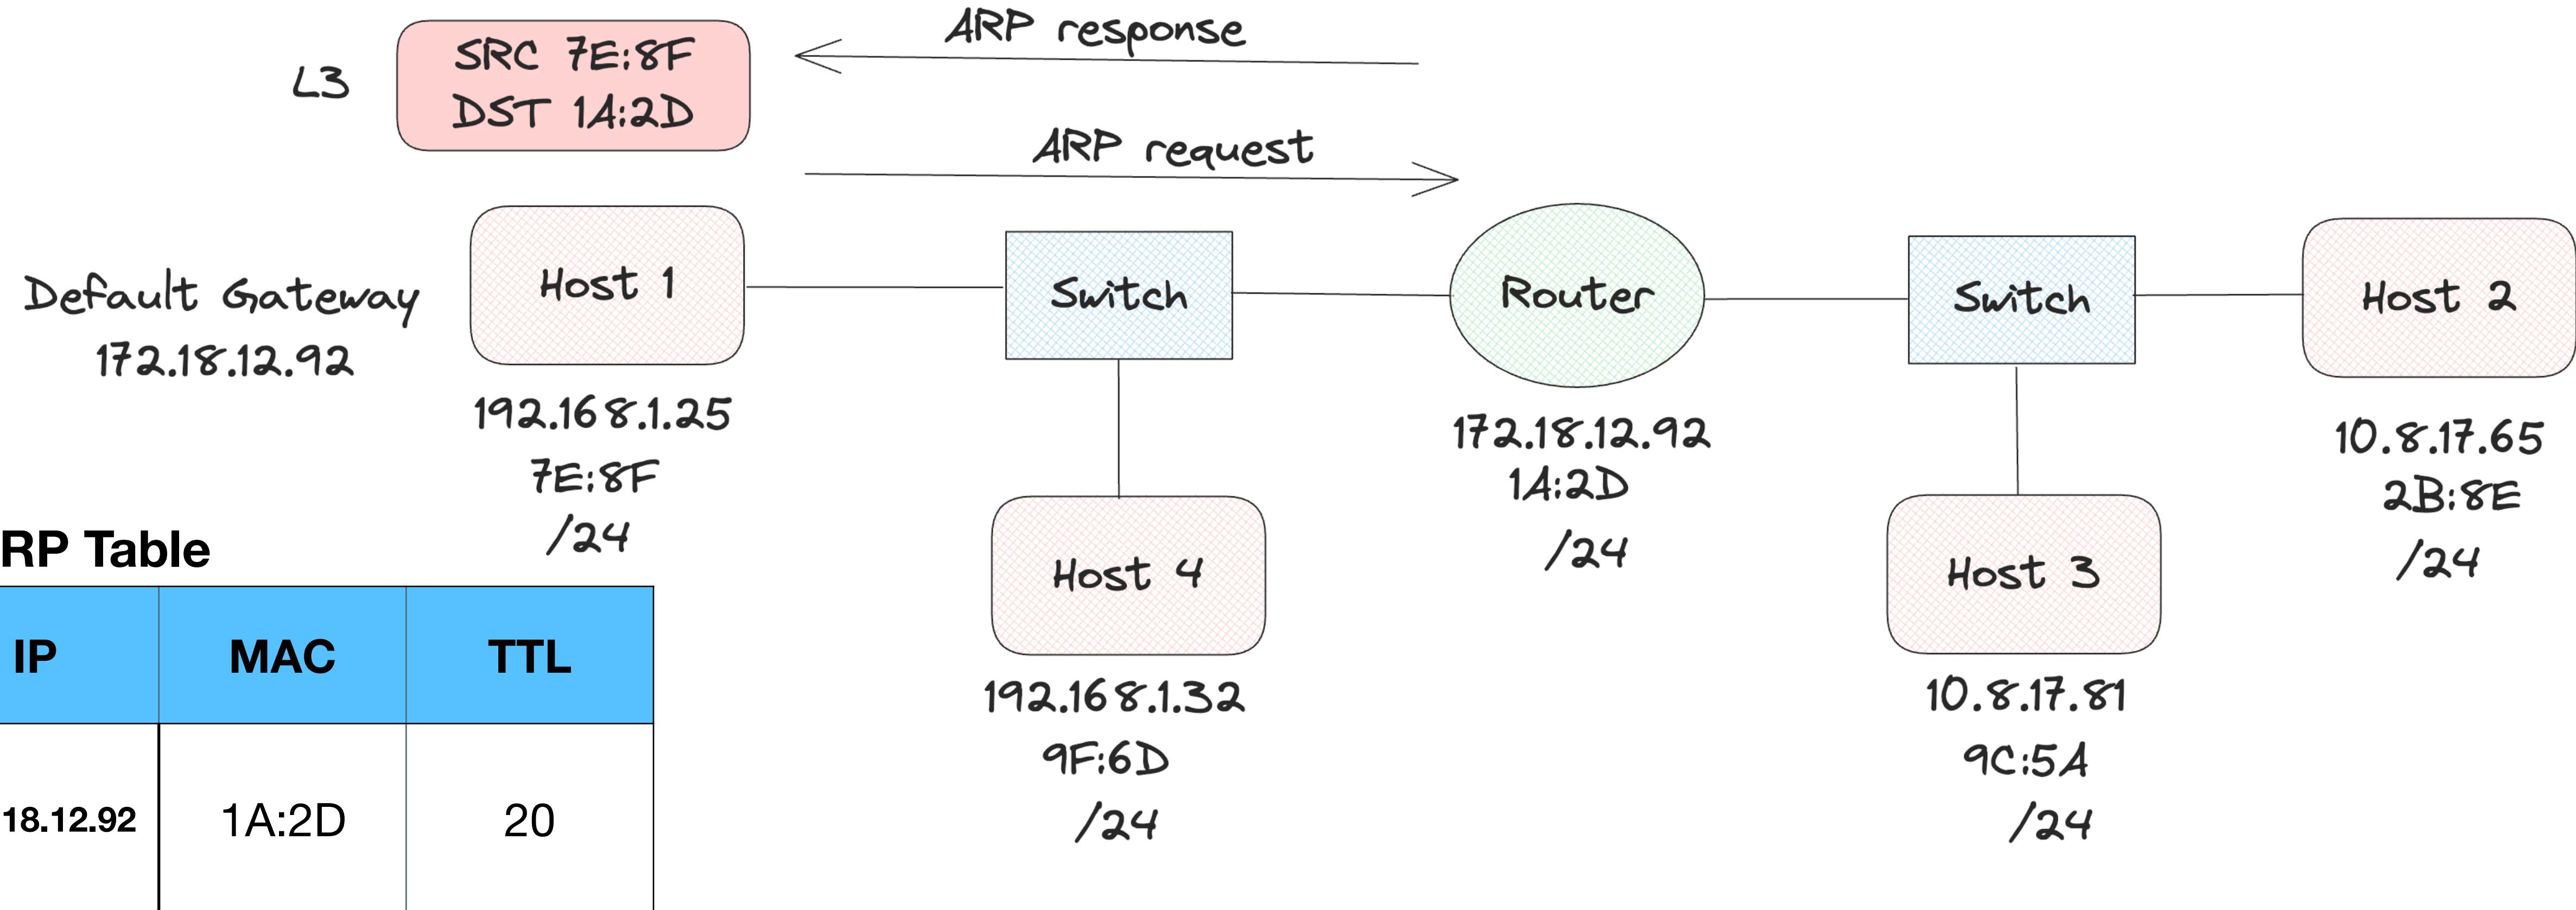


ARP Query

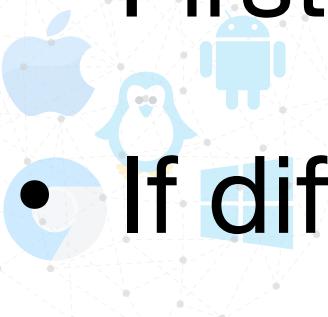
- When ARP query is sent initially, it is broadcast to all the nodes in the network
- The request includes senders IP address and MAC address
- It also includes the target IP address
 - Destination MAC is set as **FF:FF:FF:FF:FF:FF** (Reserved to send packet to all in the network)
 - If different network then send to the IP address of the gateway router
- All the nodes will have an ARP cache or ARP table
 - It stores the mapping, when the initial request is send from one host, all other hosts stores the incoming mapping as well



ARP Working



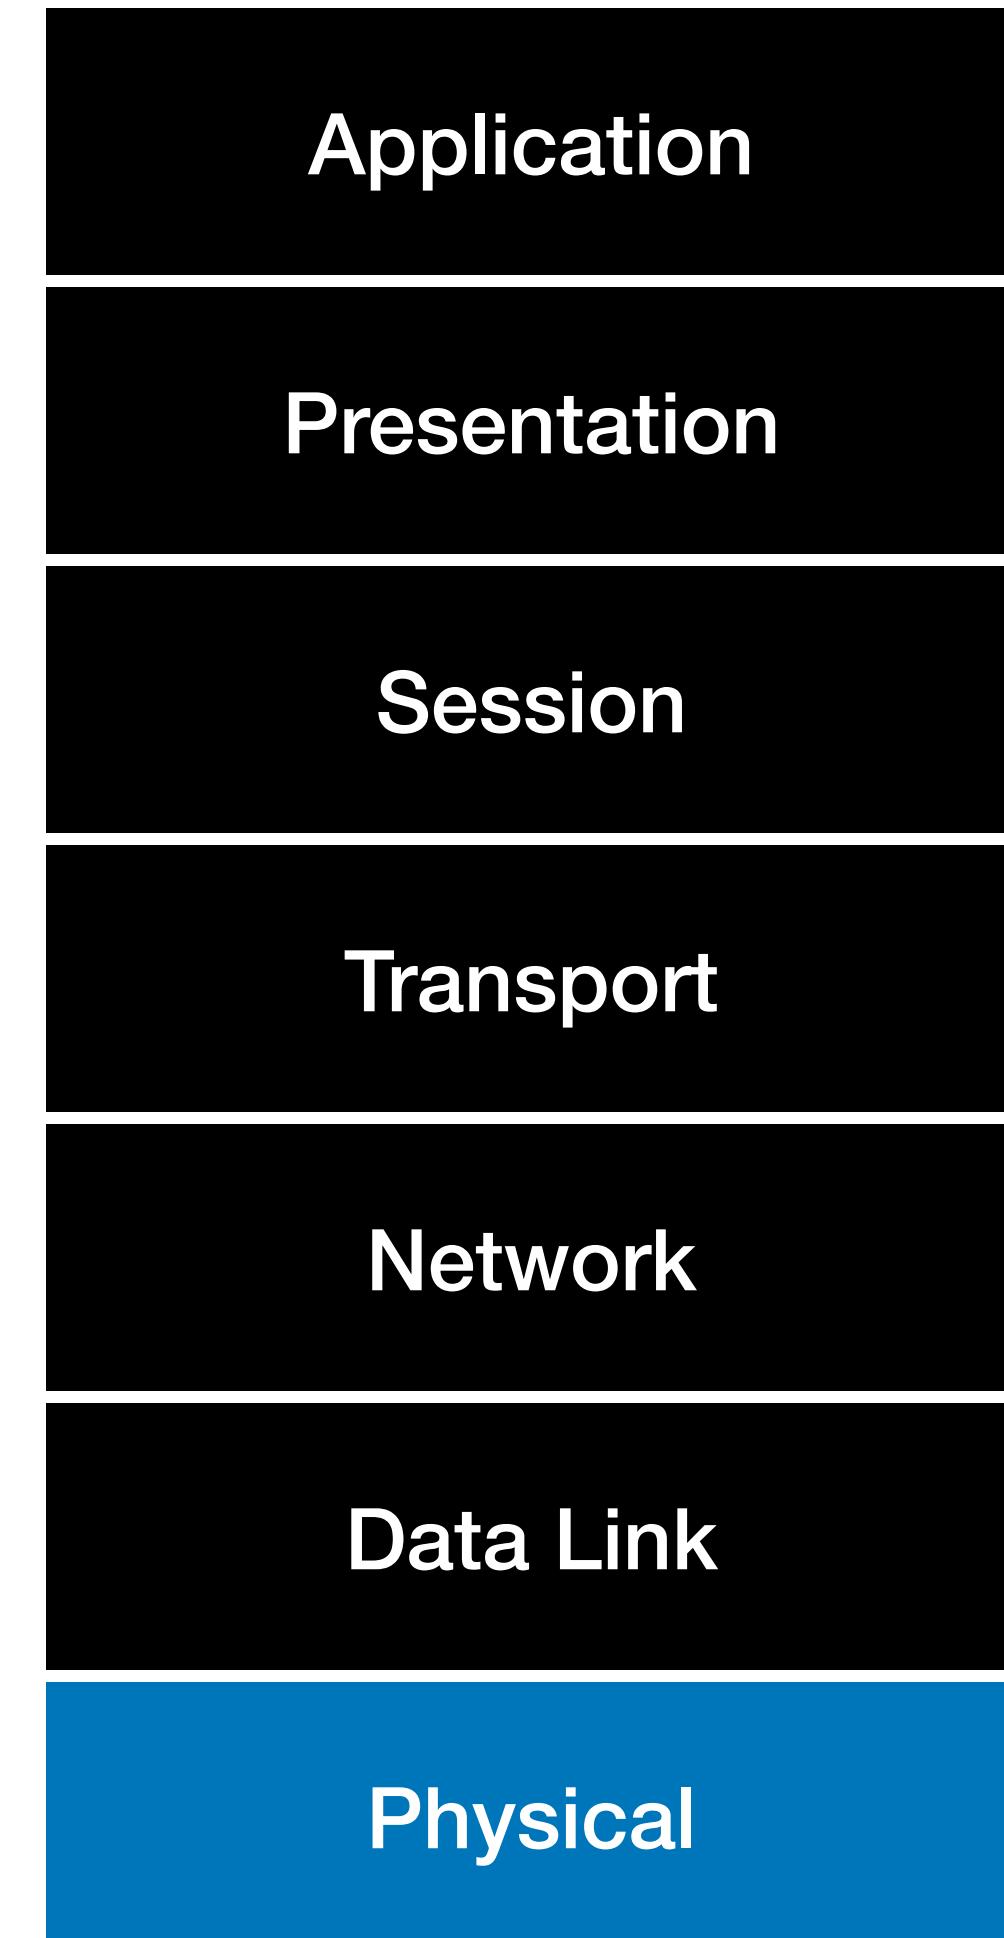
- The ARP process needs to happen only once, since router is the gateway
- First step - Check if the IP of the receiver is in the same or different network
- If different network => Send ARP to gateway else, send ARP to all nodes in the network (FF:FF....:FF)



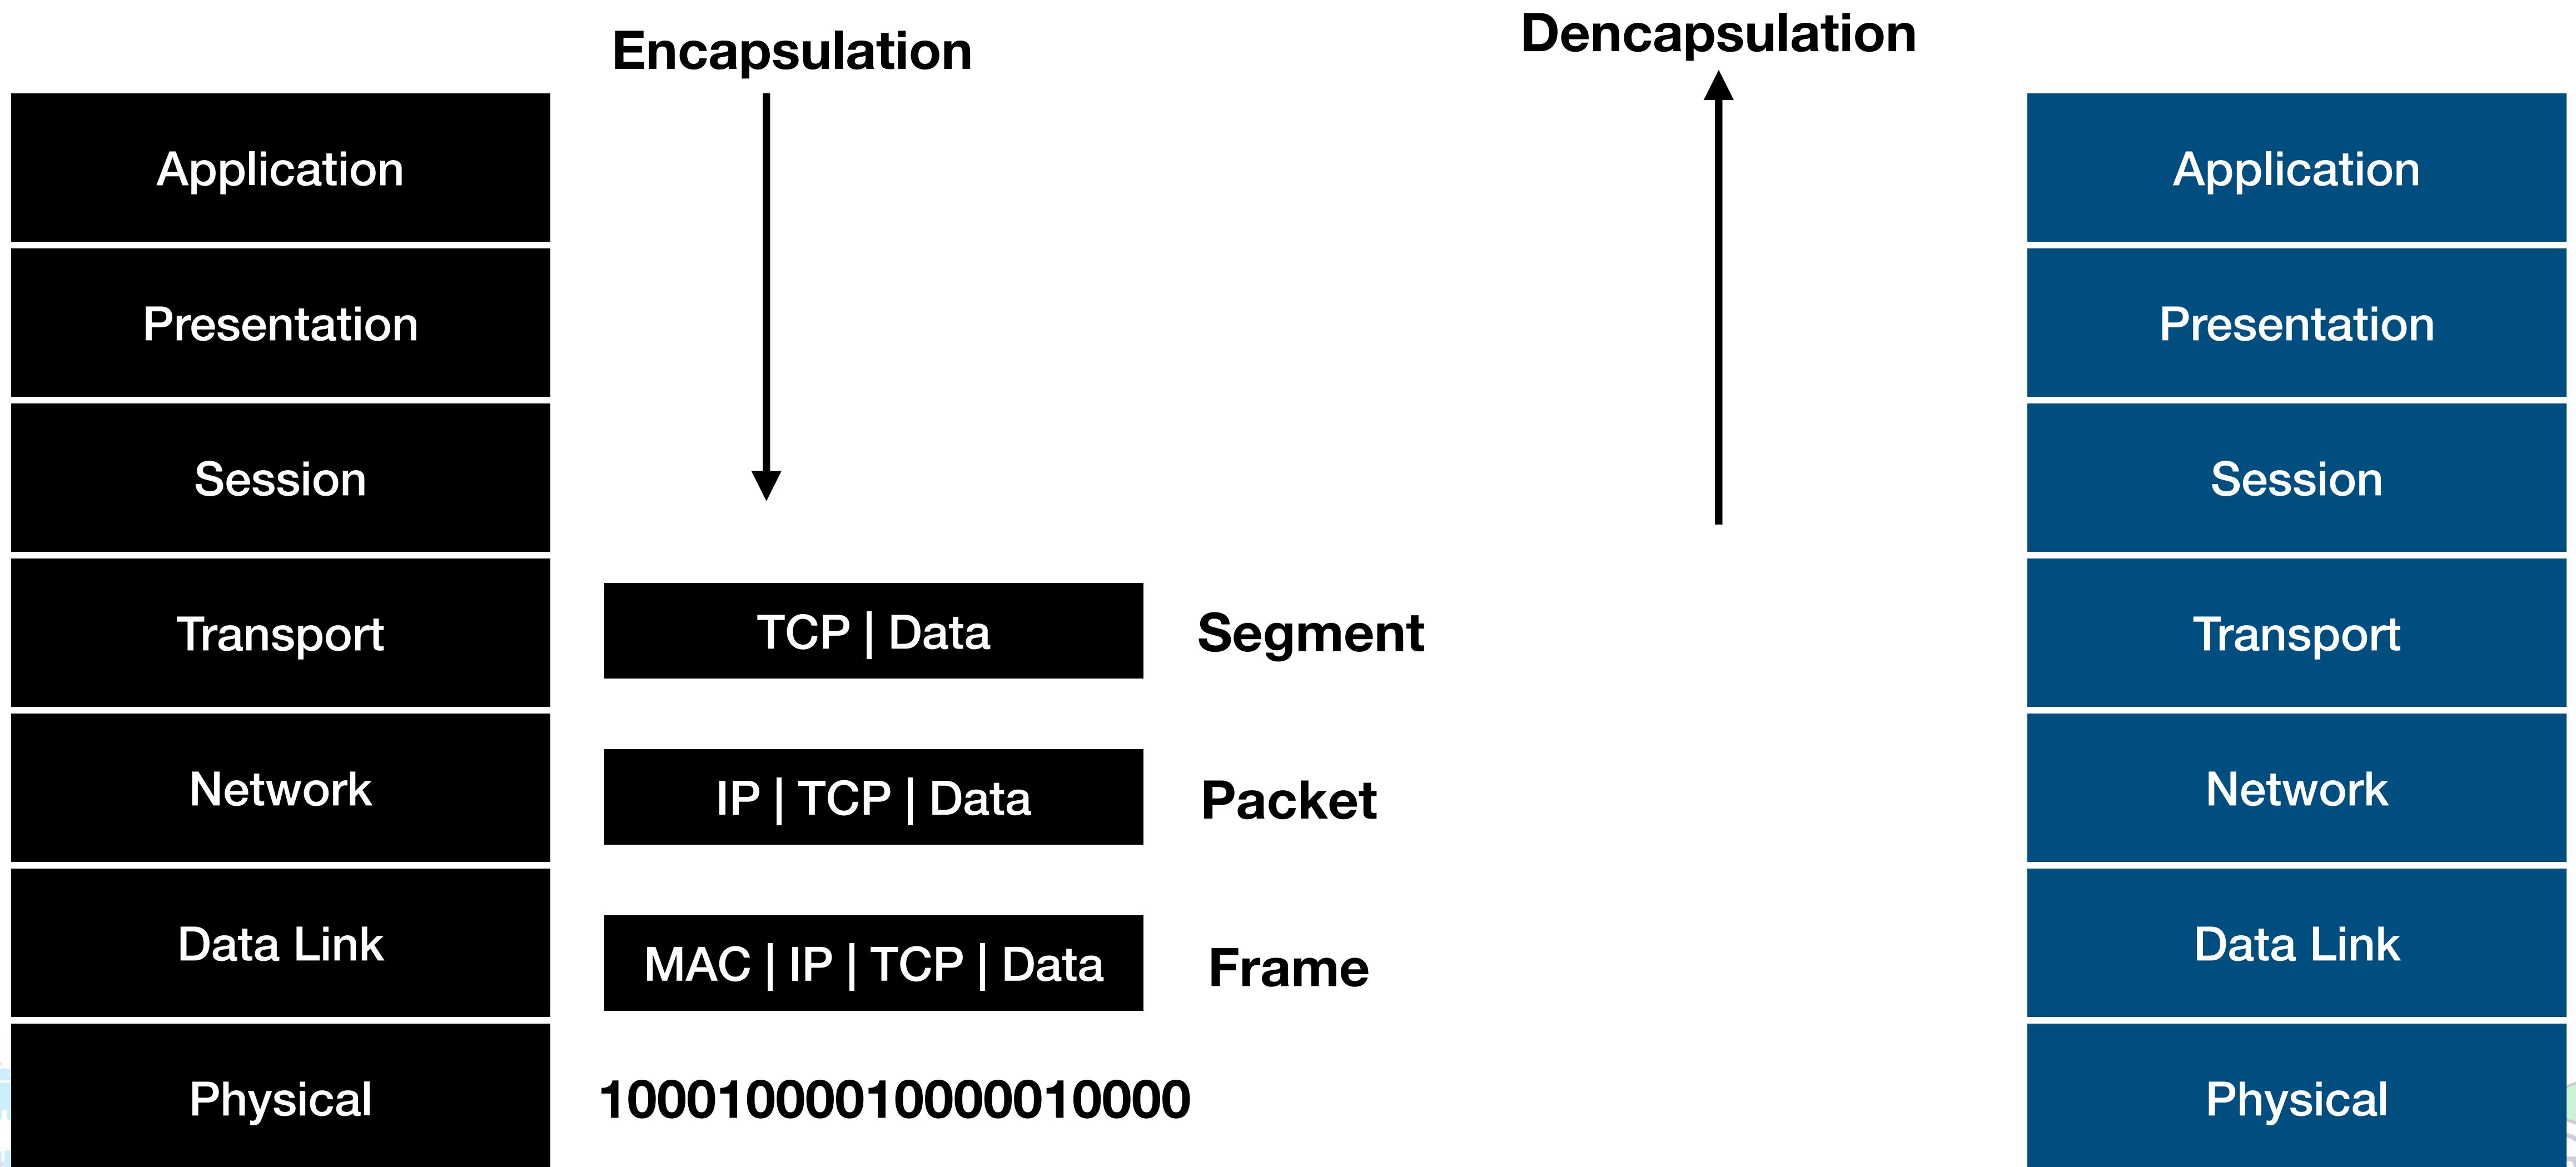
Physical Layer (L1)

Ultimately everything is 0's and 1's

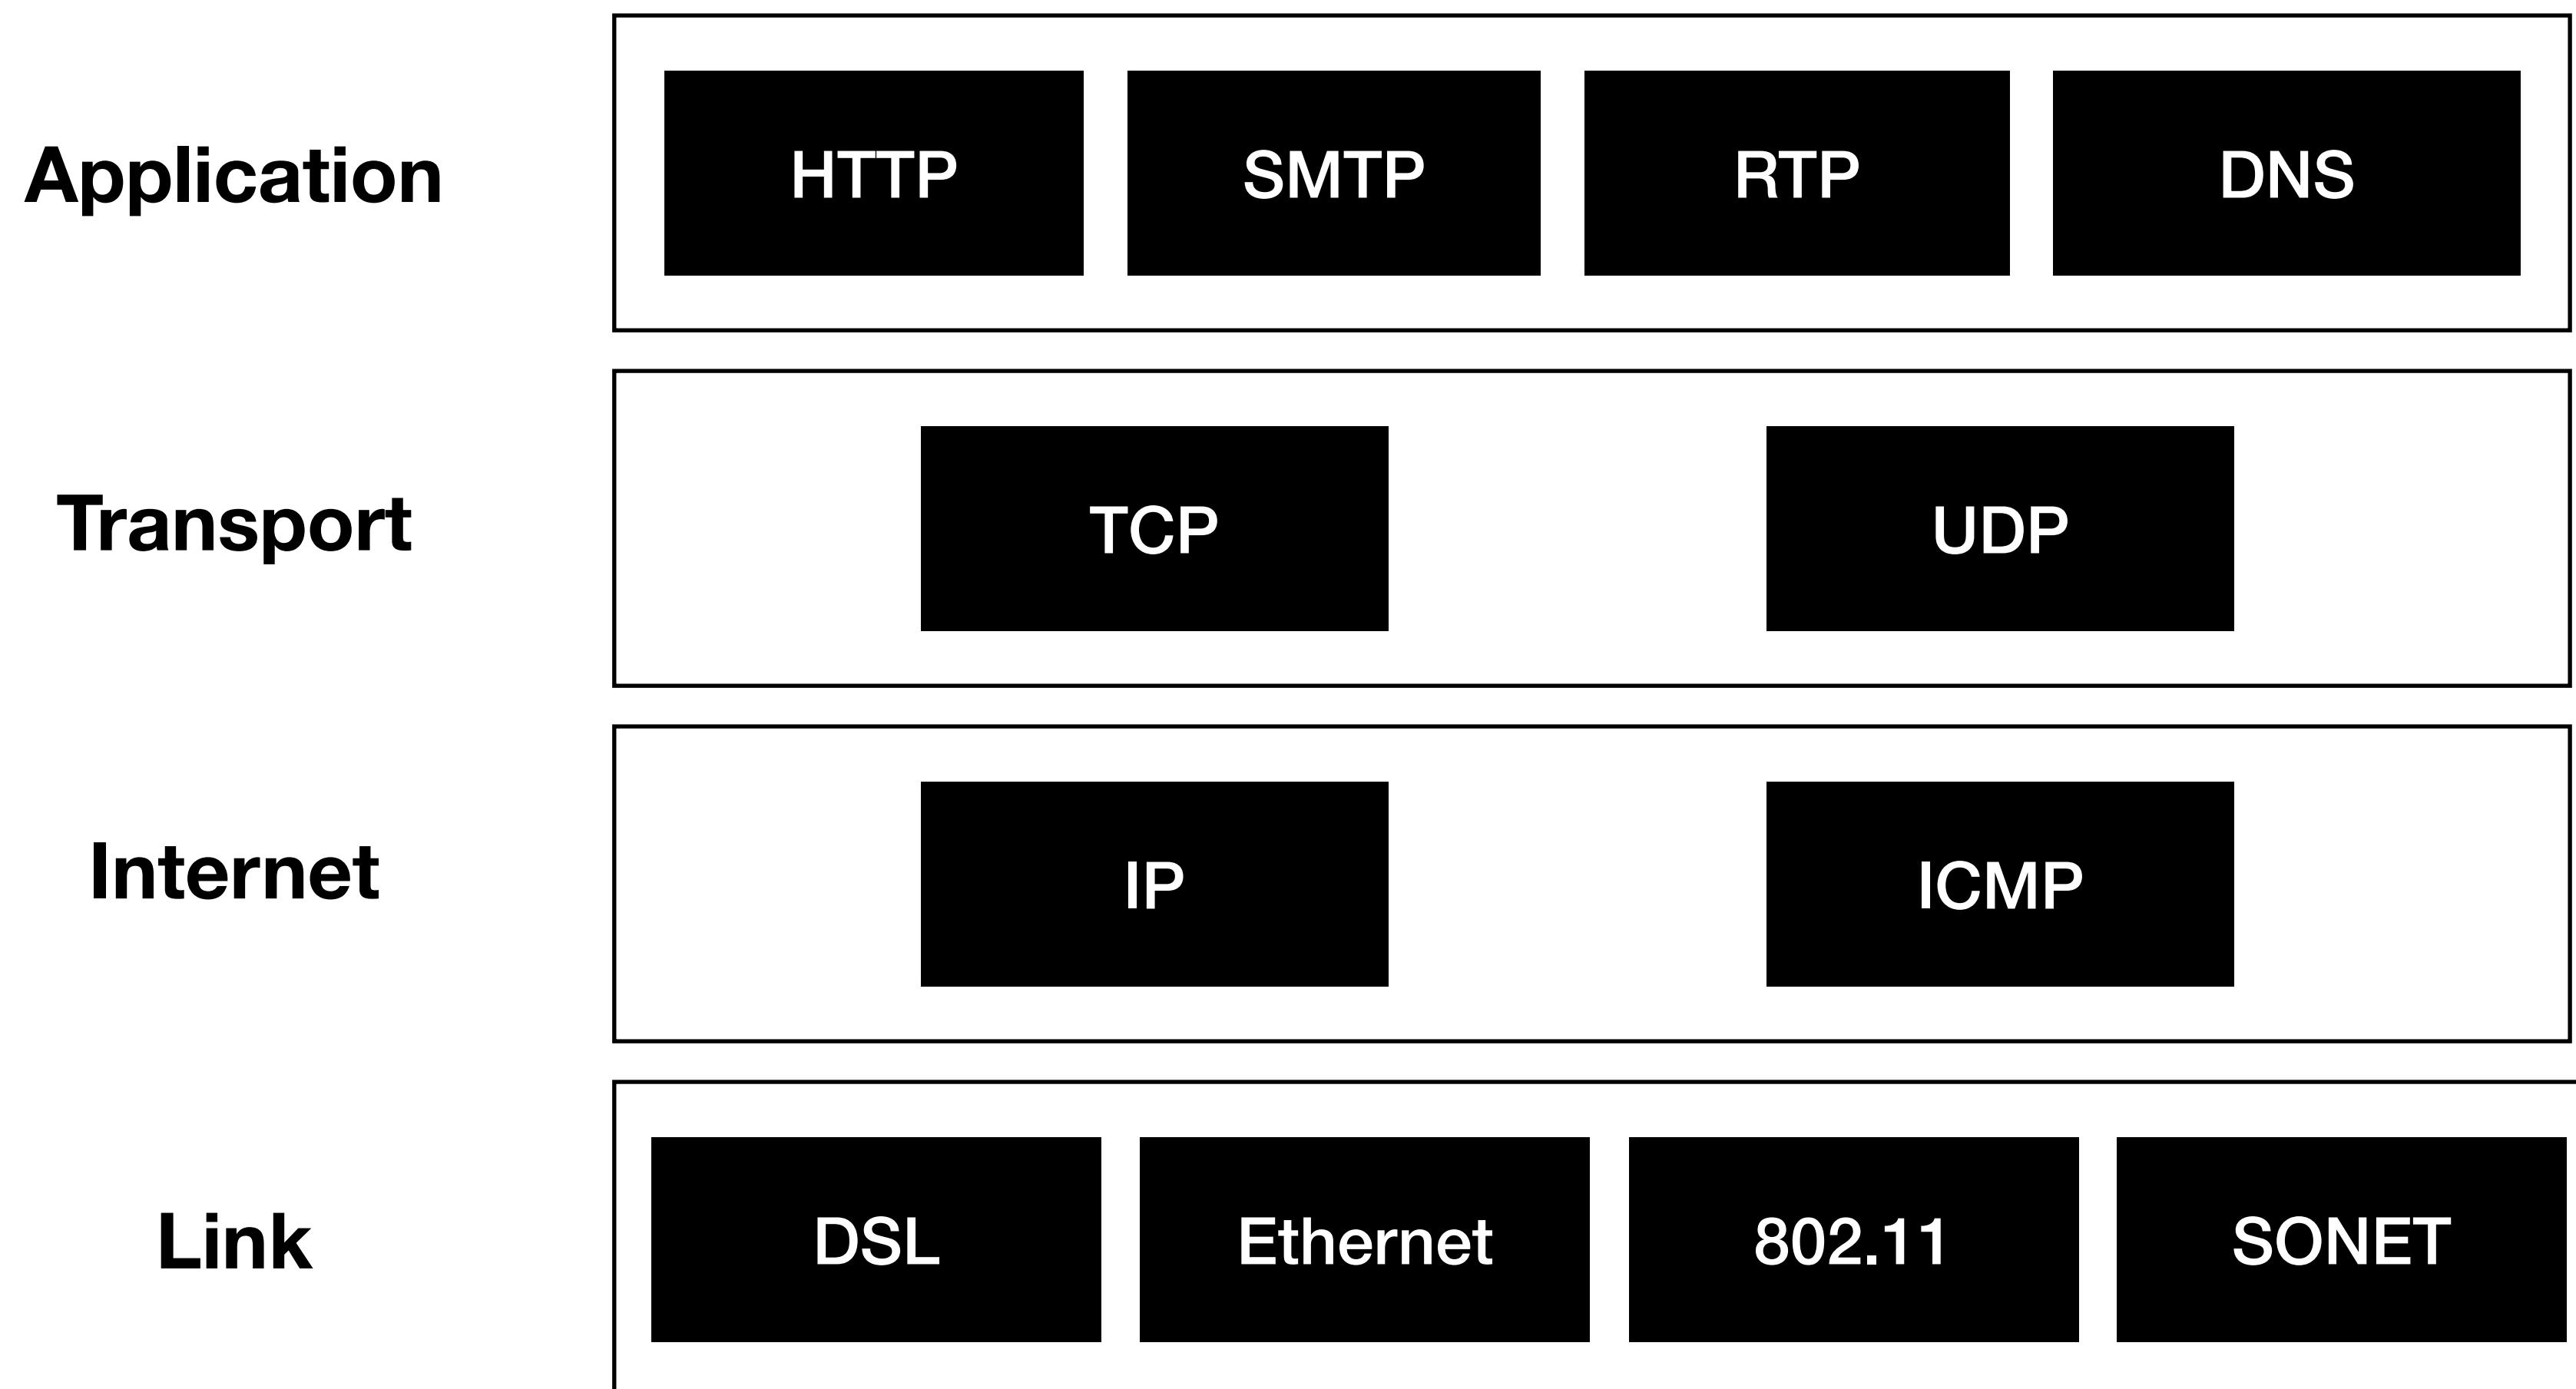
- Data is in the form of bits - 0s and 1s
- Something has to transport the bits from one machine to another - Physical layer
- Concerned with transmission of raw bits over physical medium, like a cable
- L1 technologies: Ethernet cables, Optical fiber, Coaxial cable, etc.
 - Even WiFi is L1 technology, hub, repeater, etc.



End-to-End OSI View



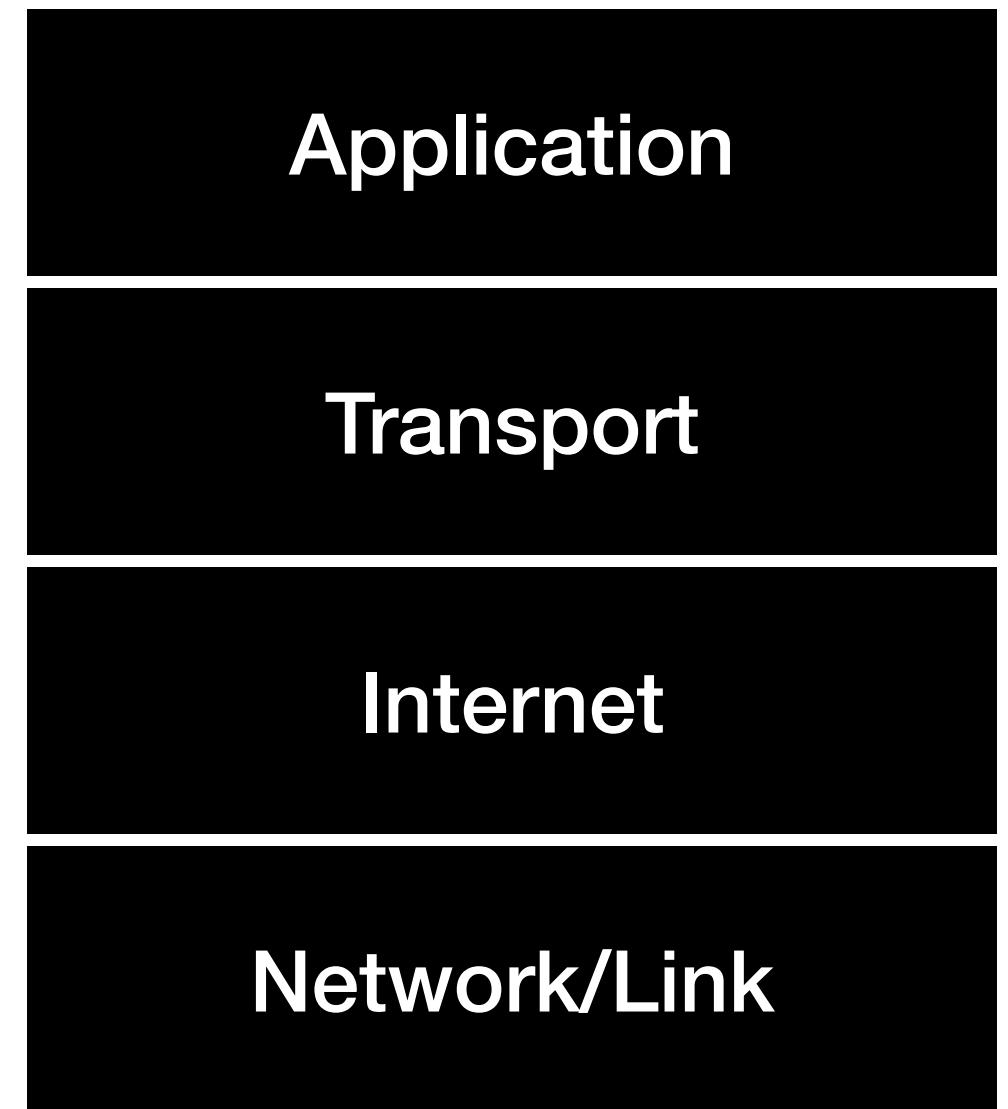
Network Protocol Stack



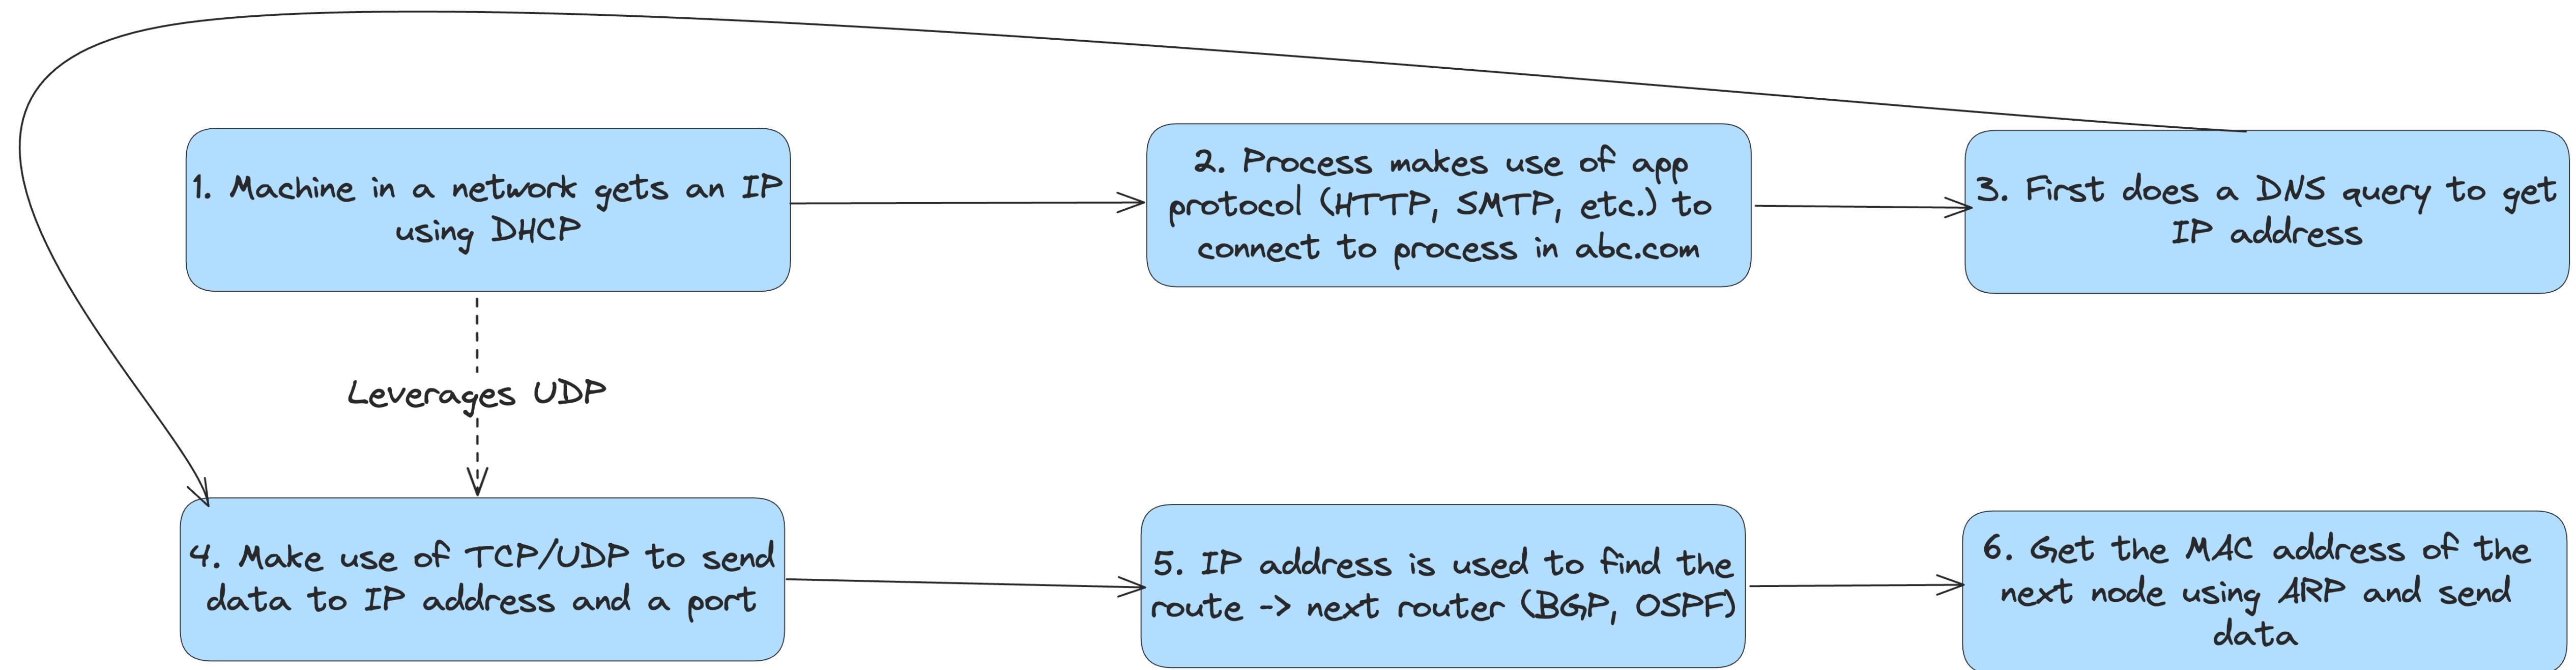
The 4 Layer Model

Internet Model or TCP/IP model,

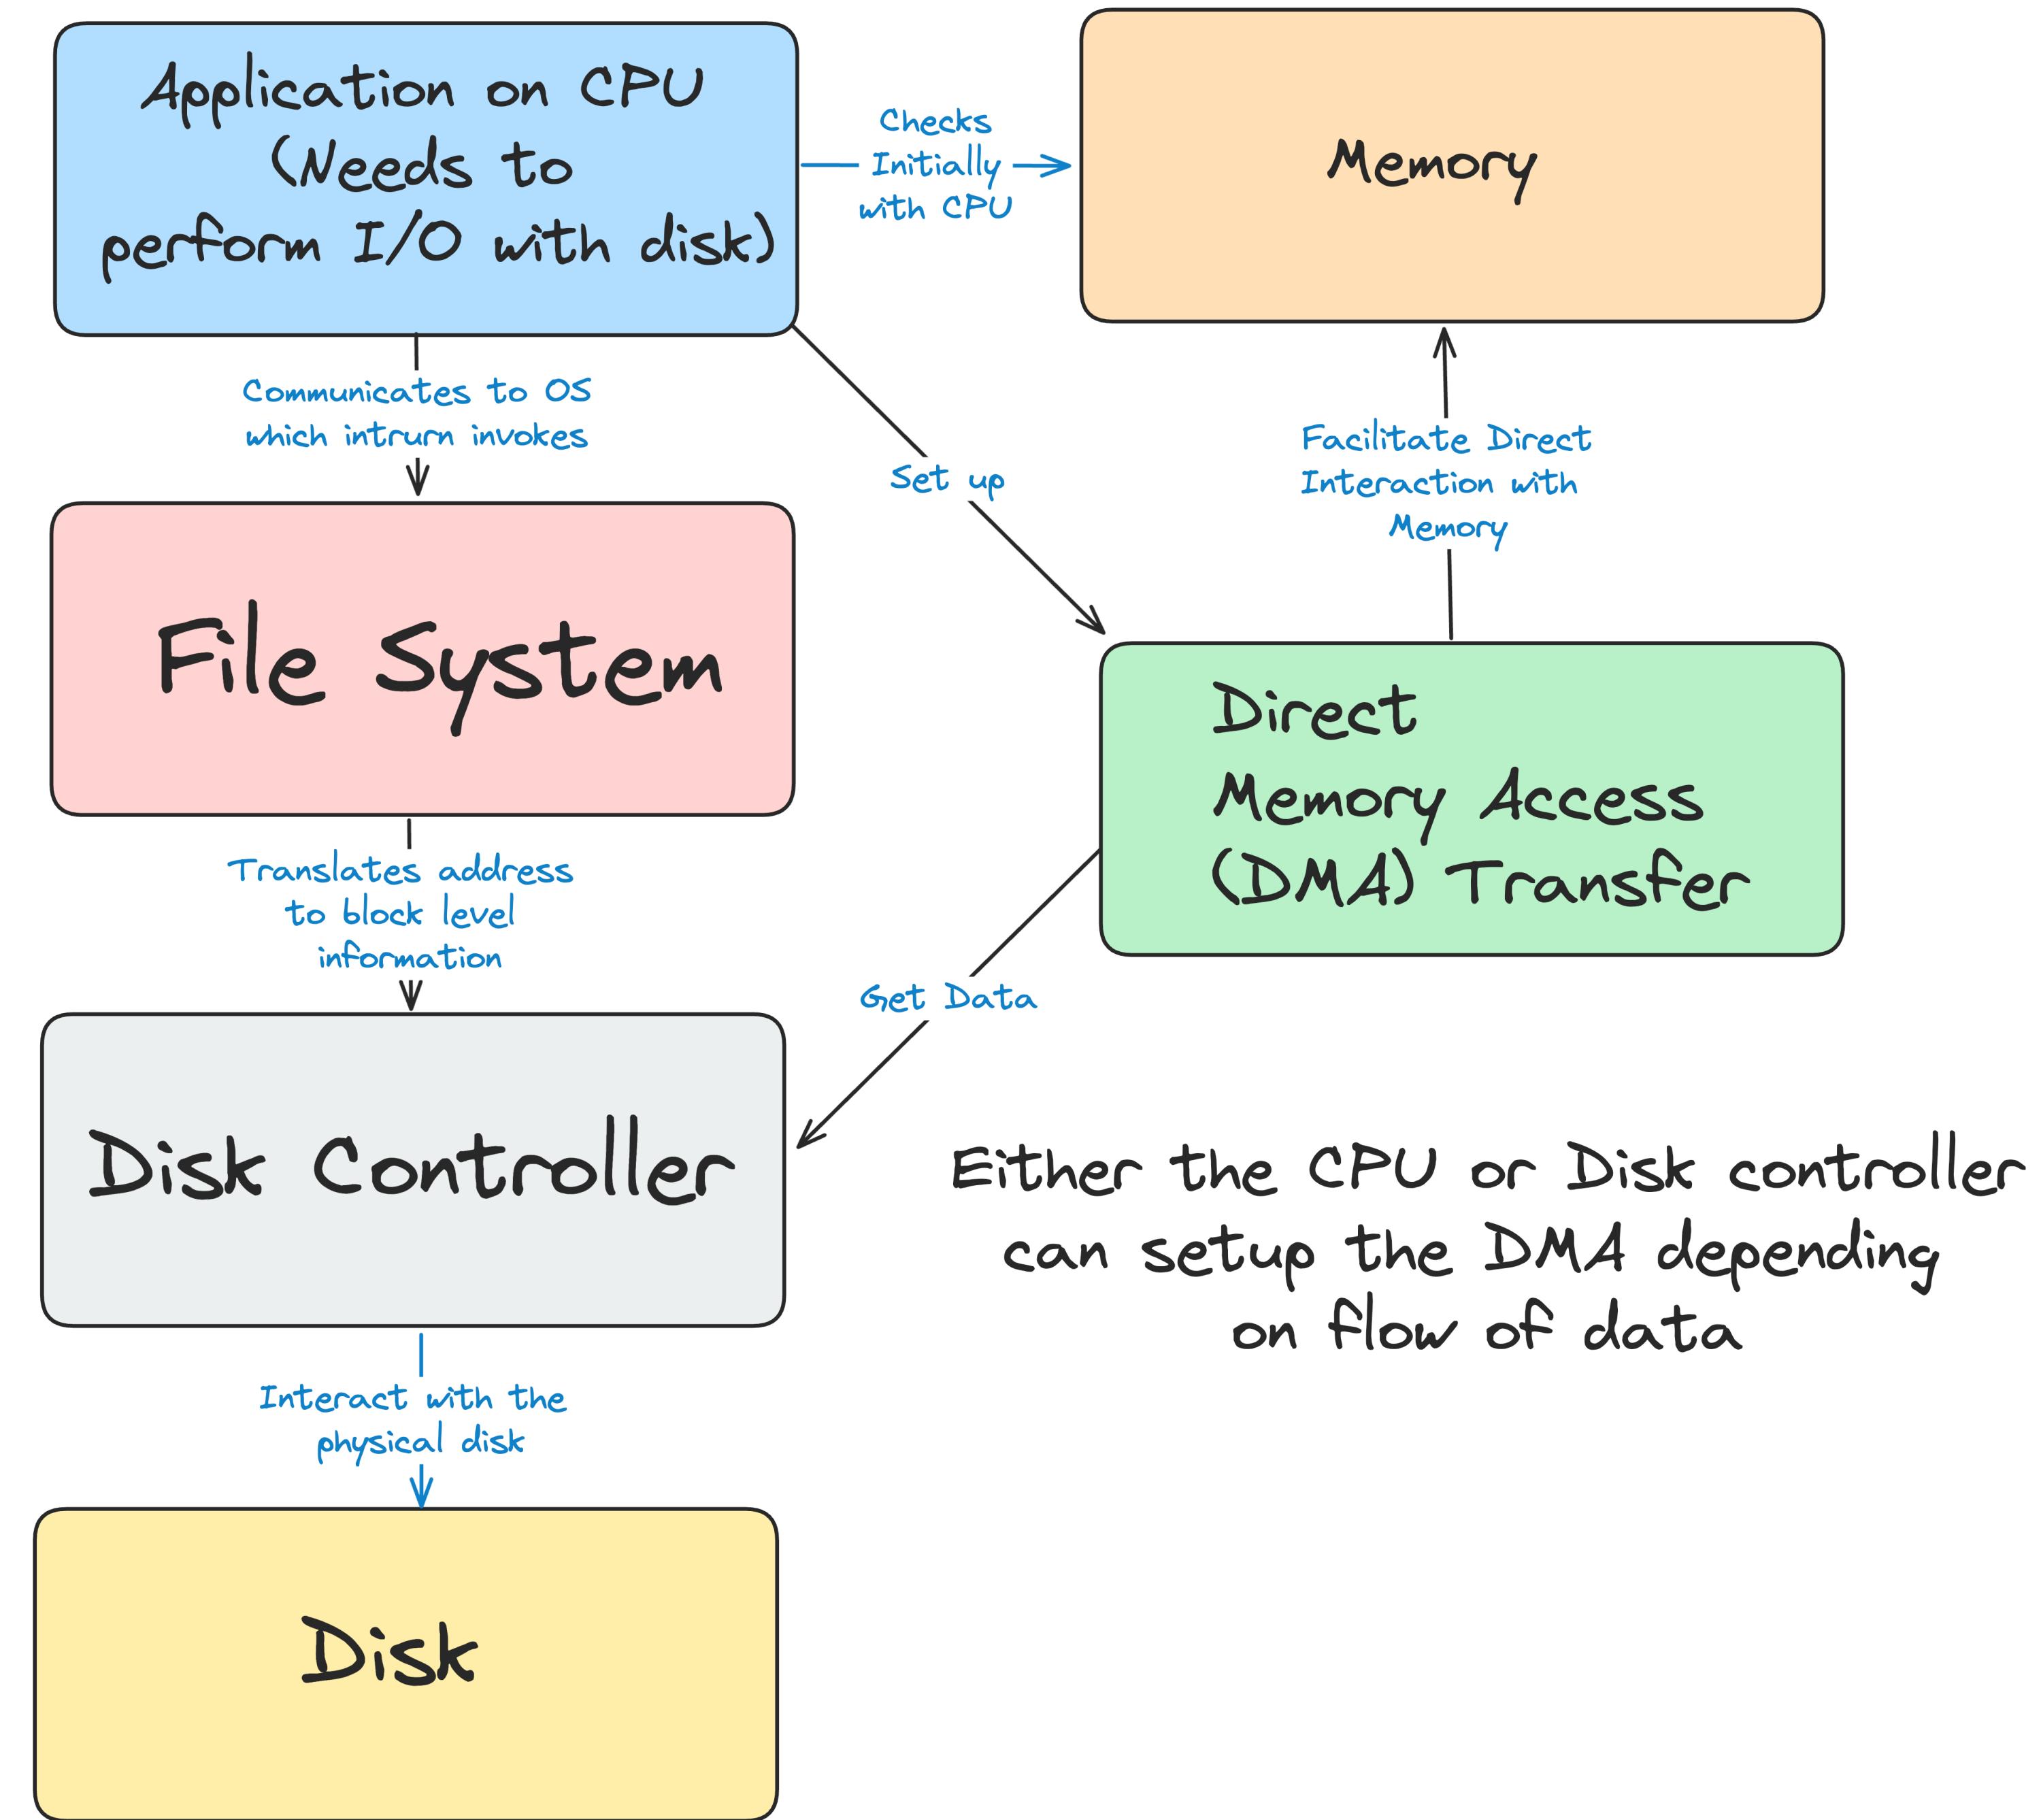
- OSI model is more educational purpose
- 4 layer model more used in reality
- Application layer - Corresponds to application, presentation and session
- Transport layer - Transport layer of OSI
- Internet layer - Network layer of OSI
- Network - Physical and data link layers of OSI



Putting it together

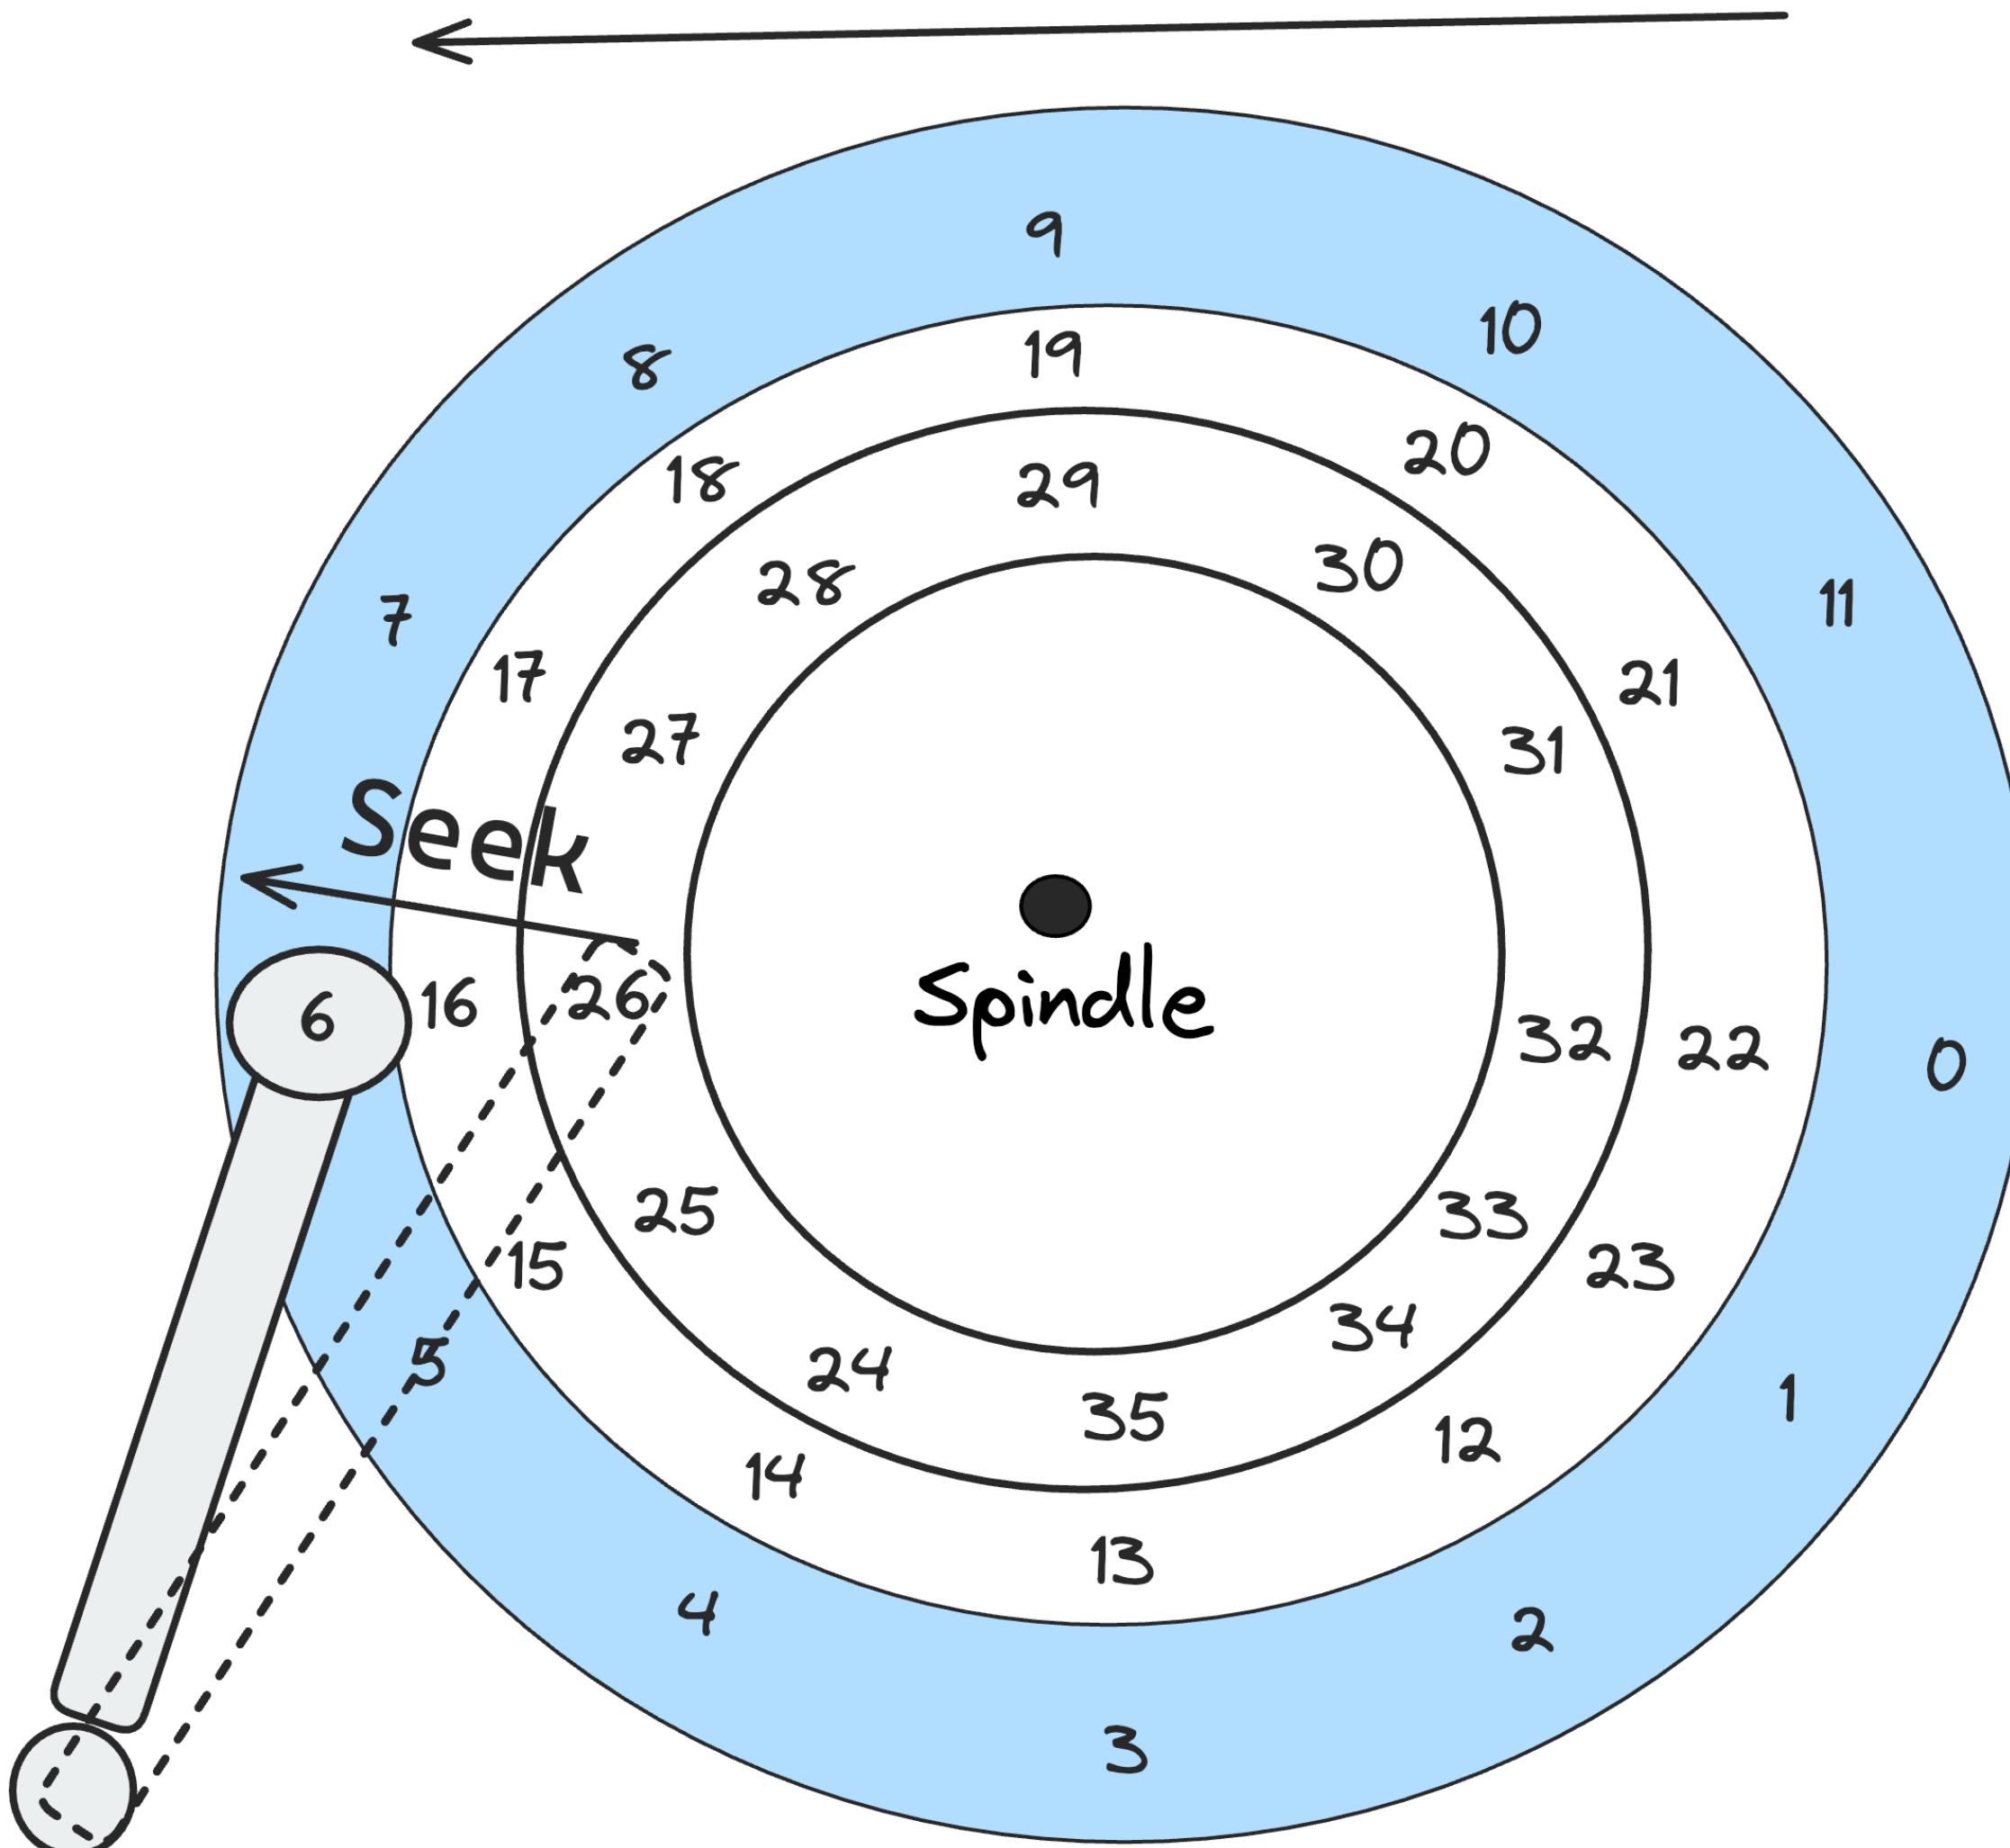


Persistence



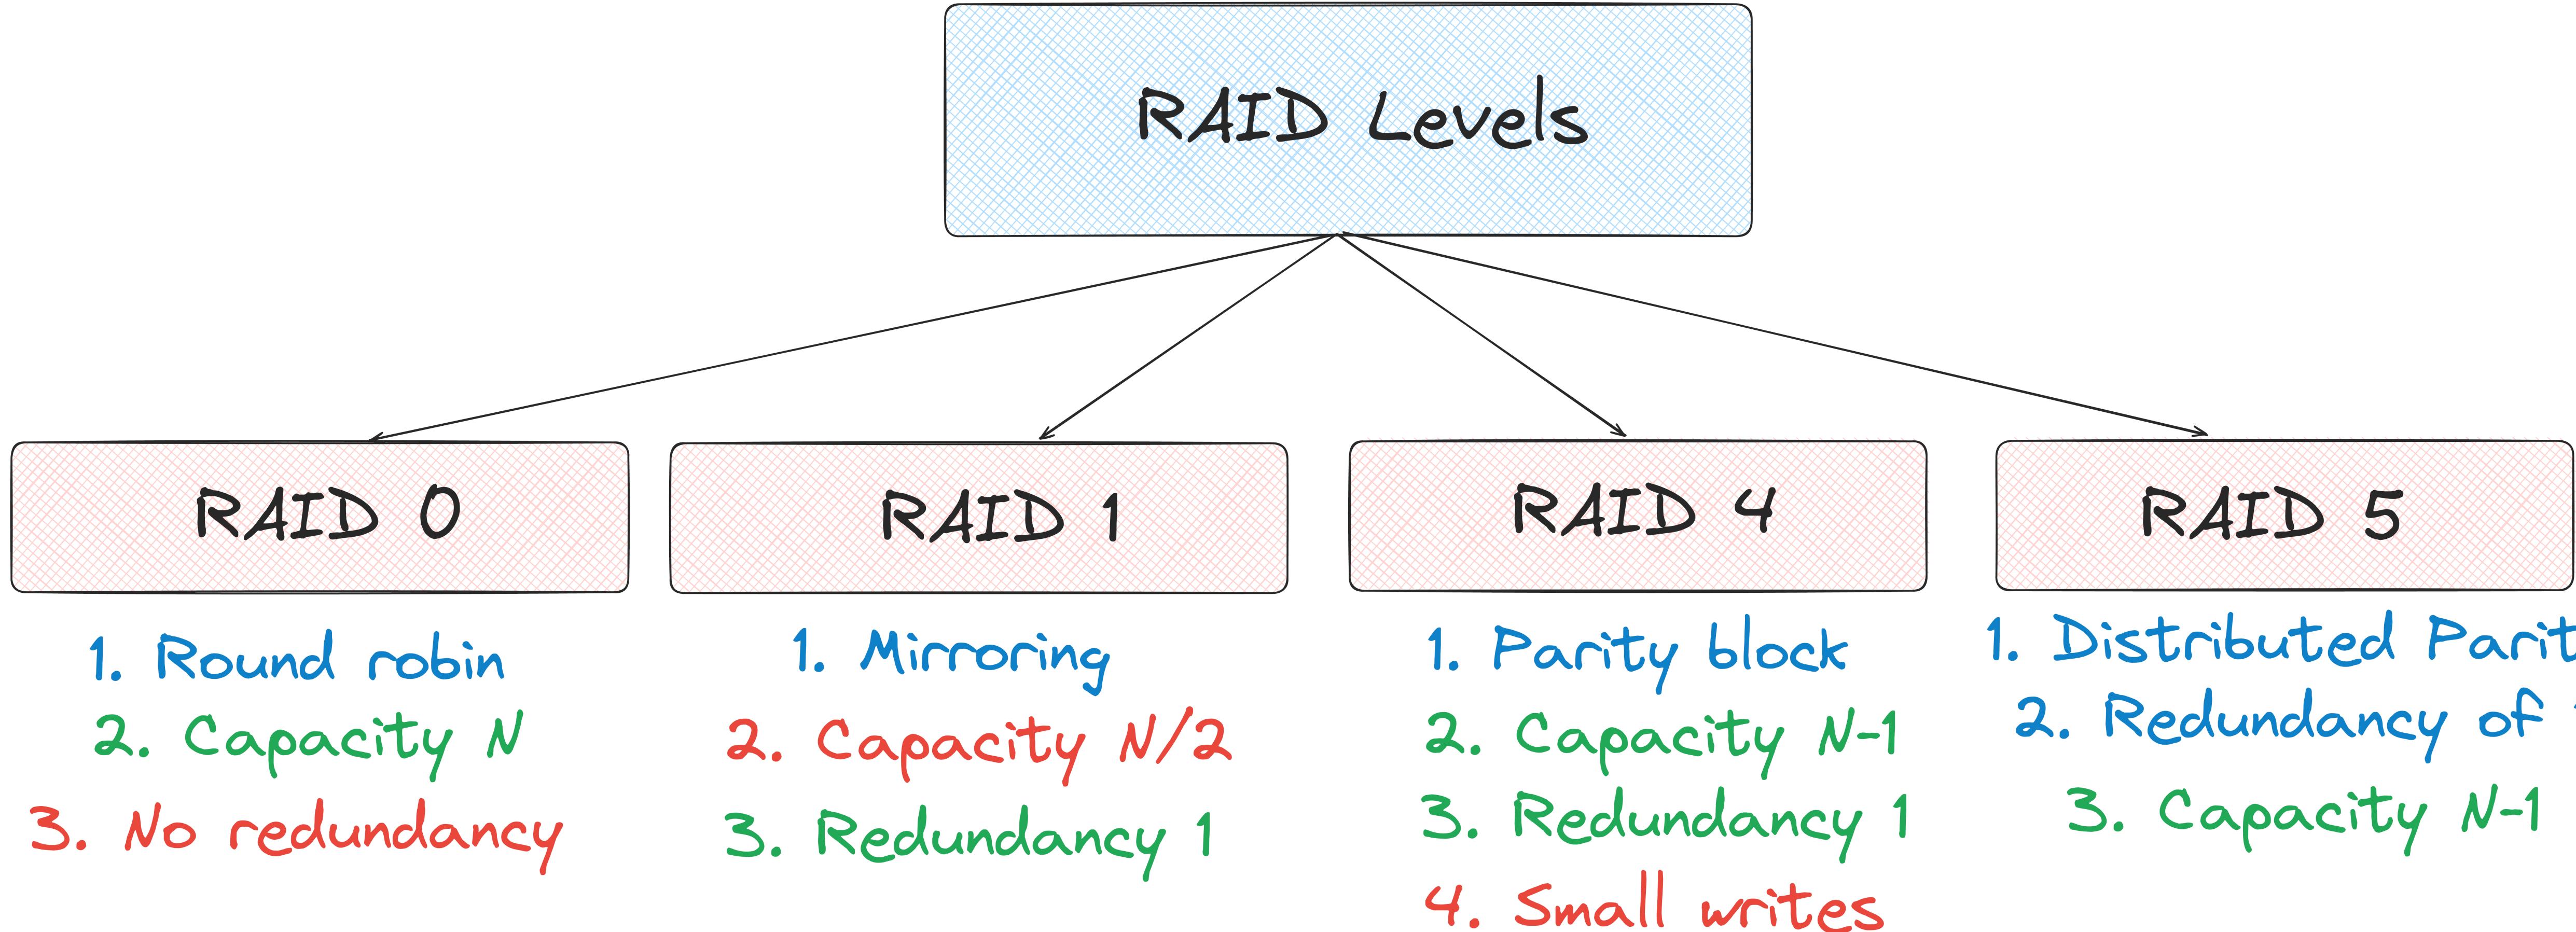
Disks: An Overview

Rotates this way



- Disk rotates on a spindle
 - The arm can move across (seek) or stay as the disk rotates
 - The head is used to read/write
- Data is arranged in tracks as blocks/sectors
- There are 100s of tracks on a single disk
- **Seek, rotate and transfer** - three key phases

RAIDs



Breaking down into two main aspects

- We worked on building a **Very Simple File System (VSFS)**
- In any FS, two key things make the difference

Data Structures

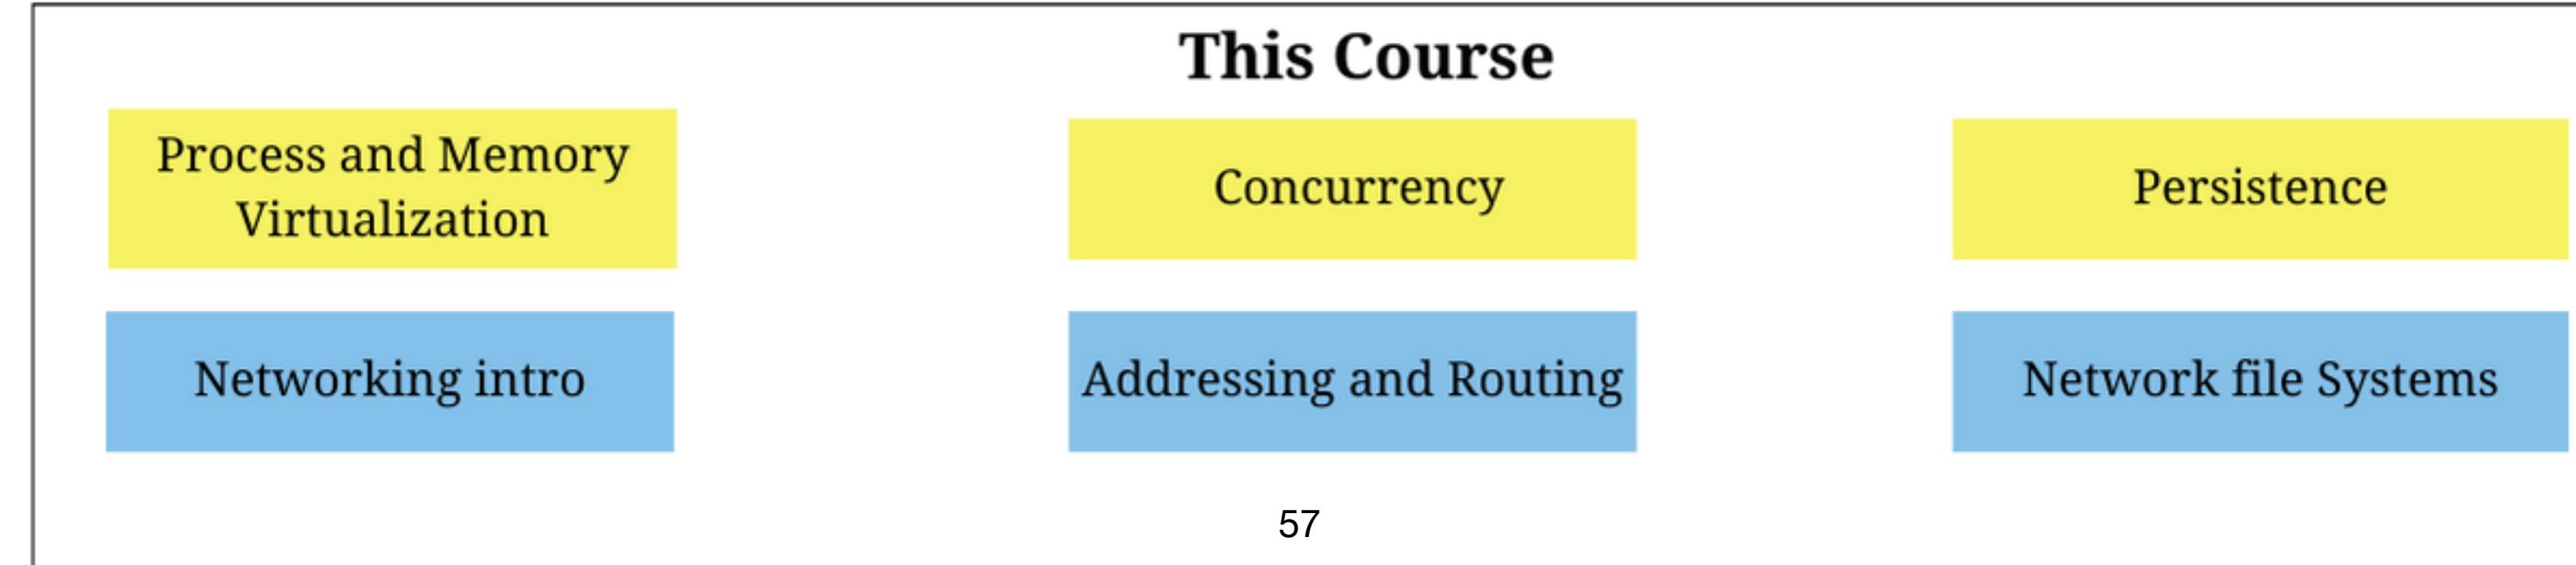
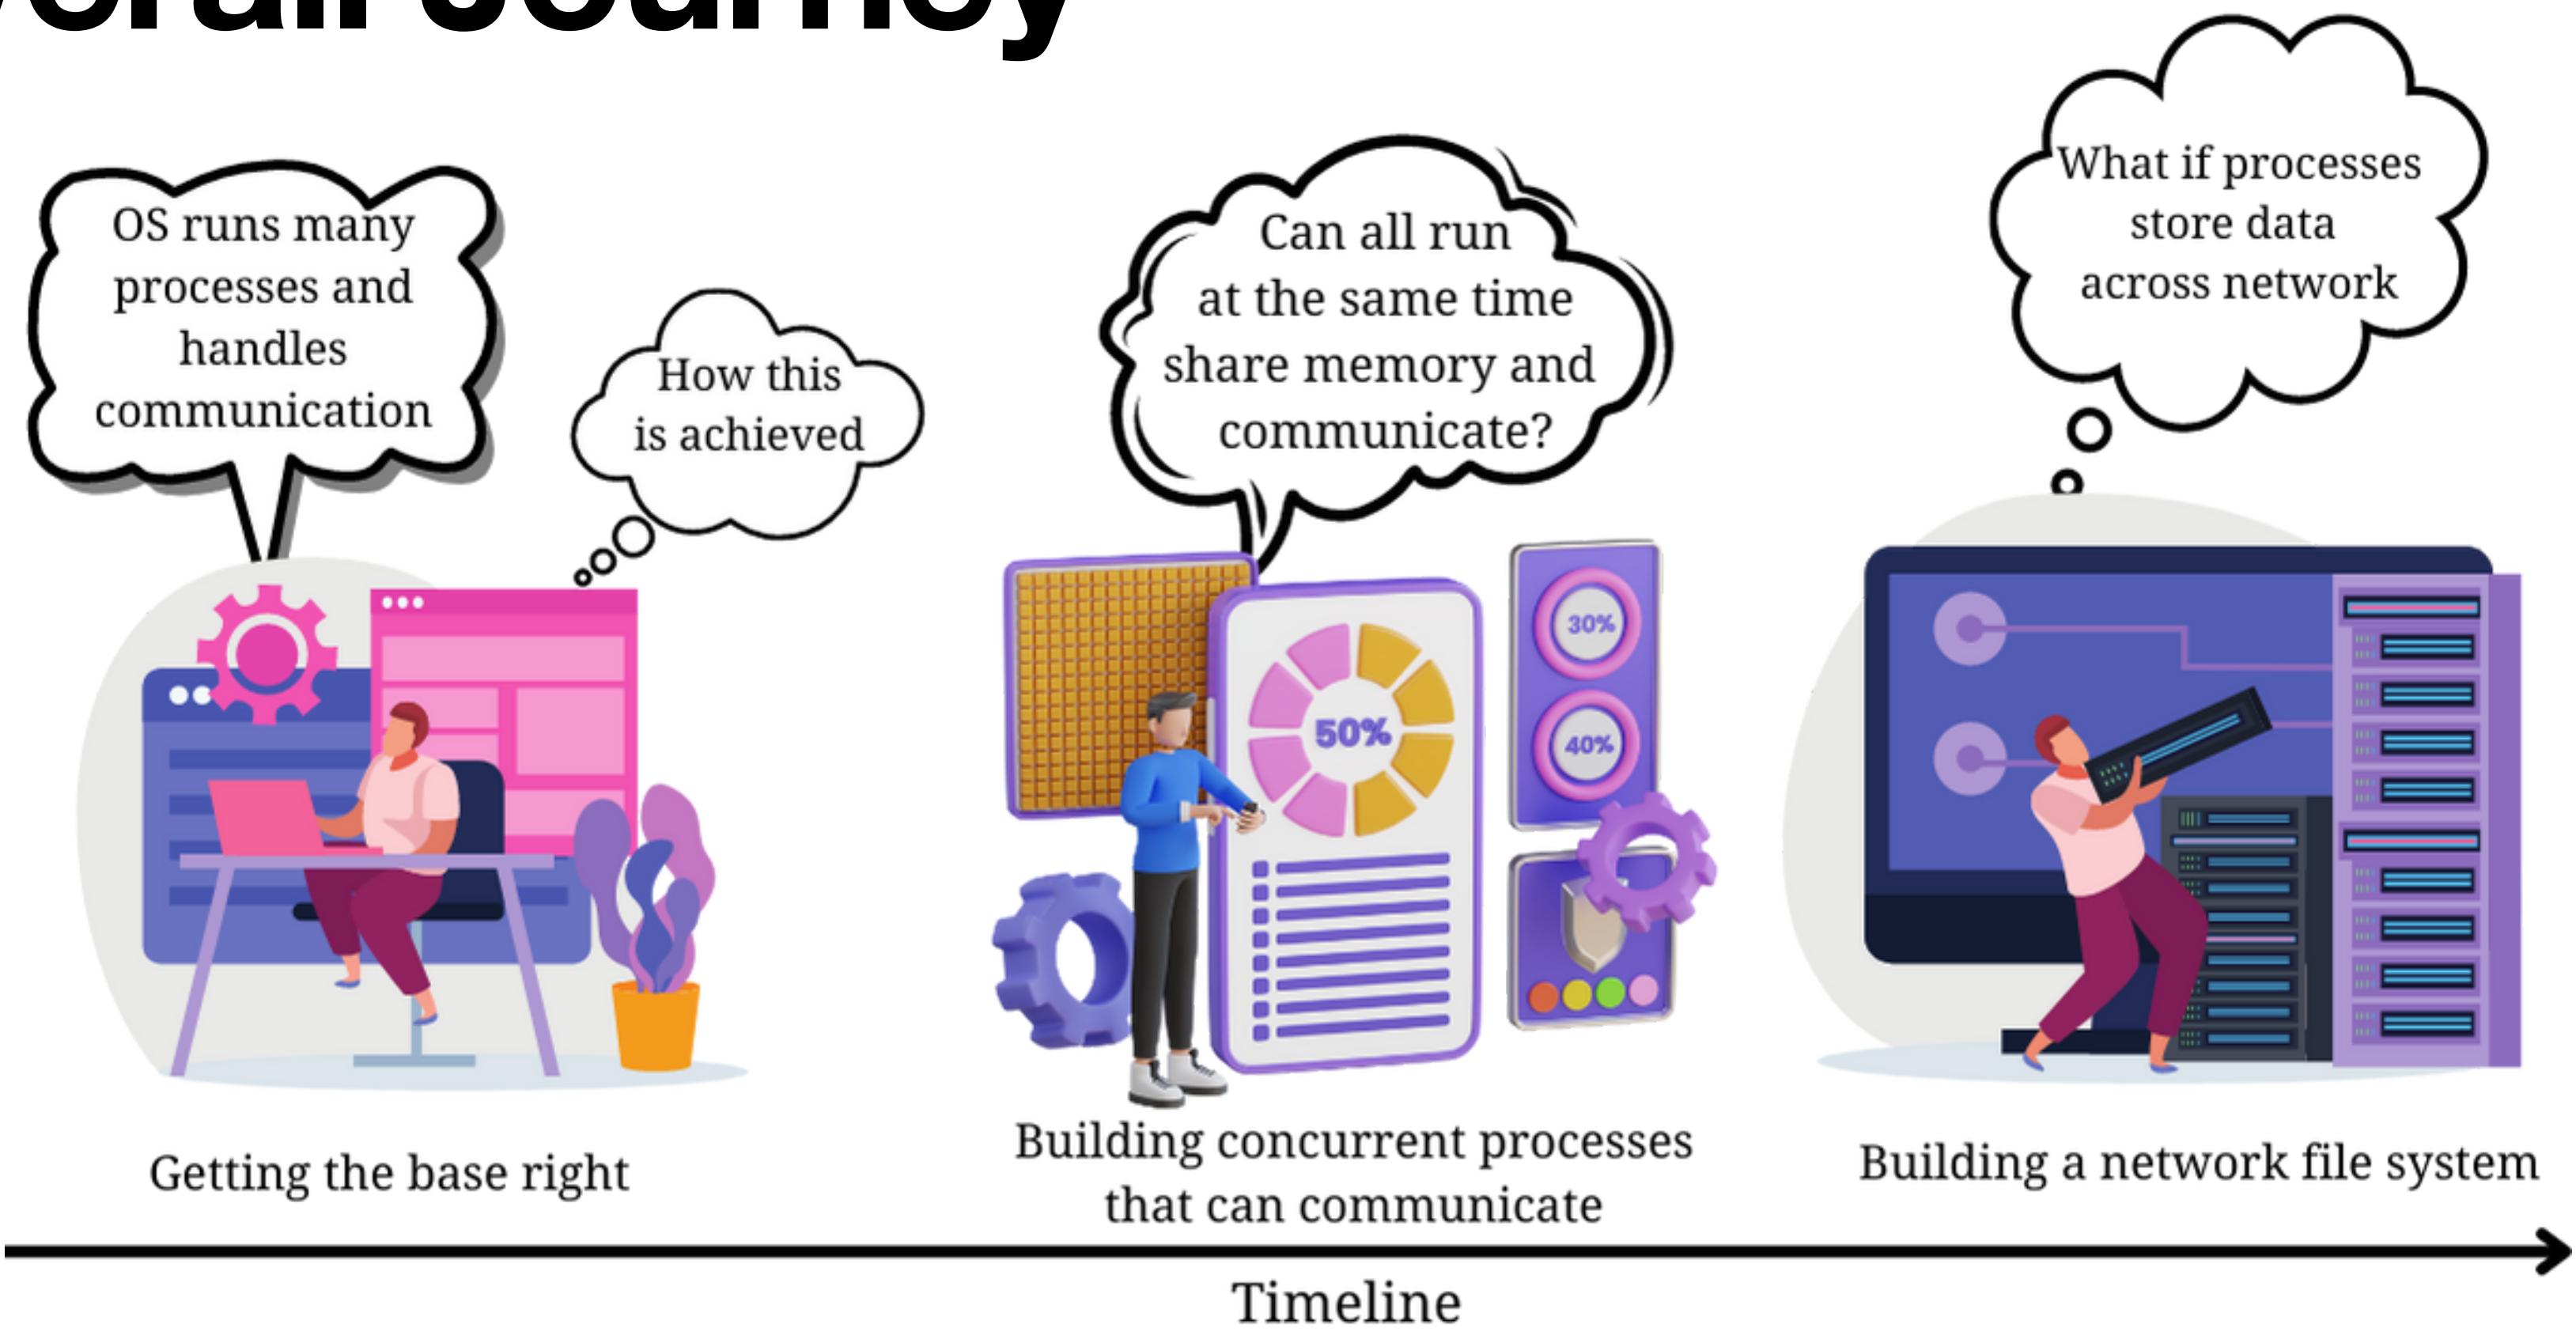
- Inode - Data structure for each file
- Store inodes, data, mapping to inodes, etc in a large array

Access Methods

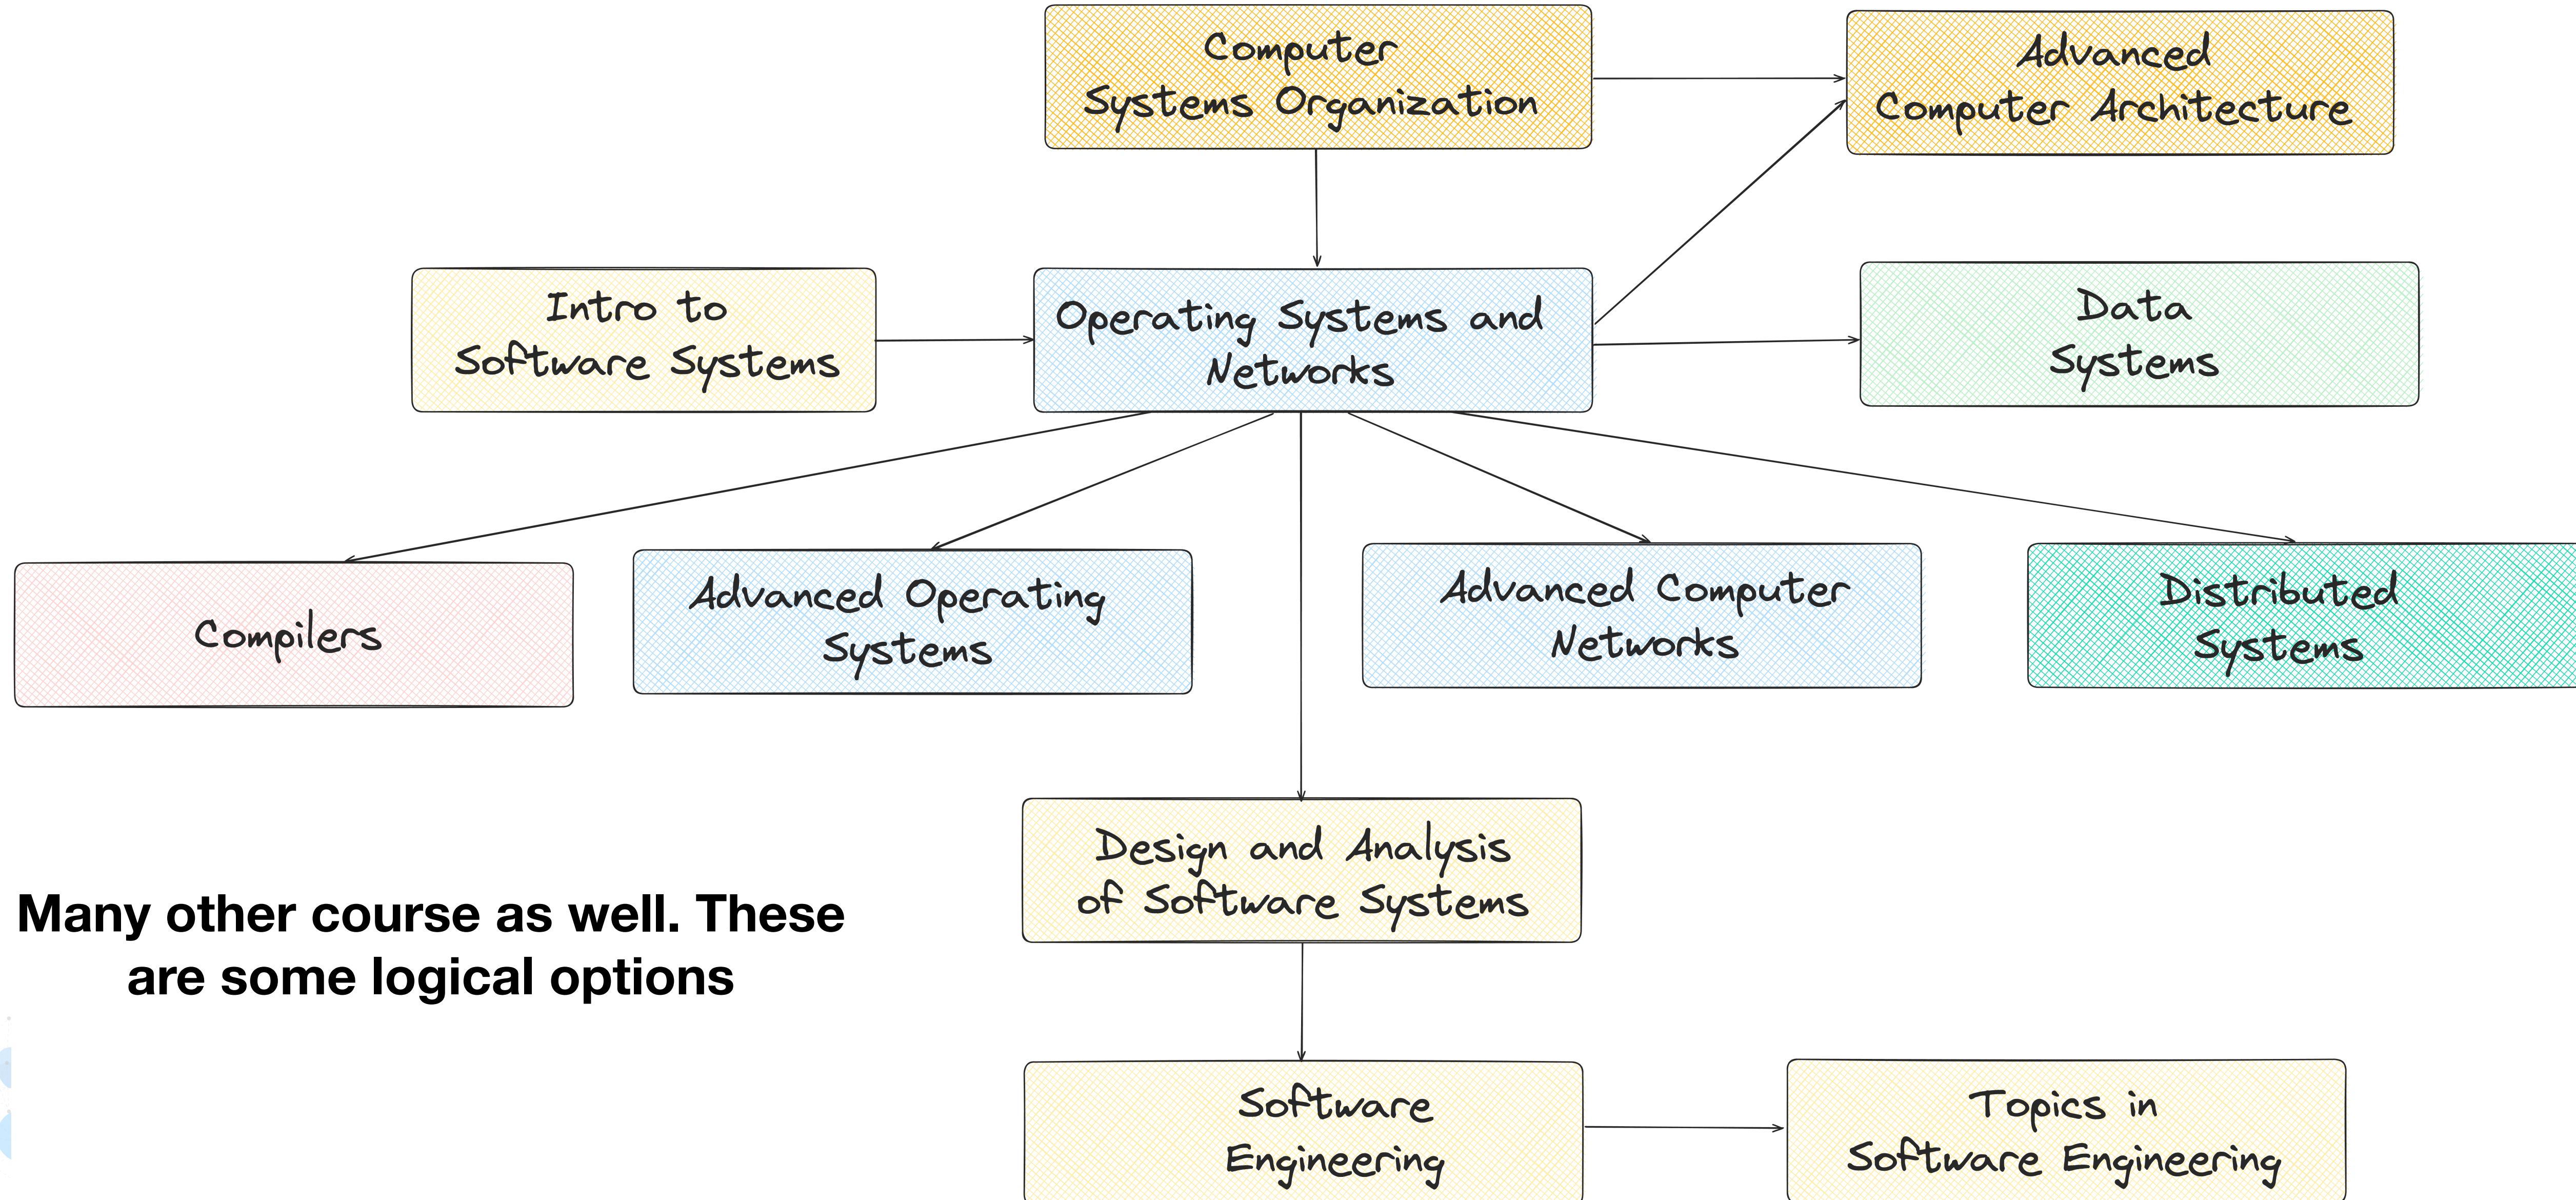
- Start with the root
 - Traverse through the path using inode mapping
- Caching can be used to improve efficiency



The overall Journey



What next?



Course Restructuring

- **Restructured little more compared to last year! - Course Calendar is more intact!**
- Changes in course logistics and planning
 - Modified the grading scheme
 - **Modified groups in final project**
 - 2 Mini projects - A Lot of changes
 - Almost every project had a network component
 - OS + Networks were kept **more intertwined**
 - Regular meetings with TAs - Increased tutorials and TA meeting hours
 - Feedbacks are always welcome!!



Thanks to the team!



Akhila Matathammal



Ananya Halagatti



Anirudh Vempari



Aviral Gupta



Karthik Vaidhyanathan



Divijh Mangtani



Eshaan Sharma



Kriti Gupta



Prakhar Jain



Prasoon Dev



Miryala Sathvika



Shlok Sand



Shubham Goel



Varun Gupta



Thanks to all of you!





Thank you

Course site: karthikv1392.github.io/cs3301_osn

Email: karthik.vaidhyanathan@iiit.ac.in

Twitter: [@karthyishere](https://twitter.com/karthyishere)

